

Containerized In Action (From Docker to Kubernetes)

Siam Chamnankit Co.,Ltd

Who We Are



Joe

Software Delivery
Experience in Test (SDET)
at SCK, Specialty (Agile,
DDD, Java, Software
Design)



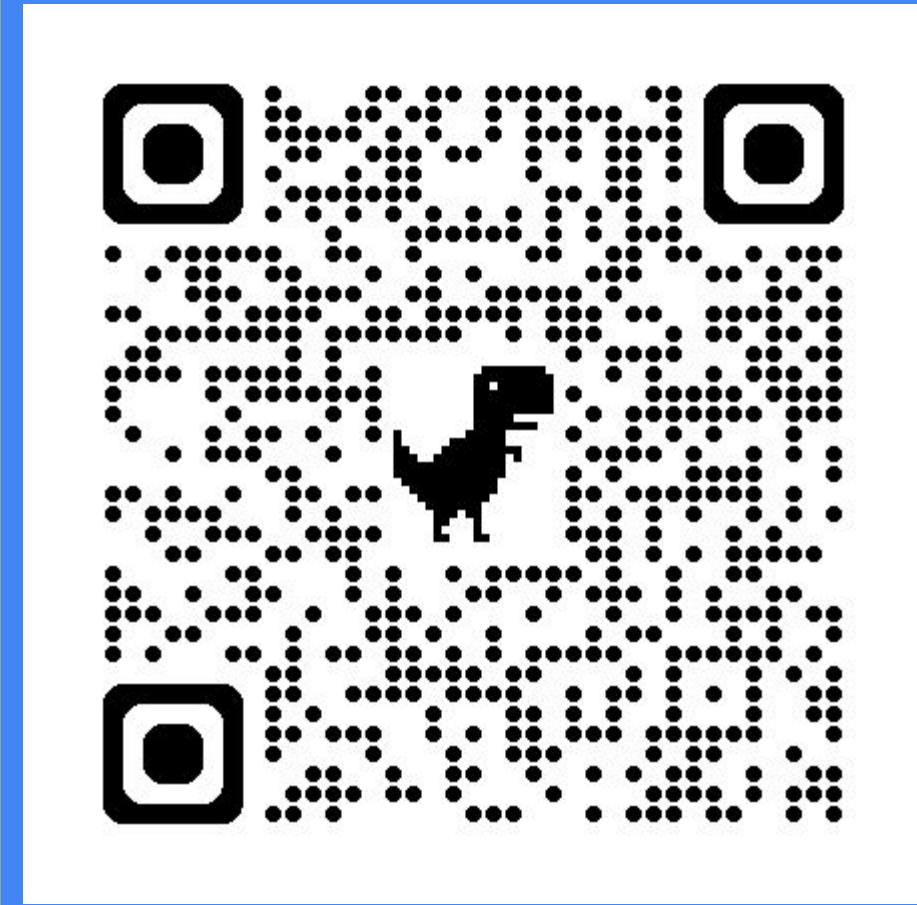
Meng

Software Delivery
Experience in Test (SDET)
at SCK, Specialty (DevOps)

Our Playground

<https://miro.com/app/board/uXjVKHU6Ggl=/>

<https://tinyurl.com/sckdojo>



Day 1

Agenda

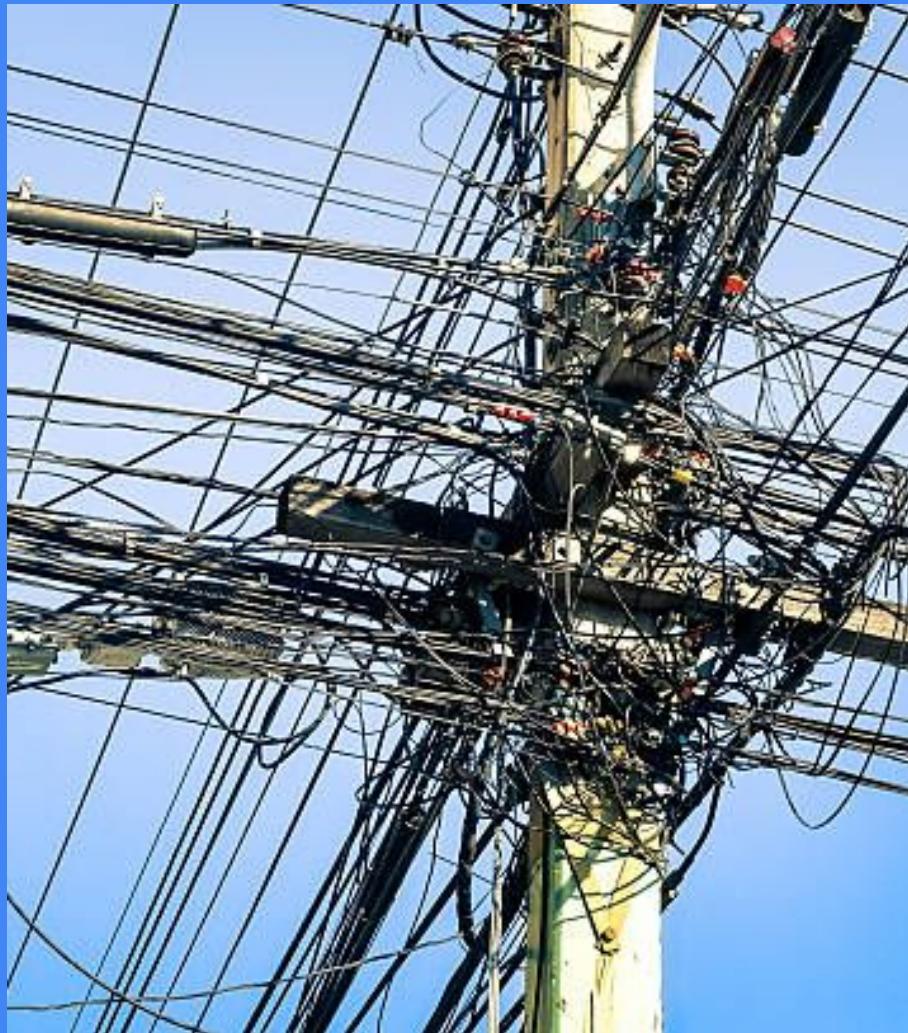
1. Introduce Container with Docker
2. Docker System (a.k.a Engine)
3. Docker Container
4. Docker Image
5. Docker Volume
6. Docker Network
7. Docker Compose
8. Principle for Design Application

A large cargo ship, the MSC ZEPHYRUS, is shown from a low angle, sailing on the ocean under a clear blue sky. The ship's hull is dark blue with white horizontal stripes and features the MSC logo. It is heavily loaded with thousands of shipping containers stacked high along its sides. A small white pilot boat is visible in the water near the ship.

What problem is Docker trying to Solve

#1

Inconsistency Dependency And Complexity

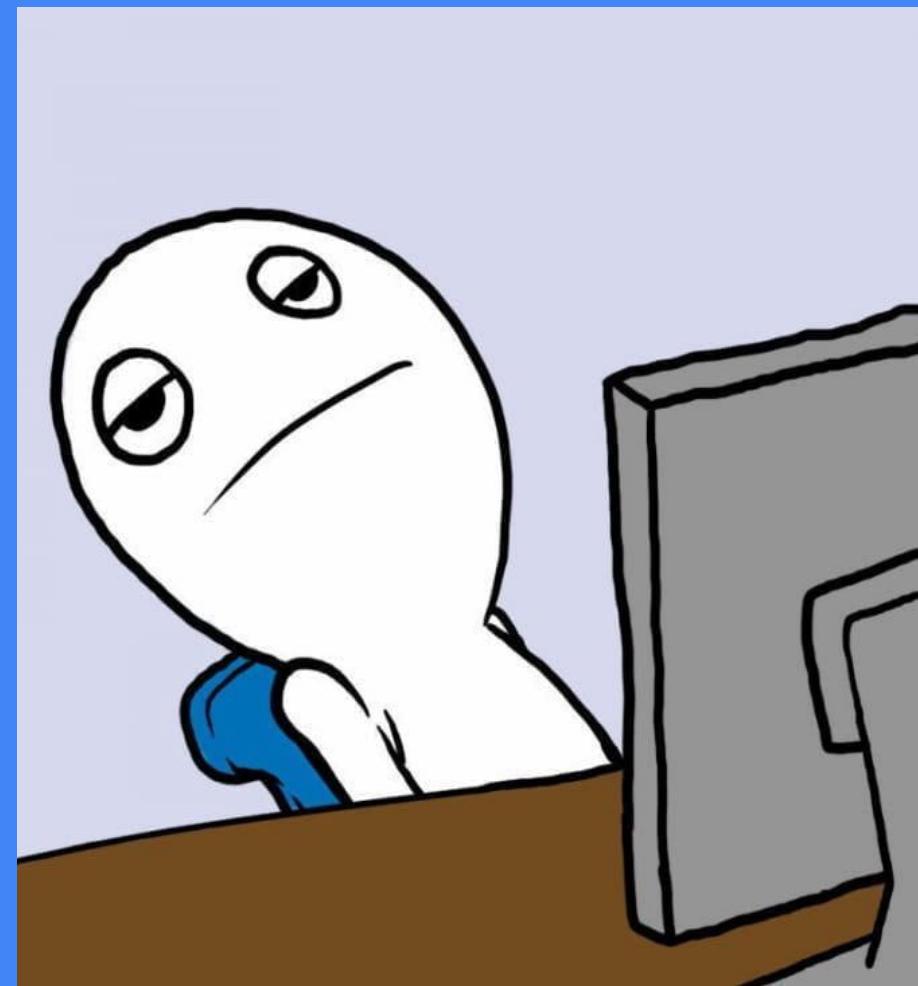


#2

It works on my machine



#3
Install,
Upgrade
Remove
and repeat.



#4 Security?

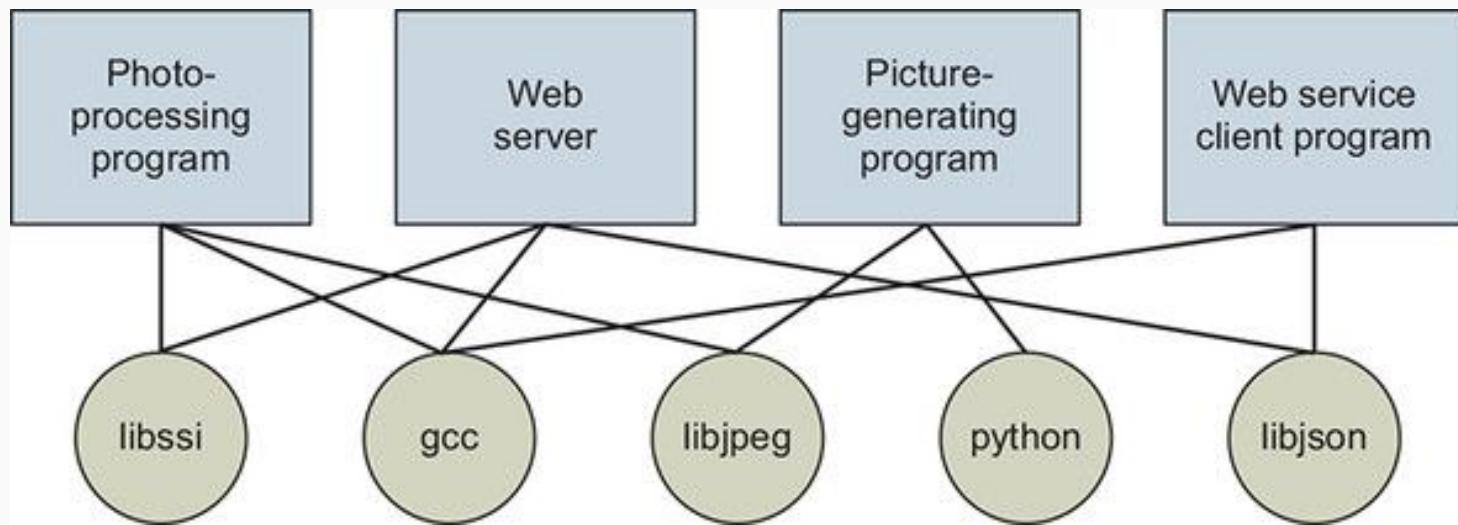


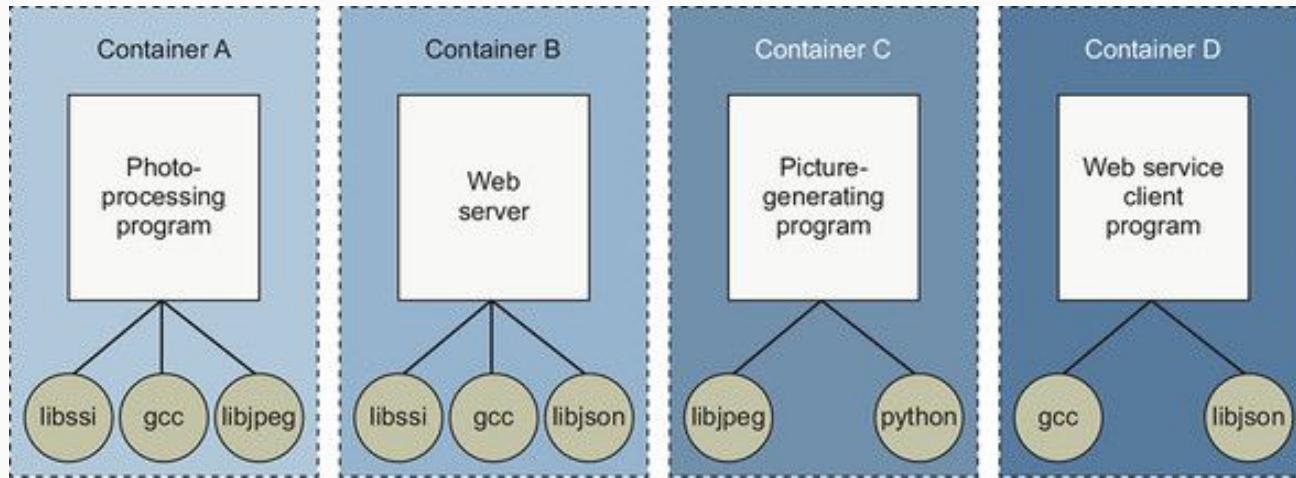
#5 Testing?

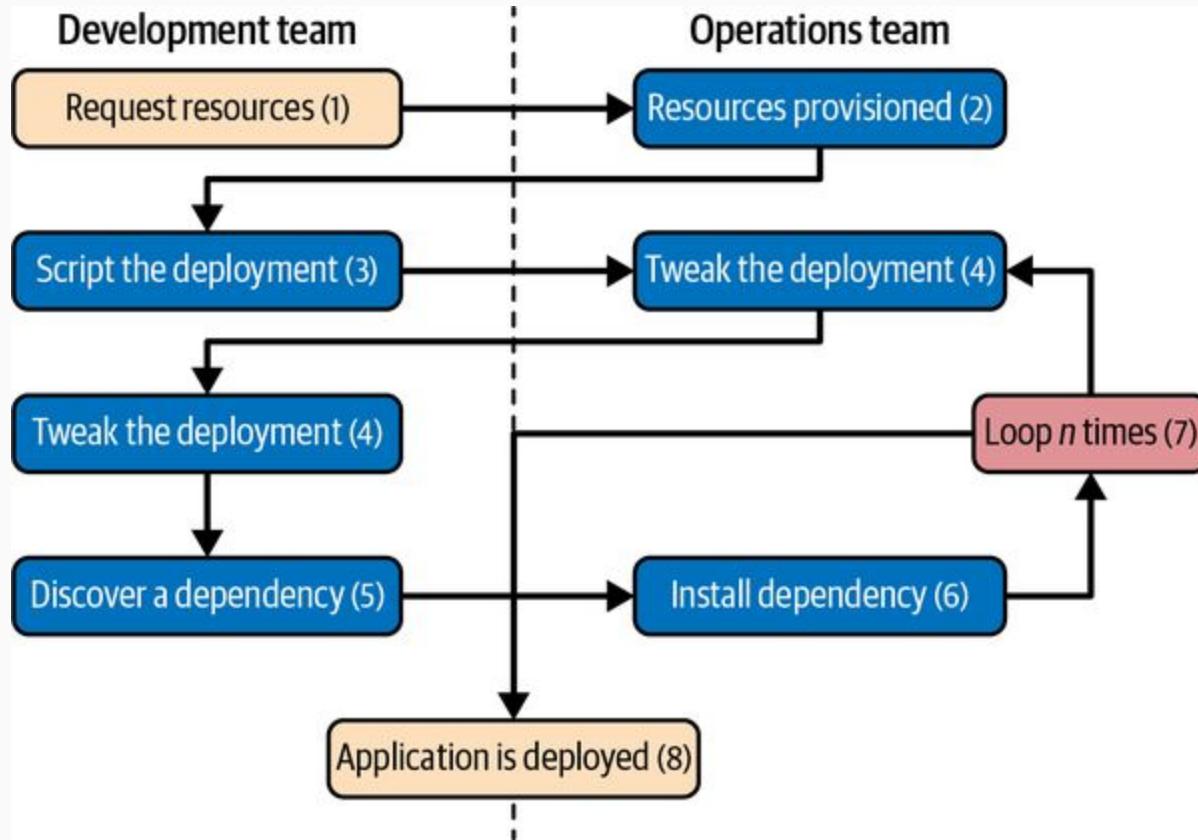


**TEST
ON YOUR
MACHINE**

**TEST IN
PRODUCTION**







Development team

Build image

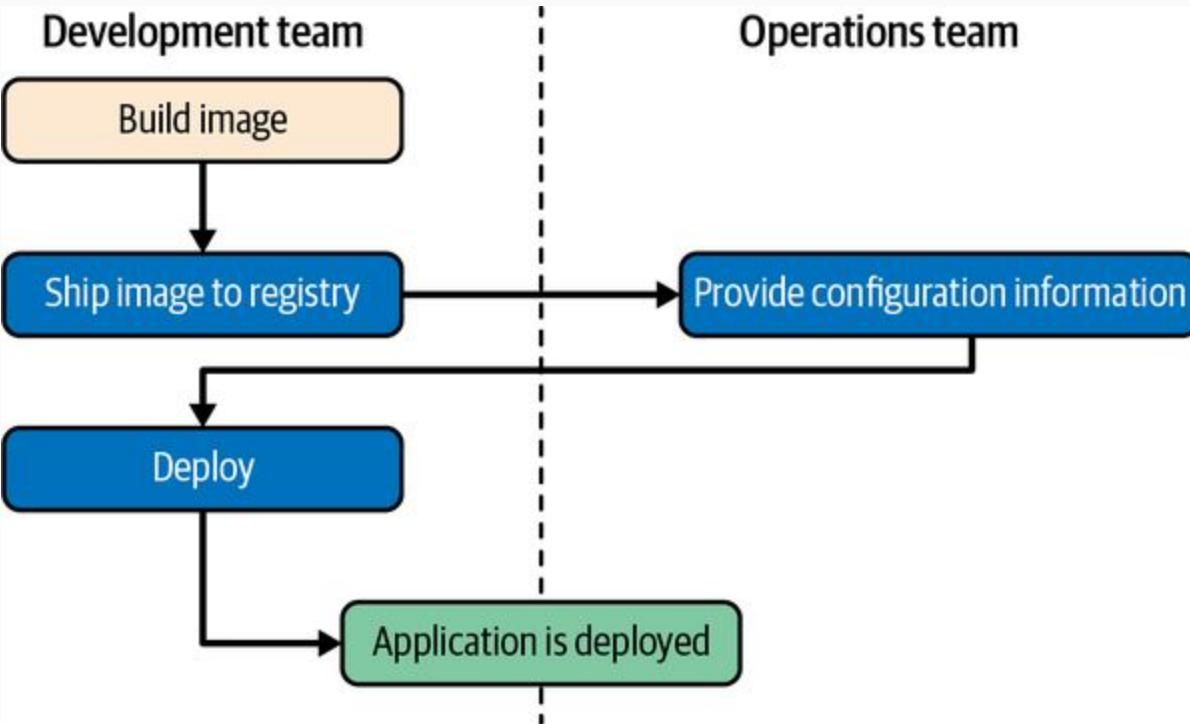
Ship image to registry

Deploy

Operations team

Provide configuration information

Application is deployed



Docker Container



Workshop

Hello Docker



Exploring Docker Machine

\$ docker system info

\$ docker context ls

\$ docker context desktop-linux

```
# Run a simple docker container
```

```
$ docker image pull hello-world
```

```
$ docker image ls
```

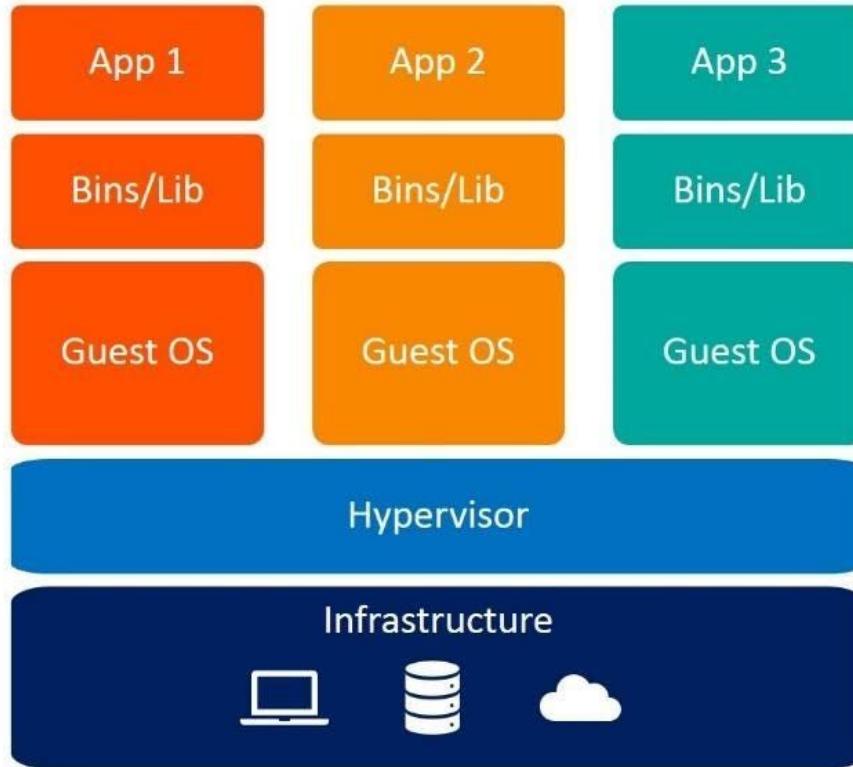
```
$ docker images
```

```
$ docker run - -name hello-world hello-world
```

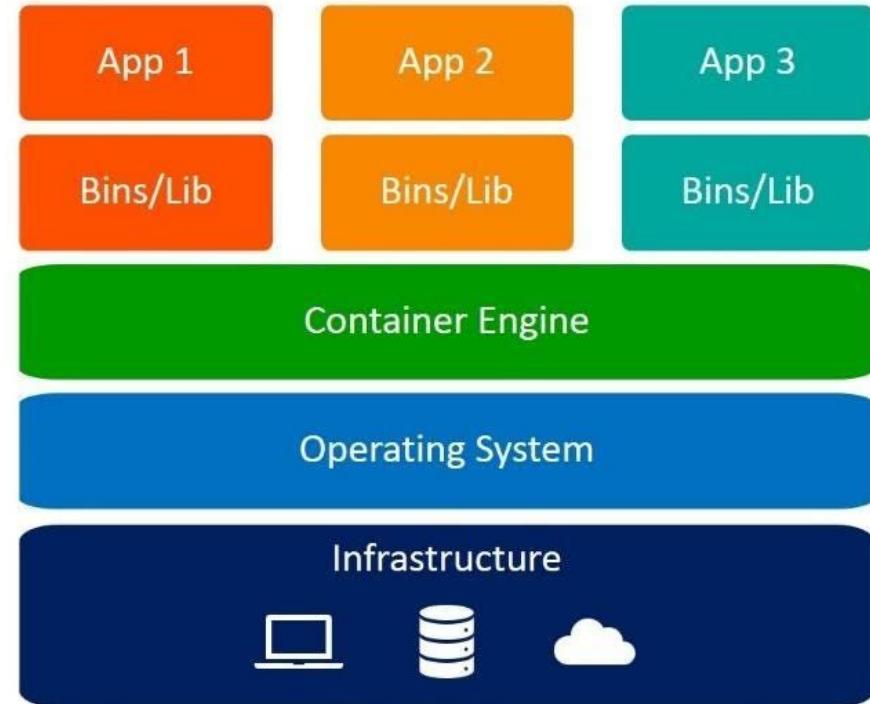
```
$ docker container ps
```

Docker Architecture Under the Hood



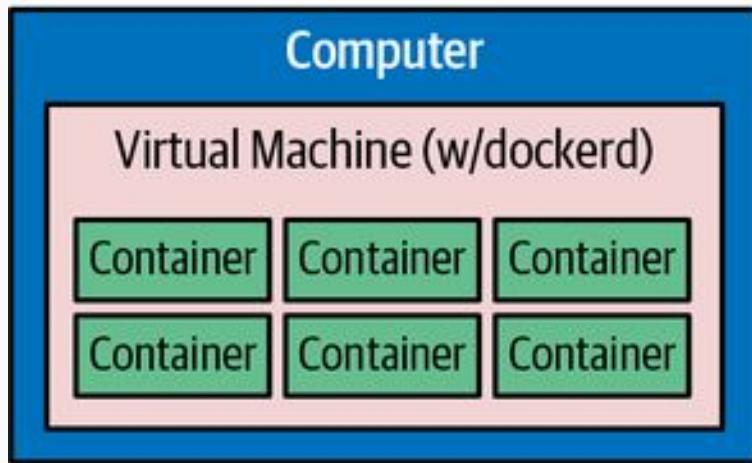


Virtual Machines

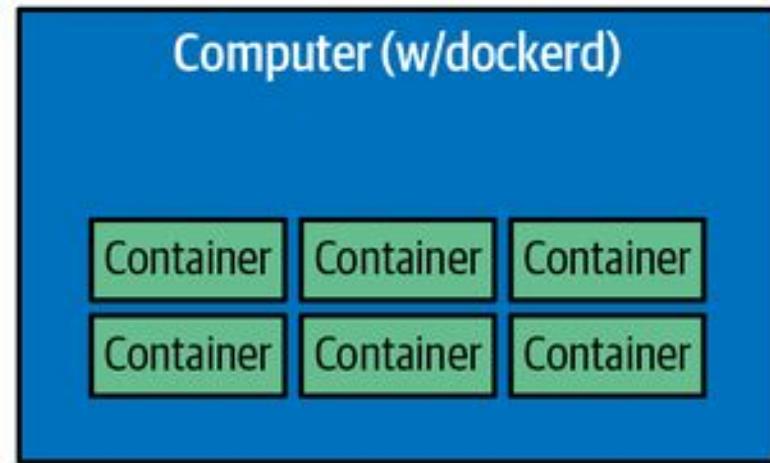


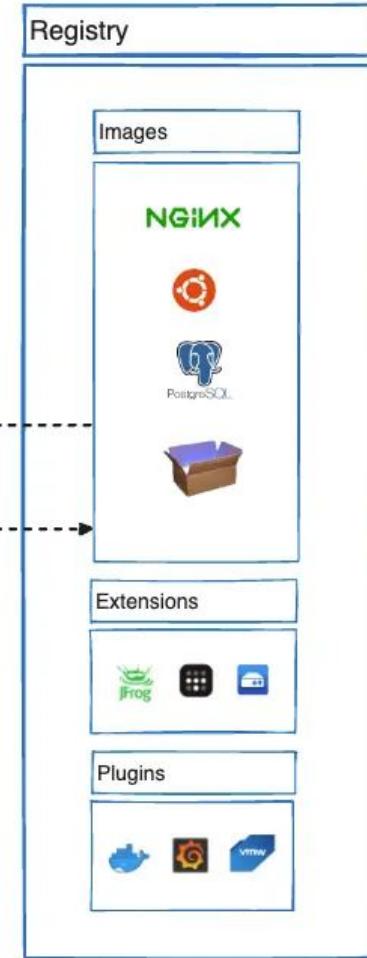
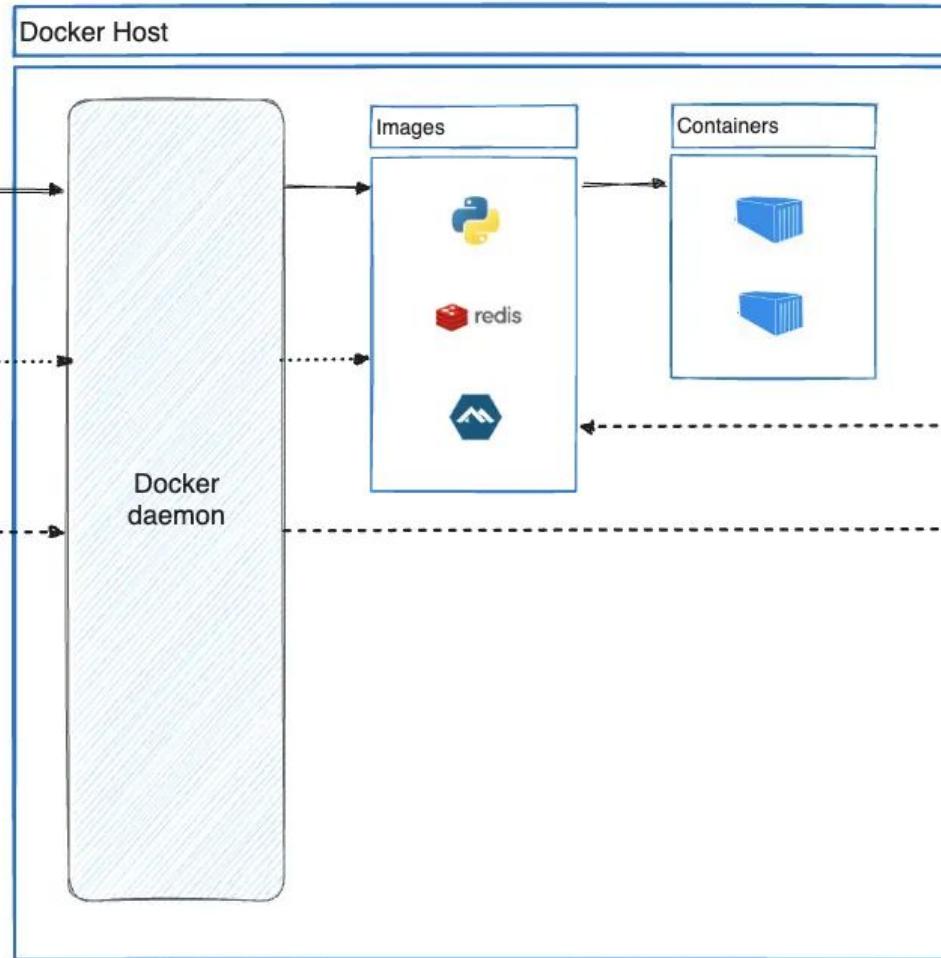
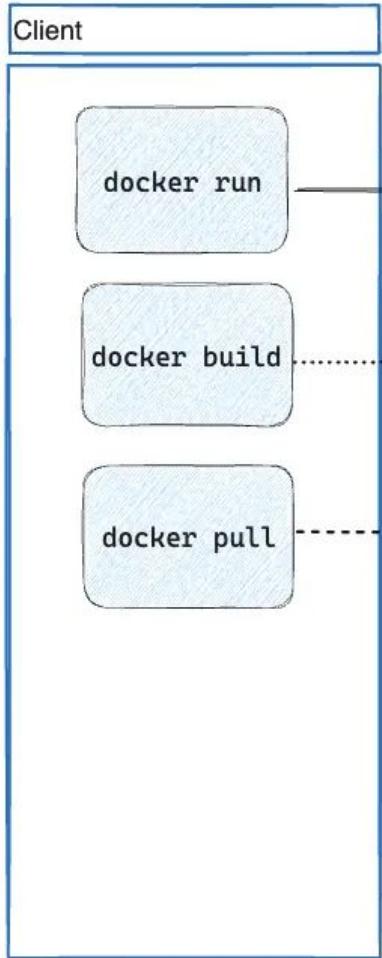
Containers

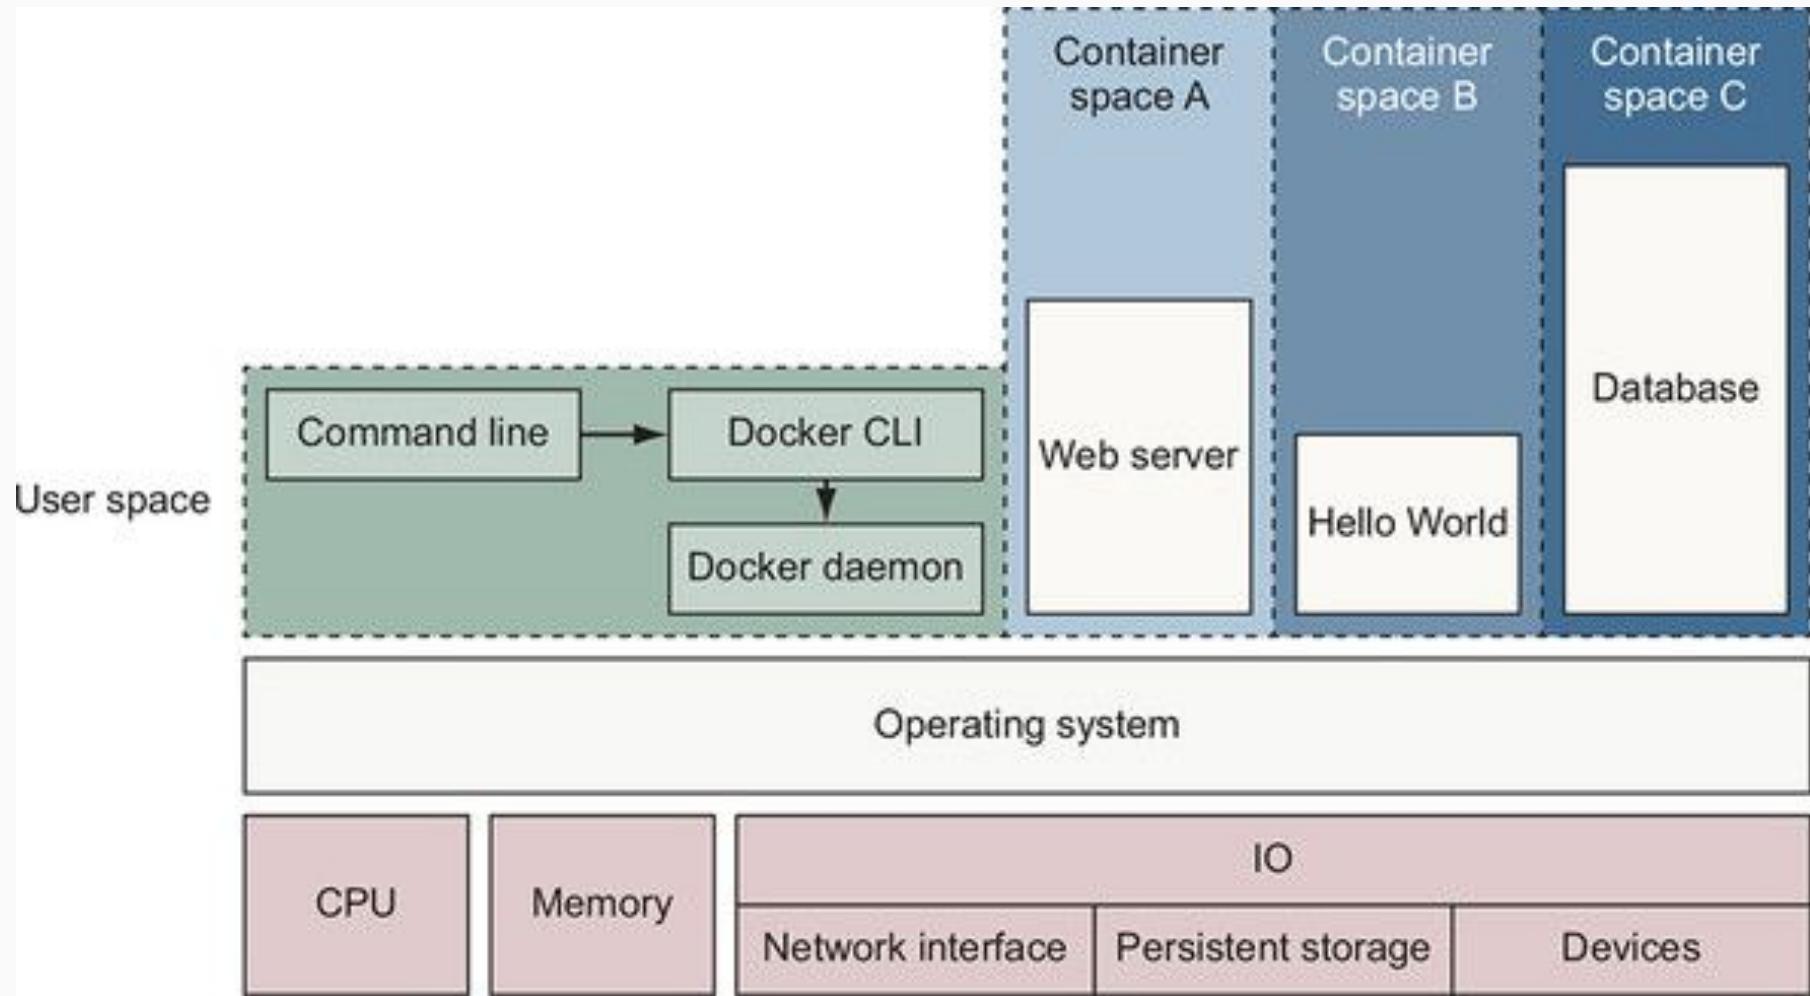
Docker with VM
(Windows, macOS, etc.)

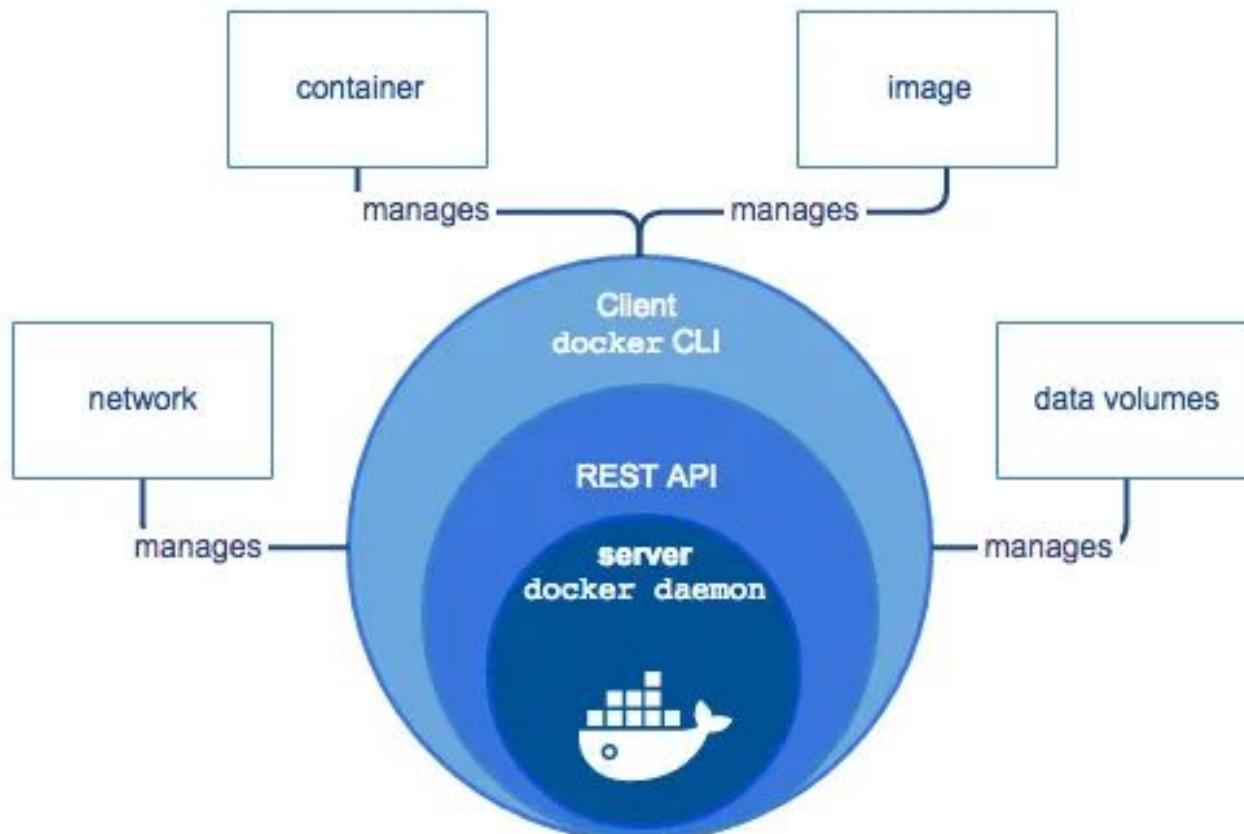


Native Docker
(Linux)









Docker Image



Workshop

Create Your Own Docker Image



```
# Create an image from cli
```

```
$ docker image pull nginx
```

```
$ docker images
```

```
$ docker run --detach --name my-nginx-container nginx
```

```
$ docker container ps
```

```
$ docker container exec -it my-nginx-container bash
```

```
# Update index.html in Container  
  
$ docker container exec -it my-nginx-container bash  
  
#root@d87900c0fa56: cd /usr/share/nginx/html  
  
# ls -l  
  
# vim index.html  
  
# apt-get update && apt-get install vim  
  
# exit
```

```
# Create an image from cli  
  
$ docker container commit \  
  -a jostars -m "make nginx to my welcome me" \  
  my-nginx-container \  
  my-nginx-image:v1  
  
$ docker image ls  
  
$ docker run -it --name my-nginx-container-v1 my-nginx-image:v1 bash  
  
$ cat /usr/share/nginx/html/index.html
```



Dockerfile

build



Docker Image

run



Docker Container

```
# Create an image from Dockerfile
```

```
$ cat <<EOF >>index.html  
<h1> HELLO, My NGINX V2 </h1>  
EOF
```

```
```Dockerfile  
FROM nginx:latest
COPY index.html /usr/share/nginx/html/index.html
```
```

```
$ docker build -t my-nginx-image:v2 .
```

```
$ docker run -d --rm -p 8080:80 my-nginx-image:v2
```

```
$ curl http://localhost:8080
```



Layers

Cache?

```
FROM golang:1.20-alpine
```



```
WORKDIR /src
```



```
COPY go.mod go.sum .
```



```
RUN go mod download
```



```
COPY . .
```



```
RUN go build -o /bin/client ./cmd/client
```



```
RUN go build -o /bin/server ./cmd/server
```



```
ENTRYPOINT [ "/bin/server" ]
```

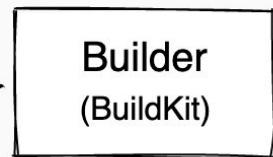


Docker Build high-level architecture

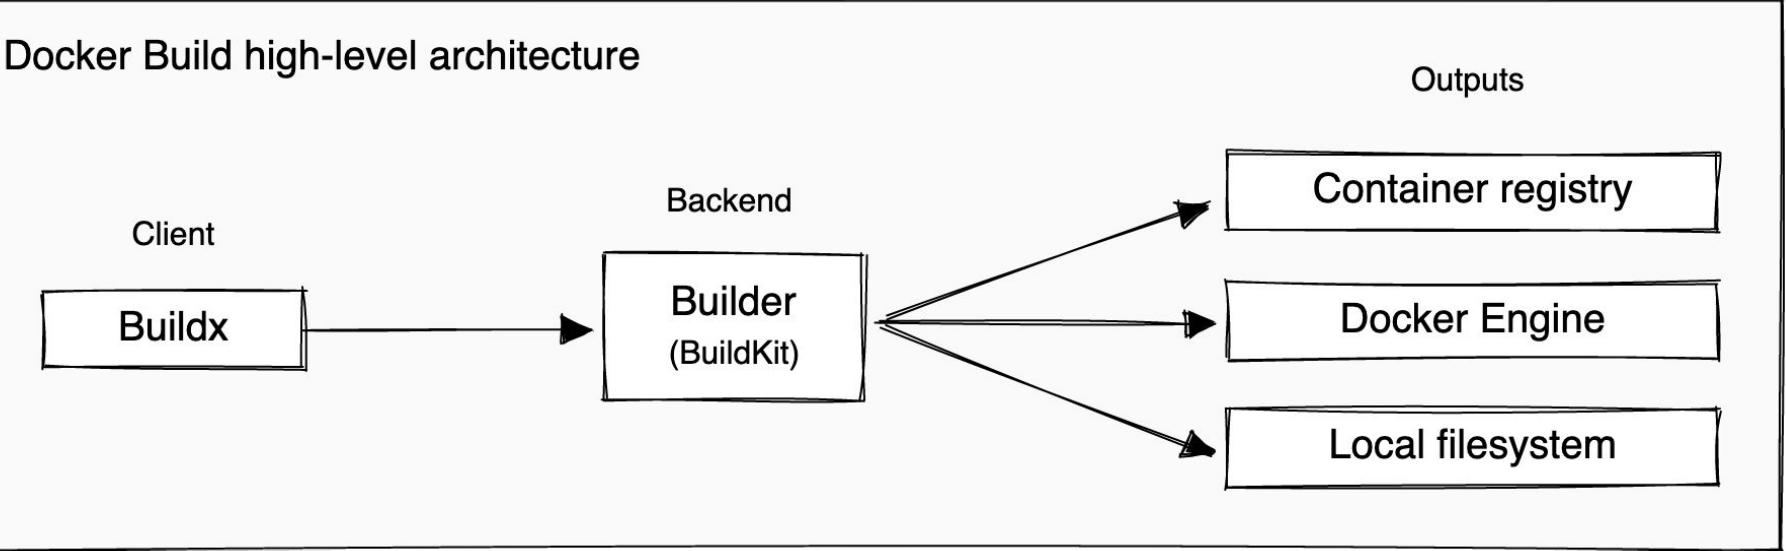
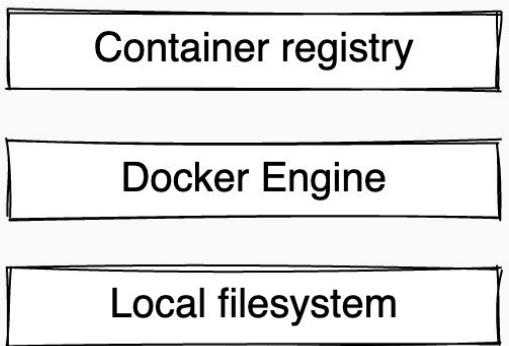
Client



Backend



Outputs



Structure of Docker Image Name

{Organization}/{Publisher}/{ImageName}:{Tag}

Ex. scrum123.com/kkosittaruk/nginx:1.0.0

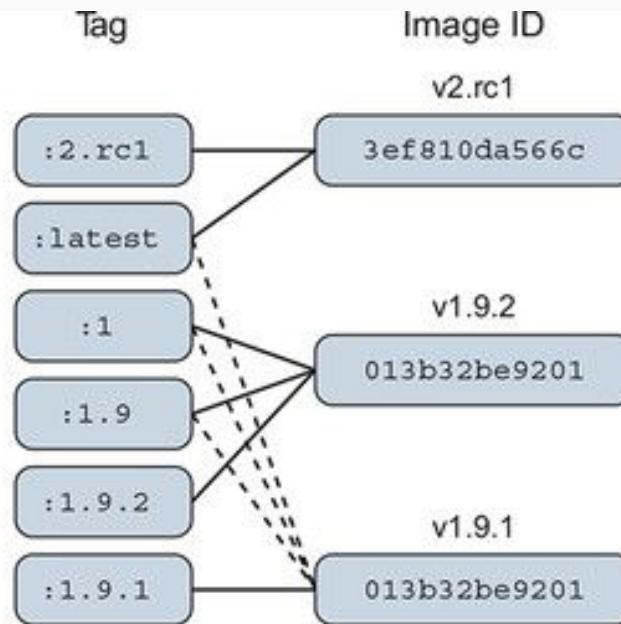
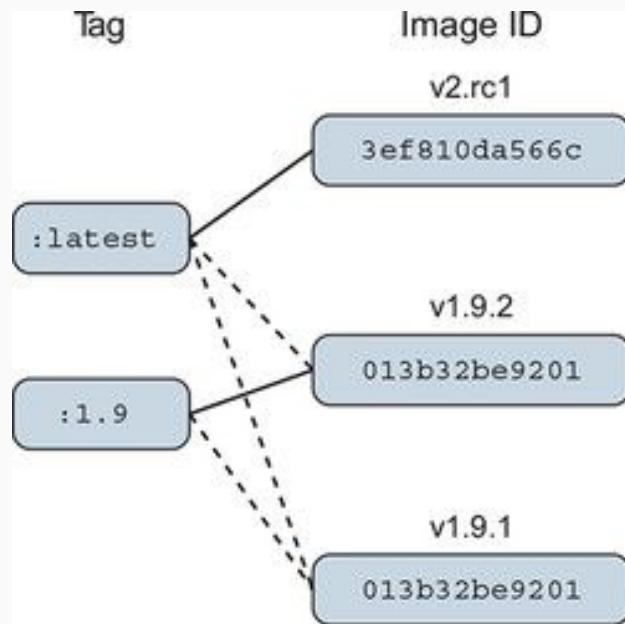
Docker Image Tag Pattern

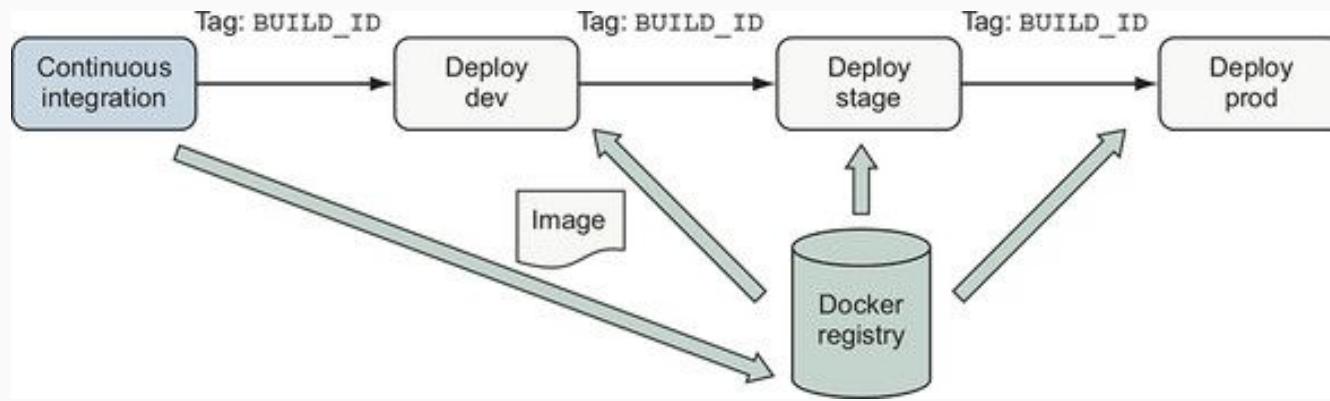
Patterns

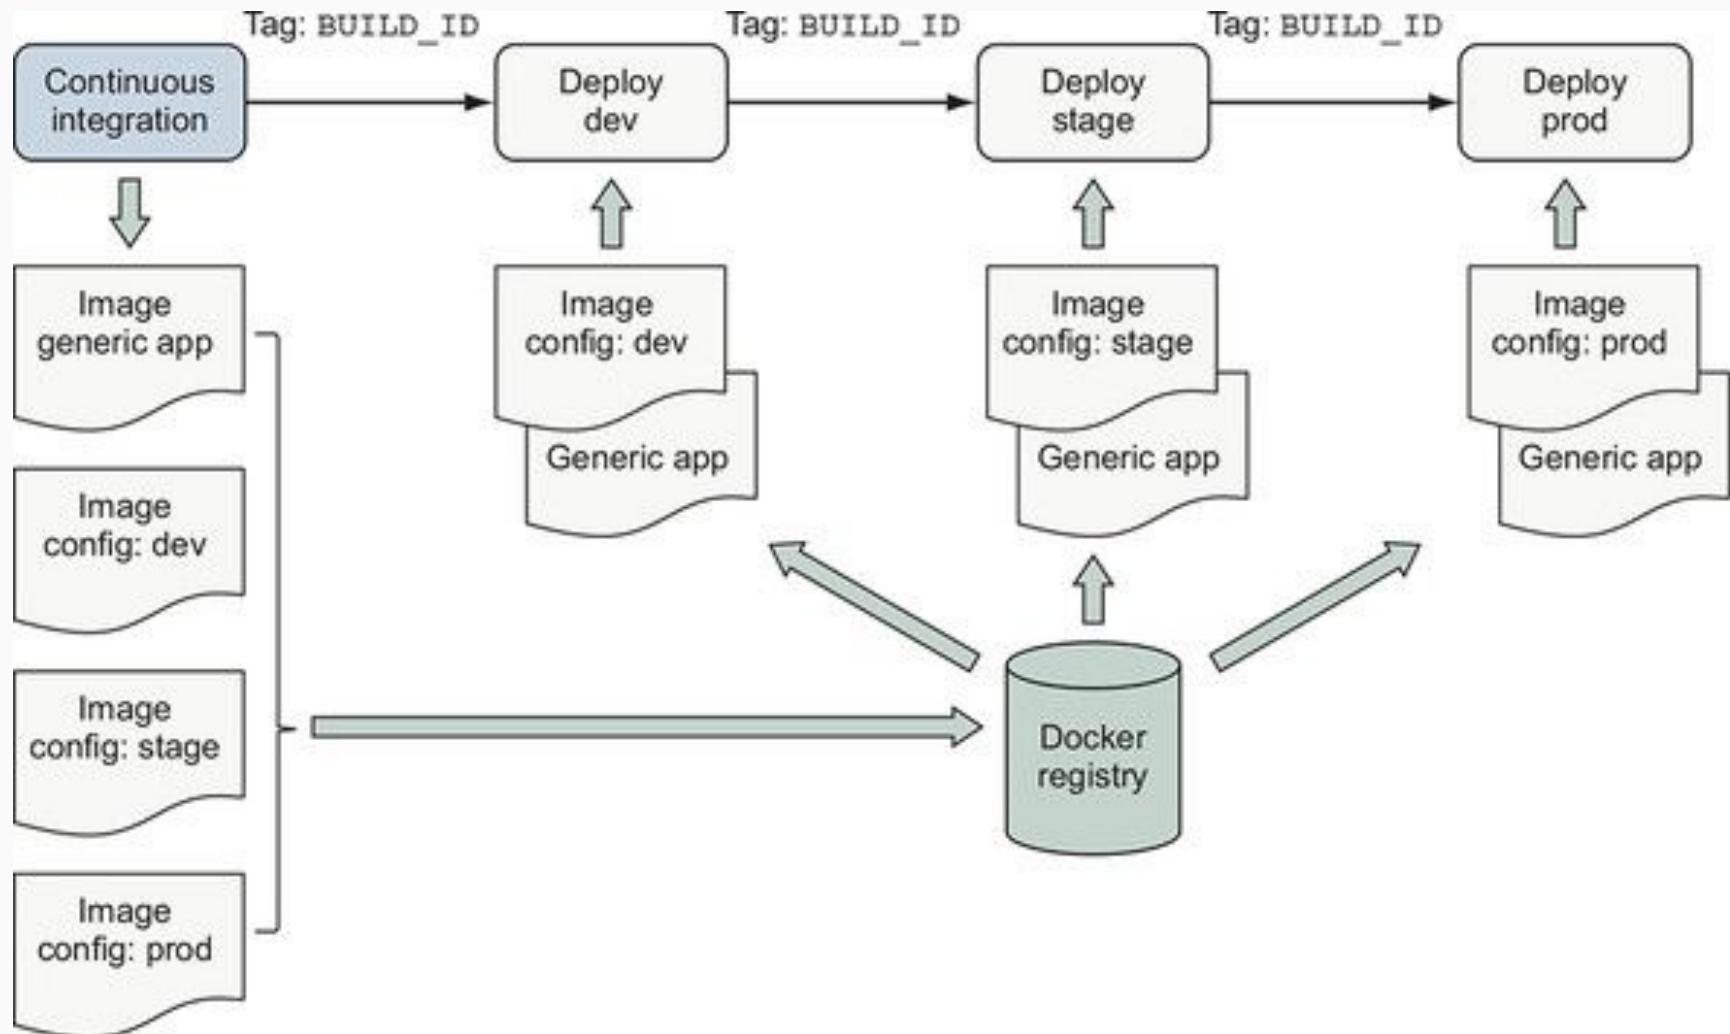
1. Semantic versioning (ex. 1.2.1)
2. Date versioning (ex. 2024.03.25)
3. Environment
(ex. nginx:dev, nginx:uat, nginx:prod)
4. Hybrid versioning
(ex. 1.2.1-dev, 1.2.0-uat, 1.2.0-prod)
5. Hash versioning (ex. nginx:cfe4ff9)

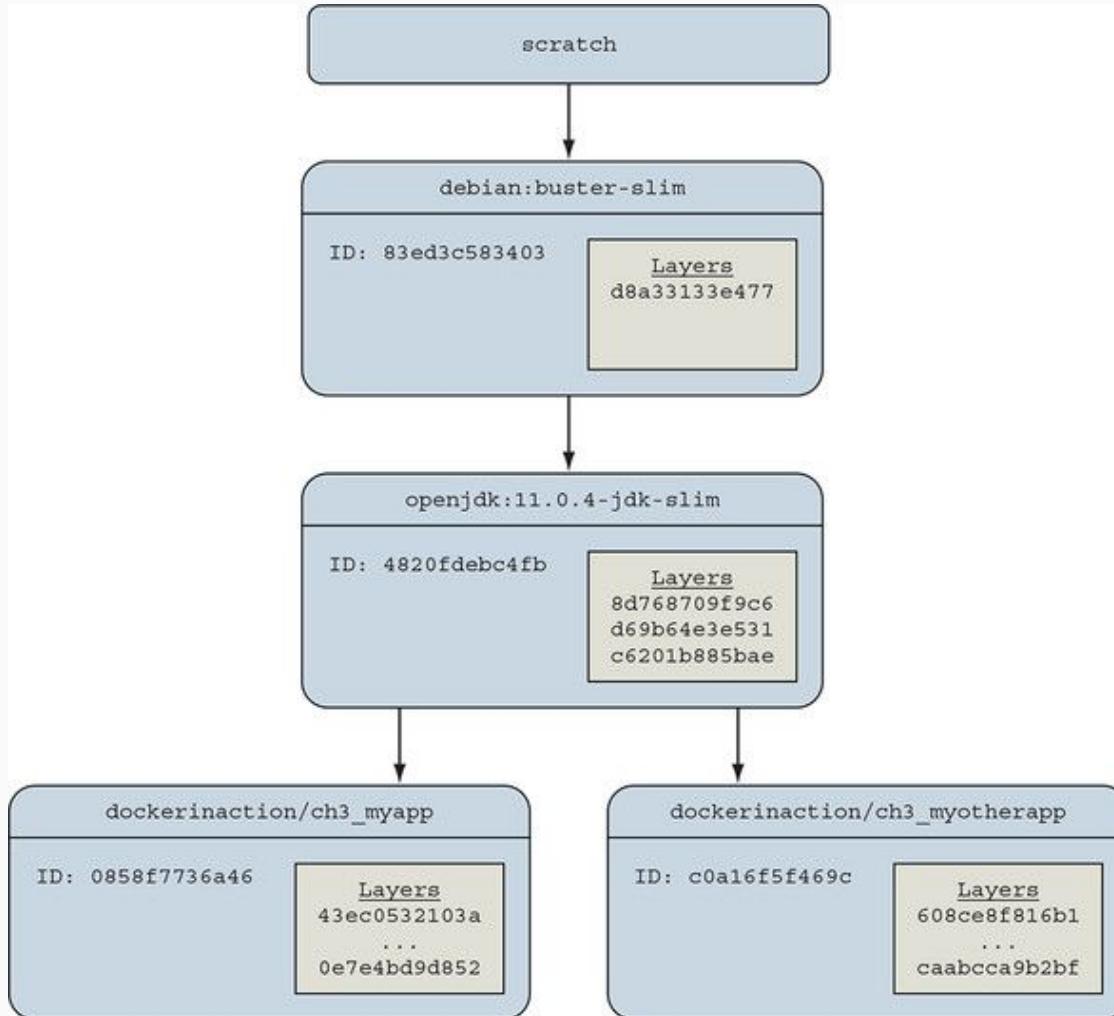
Need to know about tag

- Tags are **mutable**
- Single image can multiple tag
- **Avoid** to use tag **latest** in production and test environment









Docker Container



Workshop

Deep Dive Docker Container



| | |
|-----------------------------|--|
| # Exploring in container | # Remove Container |
| \$ docker container inspect | \$ docker container rm |
| \$ docker container ps | # Container Logs |
| # List Container | \$ docker container logs |
| \$ docker container ls | # Exec Command to Container |
| # Start Stop Container | \$ docker container exec |
| \$ docker container start | # Remove all container |
| \$ docker container stop | \$ docker container rm \$(docker container ls -aq) |

Workshop

Configure Container with Environment



```
# Set Environment via Command line
```

```
$ docker image pull alpine
```

```
$ docker image inspect alpine
```

```
$ docker run --rm alpine /bin/sh
```

```
#root: printenv
```

```
$ docker run -it --rm -e HELLO=world alpine /bin/sh
```

```
# Set Environment via Dockerfile
```

```
$ cat <<EOF >>Dockerfile.env
```

```
FROM alpine
```

```
ENV HELLO=world
```

```
EOF
```

```
$ docker build -t alpine-env:v1 -f Dockerfile.env .
```

```
$ docker run --rm -it alpine-env:v1 /bin/sh
```

```
$ docker image inspect -f "{{json .Config.Env}}" alpine-env:v1 | jq .
```

```
# Set Environment via .env file
```

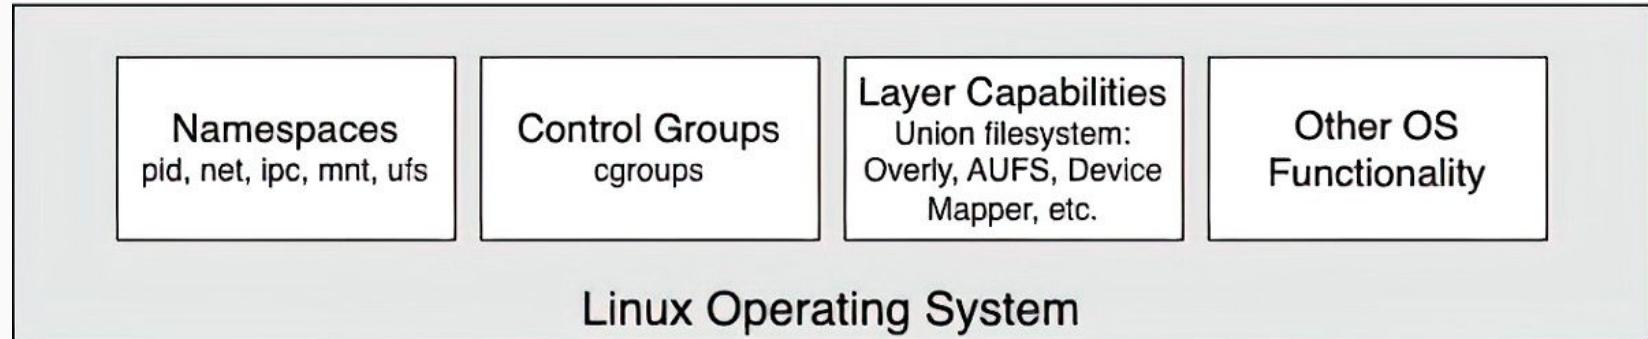
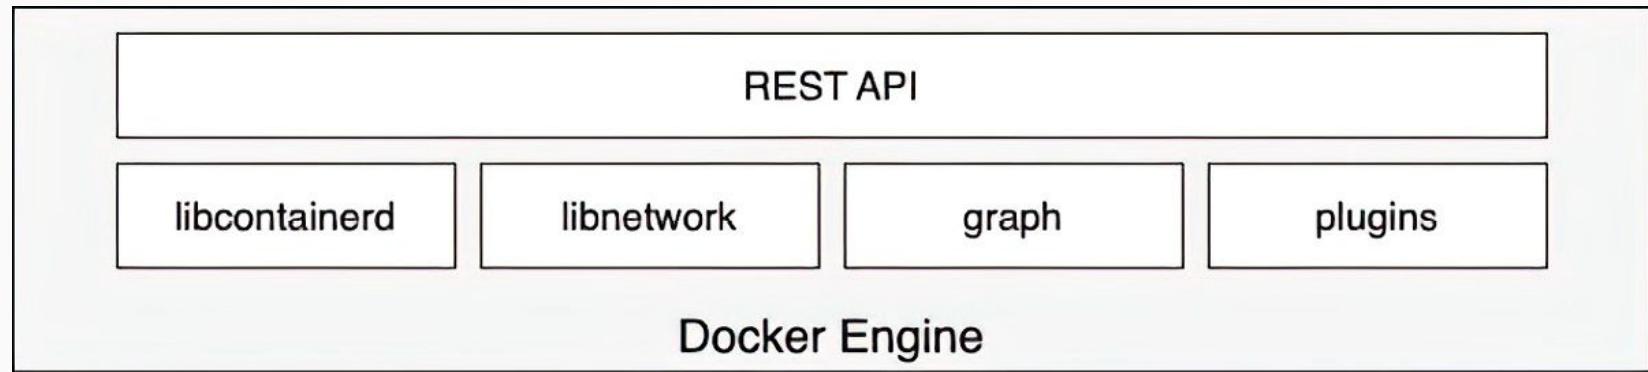
```
$ cat <<EOF >>.env
```

```
HELLO=again?
```

```
SECRET=Don't tell my wife
```

```
EOF
```

```
$ docker run --rm -it -env-file .env alpine-env:v1 printenv
```



runc

*runc is a lightweight, portable **container runtime**. It provides full support for Linux namespaces as well as native support for all security features available on Linux, such as SELinux, AppArmor, seccomp, and cgroups.*

containerd

*containerd builds on top of runc and adds higher-level features, such as image transfer and storage, container execution, and supervision as well as network and storage attachments. With this, **it manages the complete life cycle of containers**. containerd is the reference implementation of the OCI specifications and is by far the most popular and widely used container runtime.*

Check your container runtime

```
$ docker system info
```

Alternative Container Runtime



CNCF Graduated



CNCF Incubating



CNCF Incubating



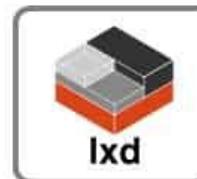
Firecracker



gVisor



kata



lxd



Nabla Containers



Pouch



OPEN
CONTAINER
INITIATIVE
runc



Singularity



SmartOS



unik

Docker Volume



Workshop

with Volume



```
# Data in container
```

```
$ docker container run --name demo \  
alpine /bin/sh -c 'echo "This is a test" > sample.txt'
```

```
# Check container change with `diff`
```

```
$ docker container diff demo
```

```
# Volume for persistent your data
```

```
$ docker volume create sample
```

```
$ docker volume inspect sample
```

```
# Mount volume to container
```

```
$ docker container run --name test -it \  
-v sample:/data \  
alpine /bin/sh
```

```
# cd data && touch data1.txt && touch data2.txt  
# exit
```

```
$ docker container rm test
```

```
# Create new container with different OS
```

```
$ docker container run --name test2 -it --rm \  
-v sample:/app/data \  
centos:7 /bin/bash
```

```
# ls -l /app/data/
```

Where is the data stored? Let's Investigate

\$ docker volume ls

\$ docker volume inspect sample

\$ docker volume inspect -f "{{json .Mountpoint}}" sample

\$ cd /var/lib/docker/volumes/sample/_data

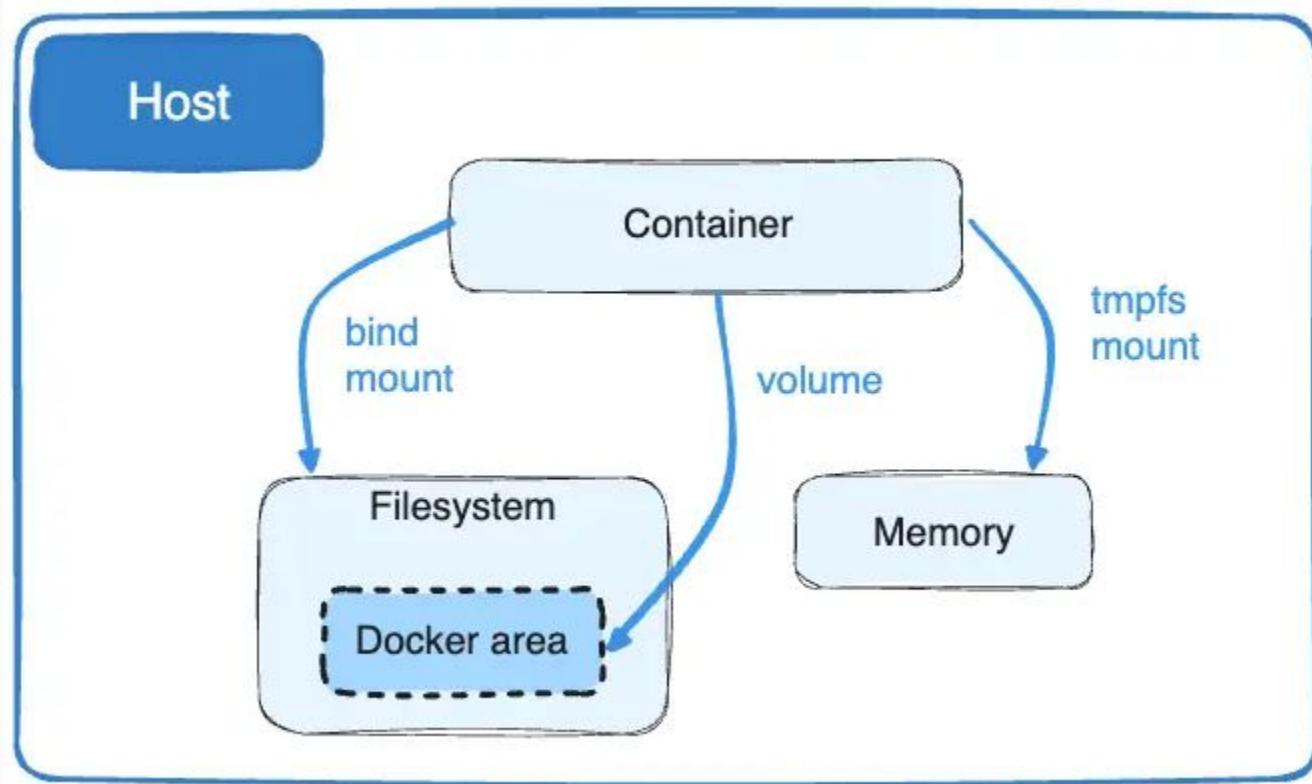
```
# no such file or directory: /var/lib/docker/volumes/sample/_data
```

WHY!!!

Answer is :

Docker does not run natively on Mac directly; instead, it runs inside a slim VM. The data for the VM that Docker creates can be found in this directory.

~/Library/Containers/com.docker.docker/Data/vms/0 folder.



Fantastic volume and where to find them"

```
$ docker container run -it --privileged --pid=host \
debian nsenter -t 1 -m -u -n -i sh
```

```
# ls -l /var/lib/docker/volumes
```

```
# ls -l /var/lib/docker/volumes/sample/_data
```

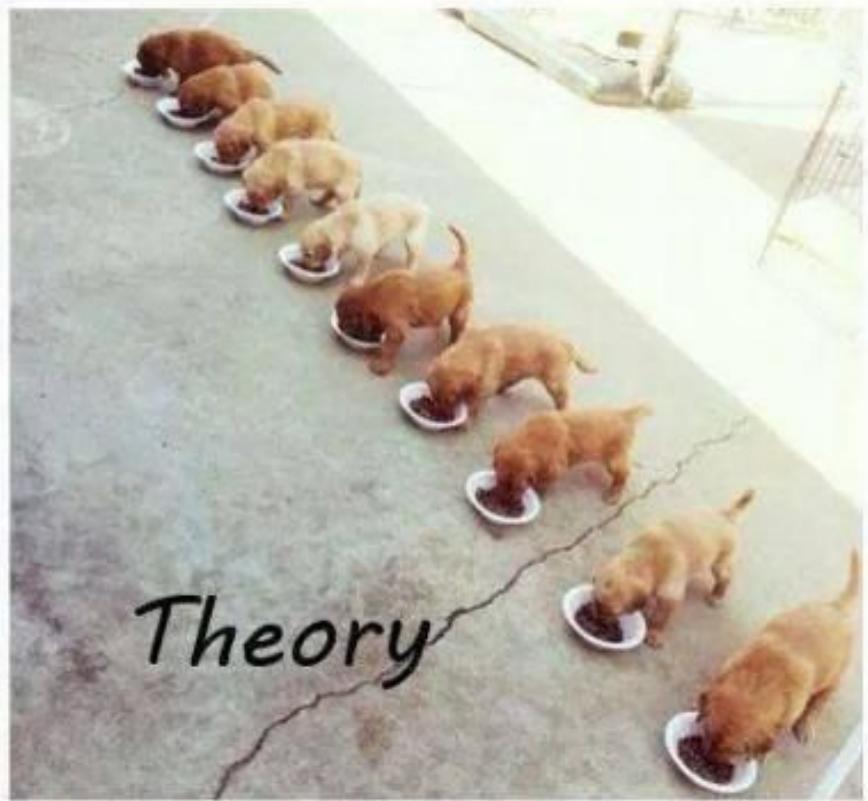
Let's add more data from another container

```
$ docker container run --rm -it -v sample:/data alpine /bin/sh
```

```
# echo "Hello world" > /data/sample.txt
```

```
# echo "Hello Other" > /data/other.txt
```

!!! Warning !!! Share Data Volume in Multiple Container



Workshop

Using Host Volume



```
# I want to host dynamic content
```

```
$ echo "<h1>HELLO, Docker</h1>" > index.html
```

```
$ cat <<EOF >>Dockerfile-volume
FROM nginx:latest
COPY index.html /usr/share/nginx/html/index.html
EOF
```

```
$ docker build -f Dockerfile-volume -t nginx-volume:v1 .
```

```
$ docker run -d --rm -p 8080:80 nginx-volume:v1
```

```
$ curl http://localhost:8080
```

```
$ echo -e "<h2> new awesome title </h2>" >> index.html
```

```
# Add new content to website
```

```
$ echo -e "<h2> new awesome title </h2>" >> index.html
```

```
# Let's get our content
```

```
$ curl http://localhost:8080
```

```
# Mount Host volume mount to container
```

```
$ docker run -d --rm -p 8080:80 \
-v $(pwd)/index.html:/usr/share/nginx/html/index.html \
nginx-volume:v1
```

```
$ curl http://localhost:8080
```

```
$ echo -e "<h3>I was made for lovin' you baby</h3>" >> index.html
```

```
$ curl http://localhost:8080
```

**DOES THIS IMAGE
REQUIRE A VOLUME?**



```
# Let's define a volume for our image.
```

```
$ cat <<EOF >>Dockerfile-define-volume
FROM nginx
VOLUME /usr/share/nginx/html/index.html
EOF
```

```
$ docker build -f Dockerfile-define-volume \
-t nginx-volume:v2 .
```

```
$ docker image inspect \
-f "{{json .Config.Volumes}}" \
nginx-volume:v2 | jq .
```

Real Life Example

```
$ docker image pull mongo
```

```
$ docker image inspect -f "{{json .Config.Volumes}}" mongo | jq .
```

```
$ docker run --rm --name my-mongo -d mongo
```

```
$ docker container inspect -f "{{json .Mounts}}" my-mongo | jq .
```

Docker Network



Workshop

Connect Container



Explore Network in Docker

```
$ docker network ls
```

```
$ docker network inspect {network id}
```

| Network | Company | Scope | Description |
|--------------------------|------------|--------|--|
| Bridge | Docker | Local | Simple network based on Linux bridges to allow networking on a single host |
| Macvlan | Docker | Local | Configures multiple layer-2 (that is, MAC) addresses on a single physical host interface |
| Overlay | Docker | Global | Multi-node capable container network based on Virtual Extensible LAN (VXLAN) |
| Weave Net | Weaveworks | Global | Simple, resilient, multi-host Docker networking |
| Contiv Network
Plugin | Cisco | Global | Open source container networking |

Explore Network in Docker

```
$ docker container run --name c1 -d --rm alpine:latest ping 127.0.0.1
```

```
$ docker container run --name c2 -d --rm alpine:latest ping 127.0.0.1
```

```
$ docker container exec -it c1 /bin/sh
```

```
# ping c2
```

```
# Connect Container with Bride Network
```

```
$ docker network create --driver bridge sample-net
```

```
$ docker network inspect sample-net
```

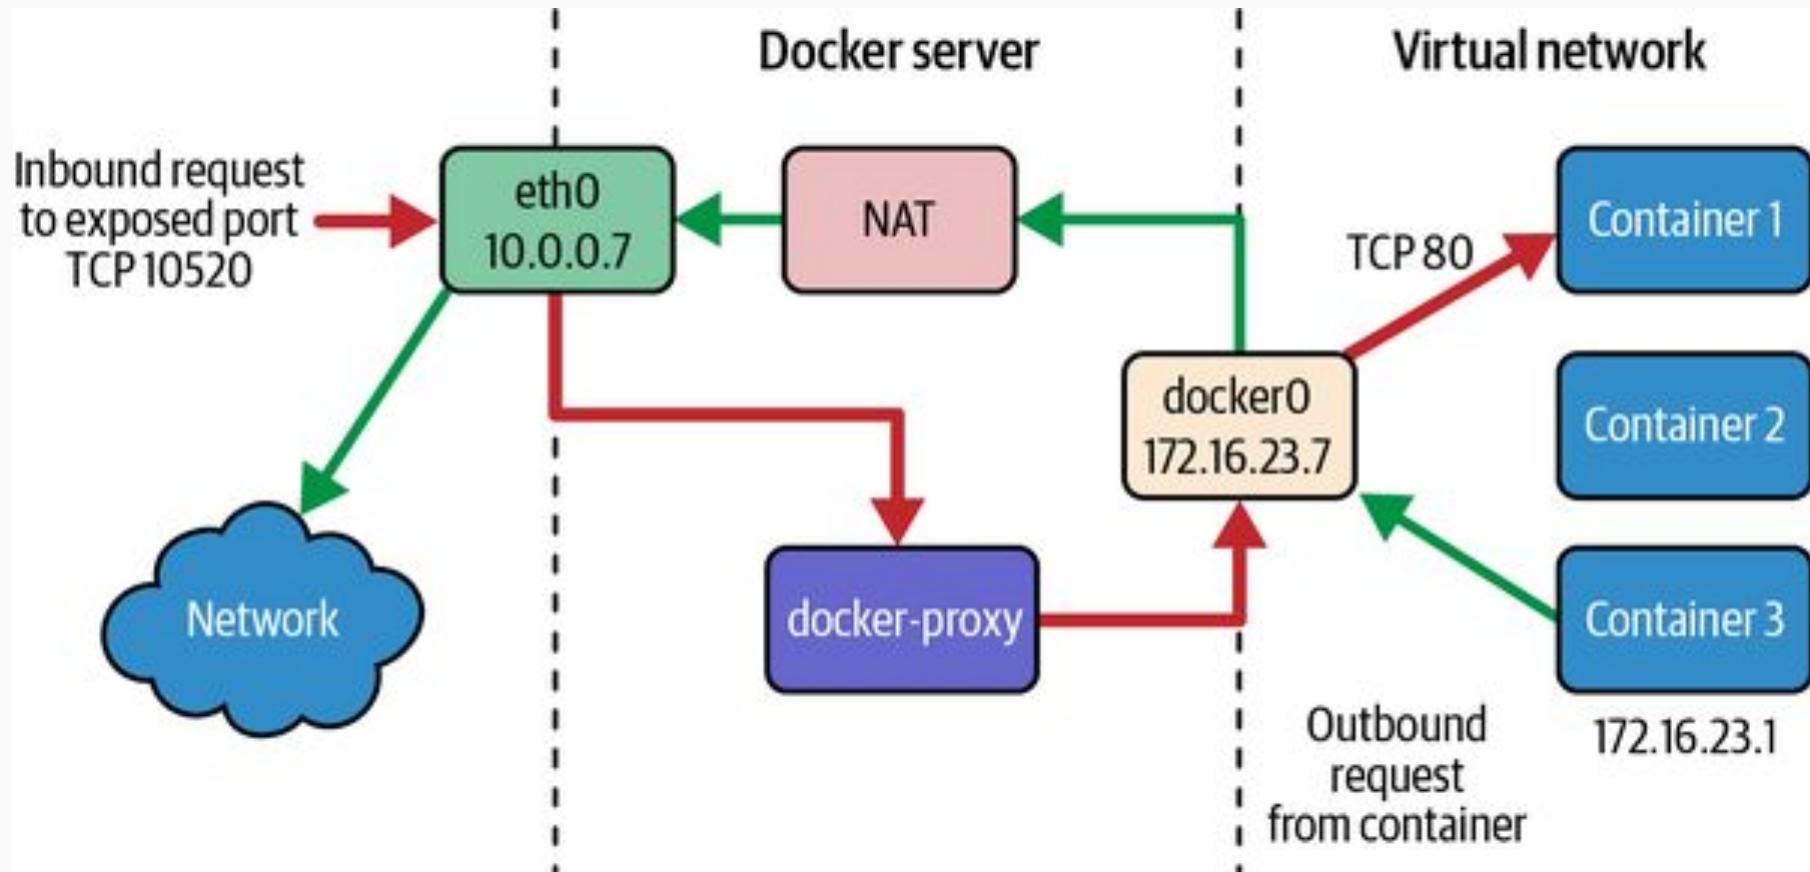
```
$ docker network connect sample-net c1
```

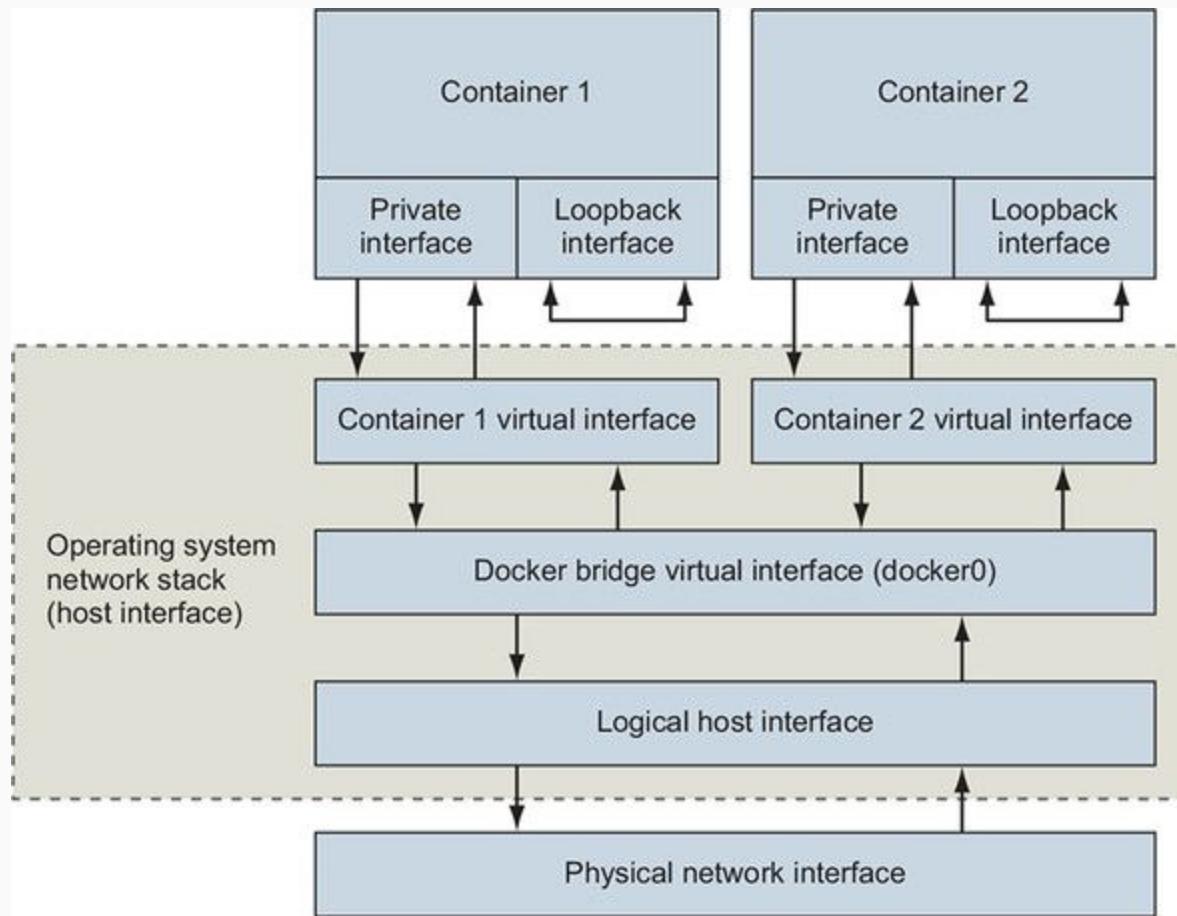
```
$ docker network connect sample-net c2
```

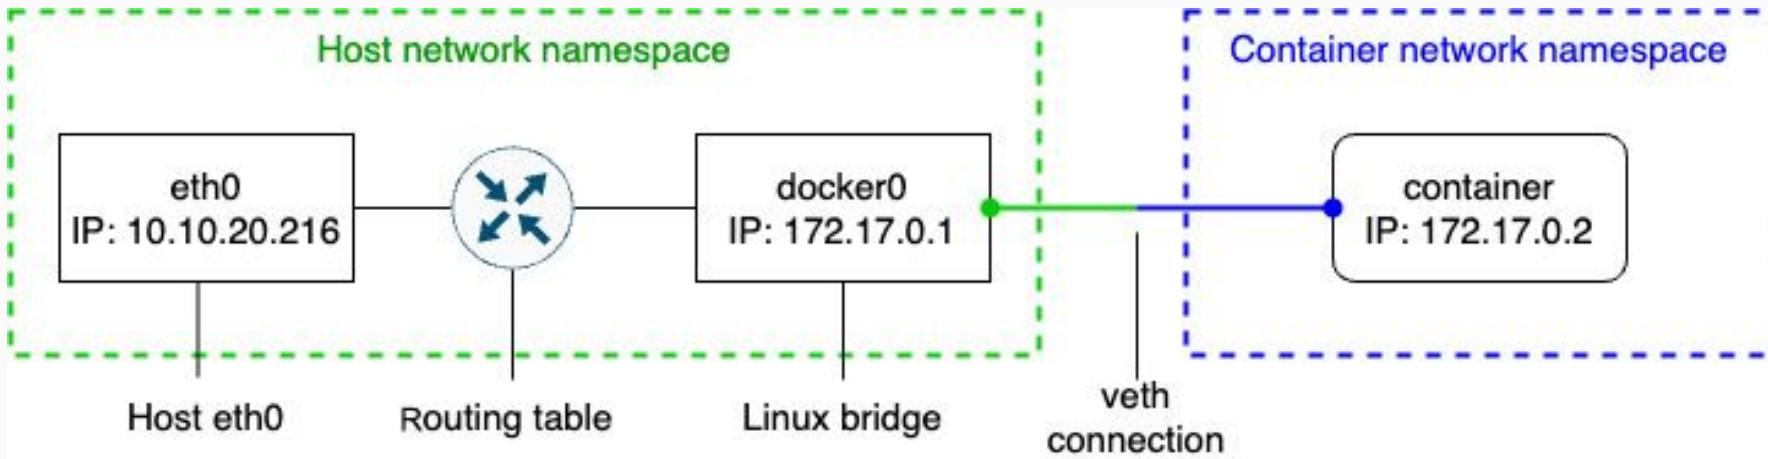
```
$ docker container exec -it c1 sh
```

```
# ping c2
```

```
$ docker container inspect c1
```







Orchestration Your Container



Workshop

Deploy Postgres Cluster



Create PostgreSQL Cluster in Docker

```
$ docker run -d --rm --name postgres \
-e POSTGRES_USER=dockeruser \
-e POSTGRES_PASSWORD=dockerpass \
-e POSTGRES_DB=pets \
-v pg-data:/var/lib/postgresql/data \
-v $(pwd)/db:/docker-entrypoint-initdb.d \
-p 5432:5432 \
postgres:alpine
```

```
# Create PostgreSQL Cluster in Docker
```

```
$ docker volume create admin-data
```

```
$ docker run -d --rm --name pgadmin \
-e PGADMIN_DEFAULT_EMAIL=admin@acme.com \
-e PGADMIN_DEFAULT_PASSWORD=admin \
-v admin-data:/var/lib/pgadmin \
-p 5050:80 \
dpage/pgadmin4
```

```
# Create PostgreSQL Cluster in Docker
```

```
# Trying to Login PGAdmin
```

localhost:5050

```
$ docker network db-net
```

```
$ docker network connect db-net postgres
```

```
$ docker network connect db-net pgadmin
```

```
# Create PostgreSQL Cluster in Docker
```

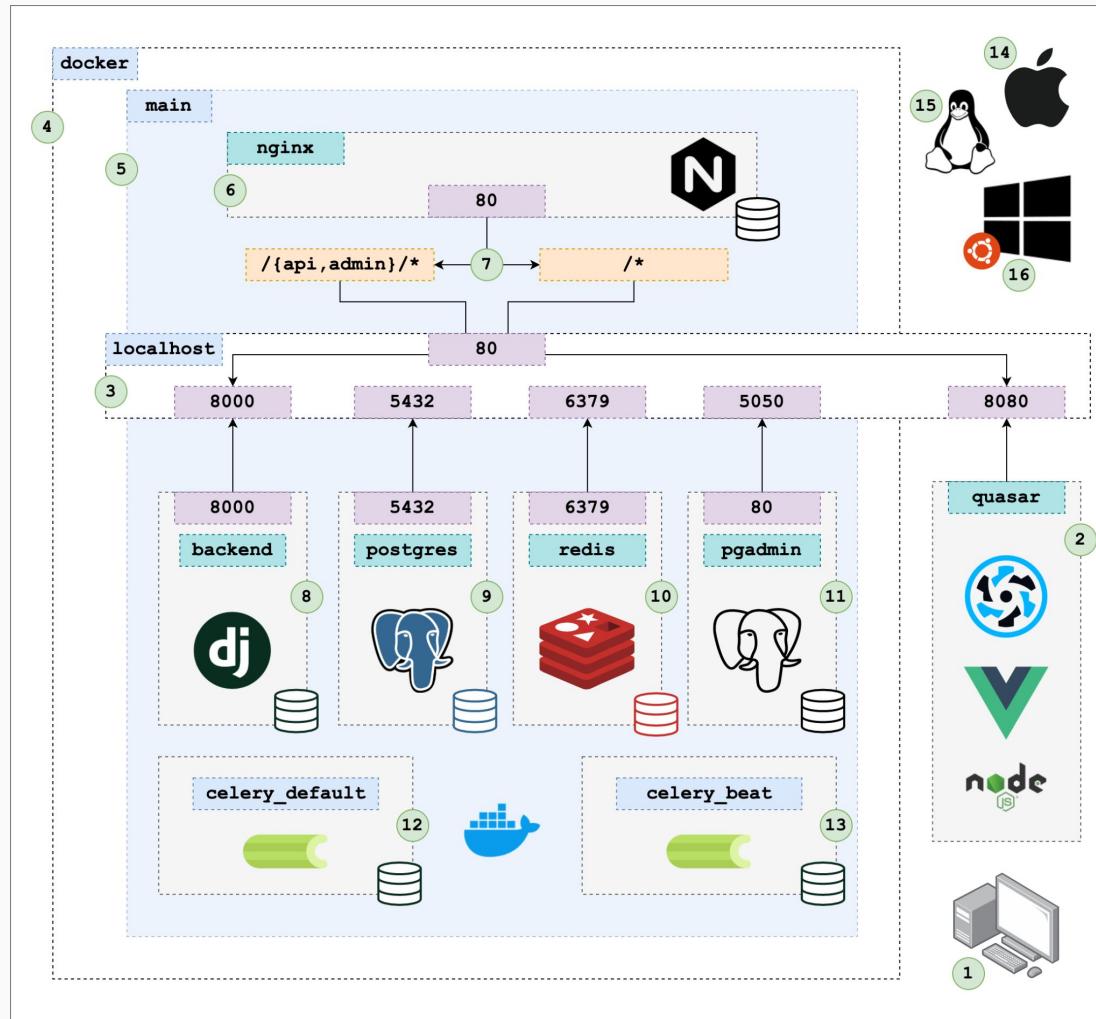
```
$ touch docker-compose.yaml
```

```
$ docker compose up
```

```
$ docker compose up -d
```

```
$ docker compose ls
```

```
$ docker container ps
```



Workshop

Using Docker Compose for Development




```
# Setup Development Cluster
```

```
$ touch docker-compose.yaml
```

```
$ docker compose up
```

```
$ docker compose up -d
```

```
$ docker compose ls
```

```
$ docker container ps
```

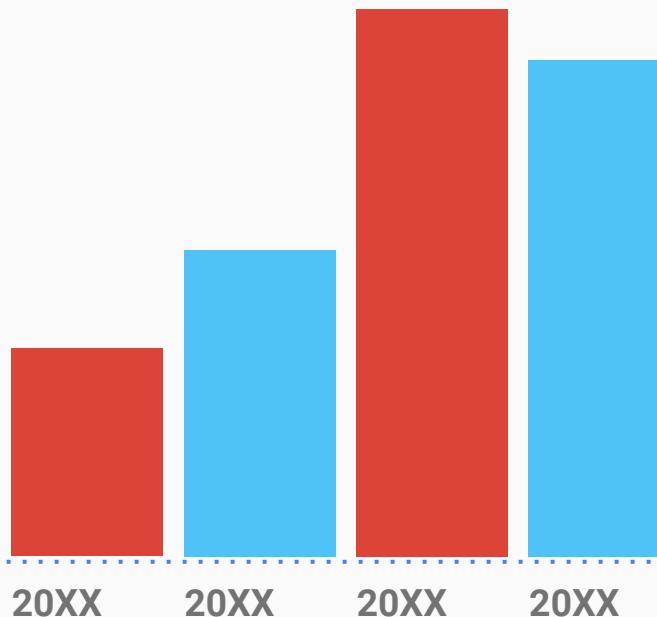
Develop Cloud Native with Twelve Factors



The problem

Frame the problem for the audience.

Quantify the scope of the problem
and connect it to your audience.





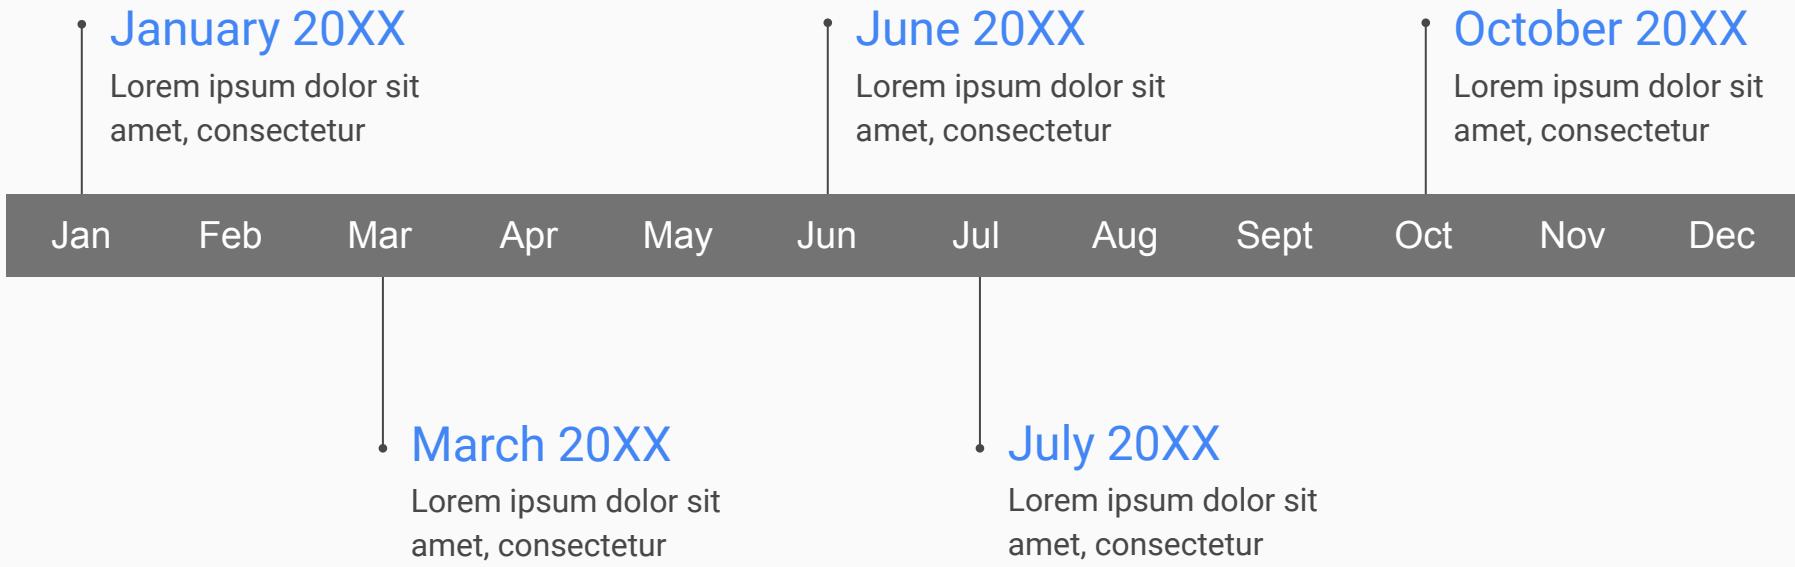
The solution

Show how you solve the problem you identified.

What will be different when the problem is solved (by you)?

Milestones

Show where you are in the process and what's left to tackle



Appendix

Show the audience you anticipated their questions.

Leave room for Q&A, but use the Appendix as a way to show that you both thought about those questions and have solid answers with supporting information. Let the audience test their understanding of the problem and the solution you've outlined - questions give them a chance to talk themselves into your approach, and give you a chance to show mastery of the subject.

How it works

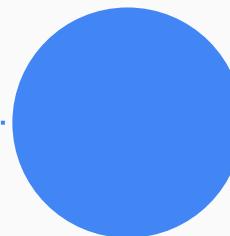
Step 1

Lorem ipsum dolor sit amet, consectetur



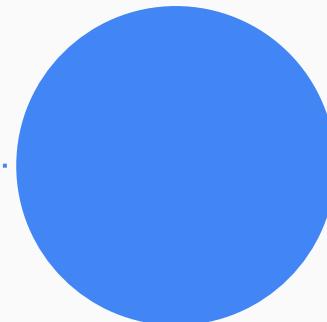
Step 2

Lorem ipsum dolor sit amet, consectetur



Step 3

Lorem ipsum dolor sit amet, consectetur

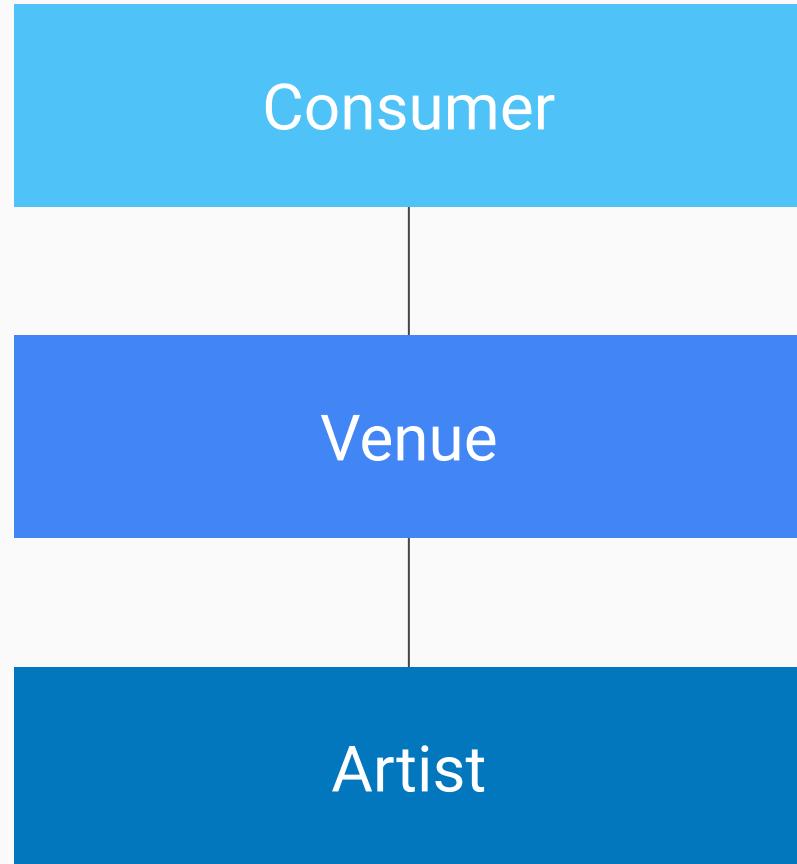


The background of the slide features a nighttime cityscape of a major metropolitan area, likely New York City, viewed from an elevated vantage point. The Empire State Building is prominent in the center-left, its Art Deco spire reaching towards a sky filled with scattered clouds. To the right, the One World Trade Center stands tall, its distinctive shape partially visible. The city is densely packed with numerous skyscrapers, their windows glowing with various shades of light. In the foreground, the dark silhouettes of buildings create a sense of depth. The overall atmosphere is one of a bustling urban environment at dusk or night.

The technology:
GPS + RFID

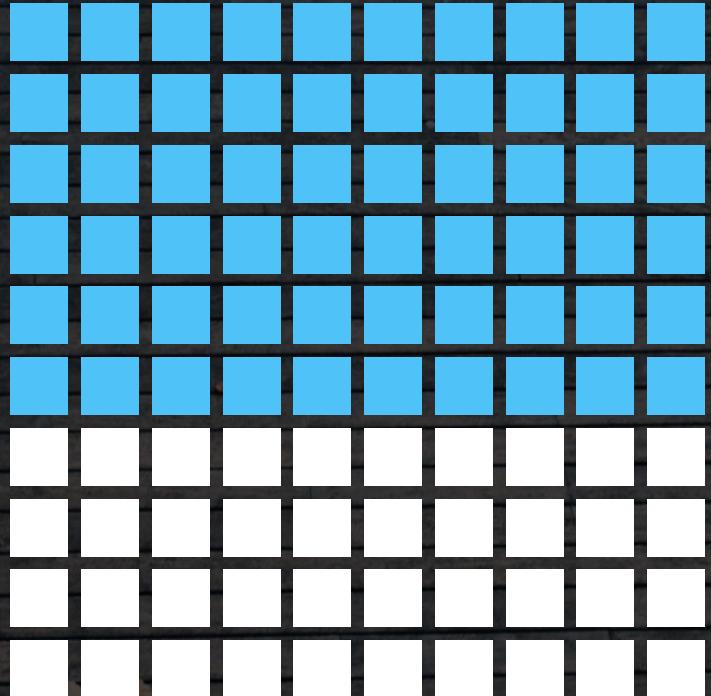
Revenue model

Lorem ipsum dolor sit amet,
consectetur adipiscing elit,
sed do eiusmod tempor
incididunt



Why now?

 Lorem ipsum dolor sit
 amet, consectetur
 adipiscing elit, sed do
 eiusmod tempor
 incididunt



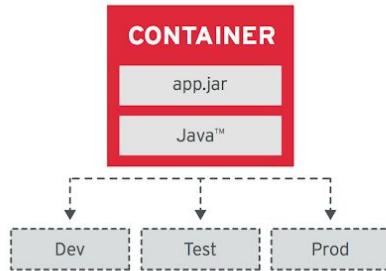
Day 2

Manage Your Enterprise Application with Kubernetes (K8S)



Principles of Container-based

Image Immutability Principle



High Observability Principle



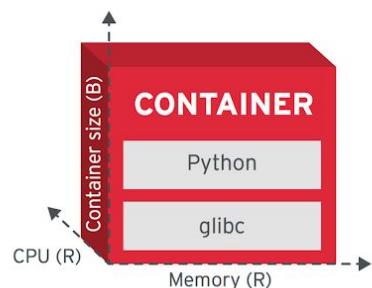
Process Disposability Principle



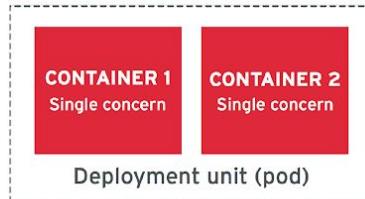
Lifecycle Conformance Principle



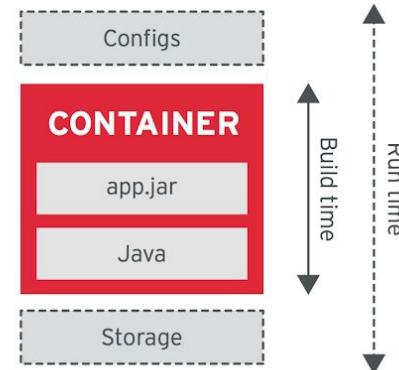
Runtime Confinement Principle



Single Concern Principle



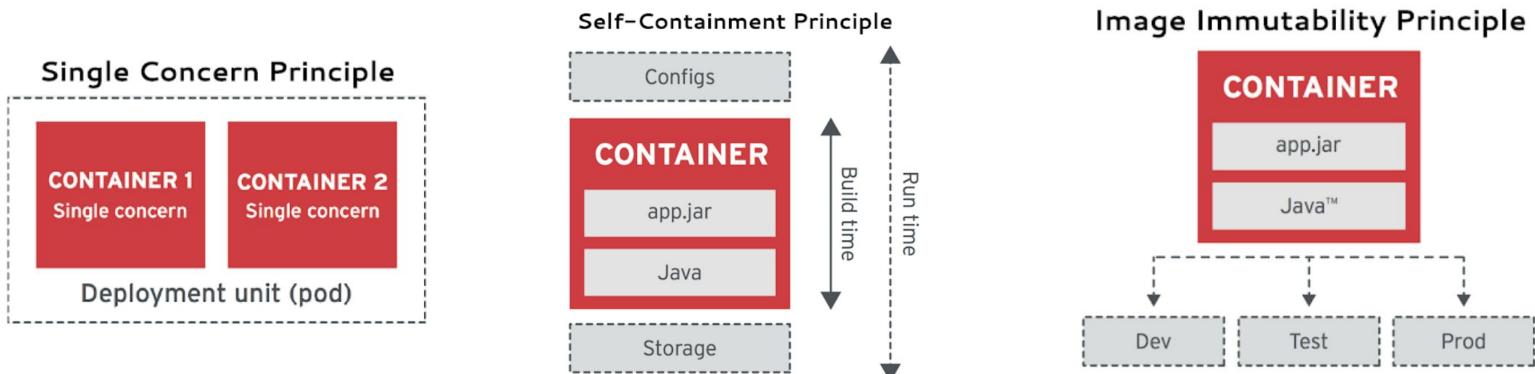
Self-Containment Principle



Principles of Container-based

Build Time

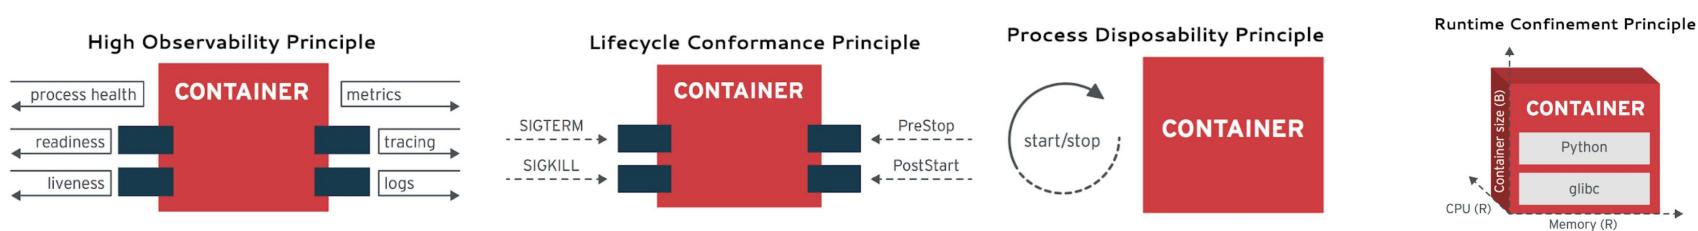
- **Single Concern:** Each container addresses a single concern and does it well.
- **Self-Containment:** A container relies only on the presence of the Linux kernel. Additional libraries are added when the container is built.
- **Image Immutability:** Containerized applications are meant to be immutable, and once built are not expected to change between different environments.



Principles of Container-based

Run Time

- **High Observability:** Every container must implement all necessary APIs to help the platform observe and manage the application in the best way possible.
- **Lifecycle Conformance:** A container must have a way to read events coming from the platform and conform by reacting to those events.
- **Process Disposability:** Containerized applications must be as ephemeral as possible and ready to be replaced by another container instance at any point in time.
- **Runtime Confinement:** Every container must declare its resource requirements and restrict resource use to the requirements indicated.





kubernetes

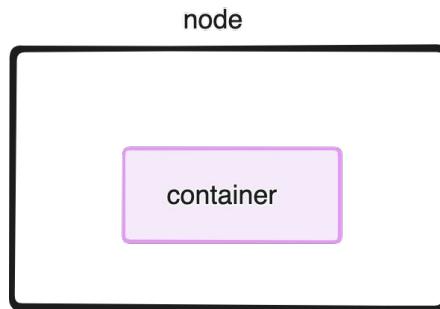
Agenda

- Introduction
- Kubernetes Architecture
- Resources / Objects
- Namespace
- RBAC (Role Based Access Control)
- HELM
- Monitoring
- HPA (Horizontal Pod Autoscaler)
- Deployment & Pipelines

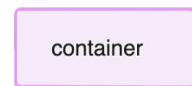
Introduction



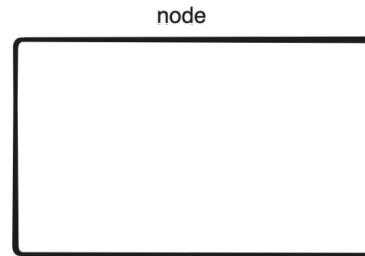
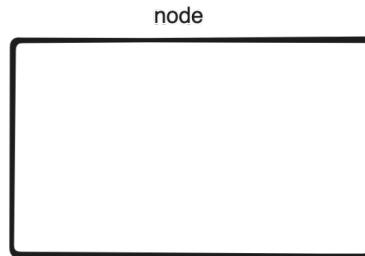
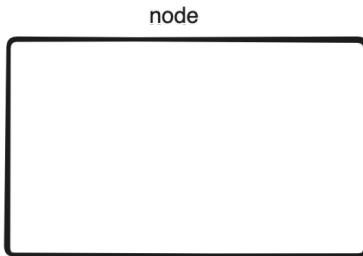
Docker on one server



Docker on multi servers

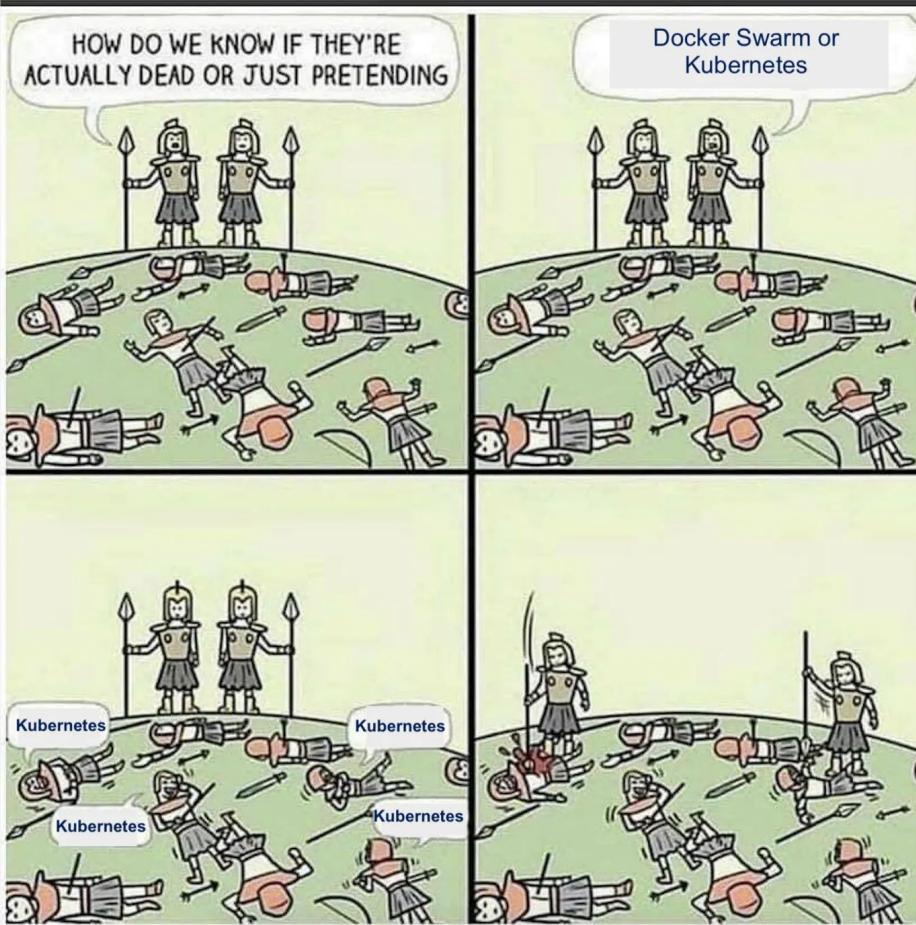


???

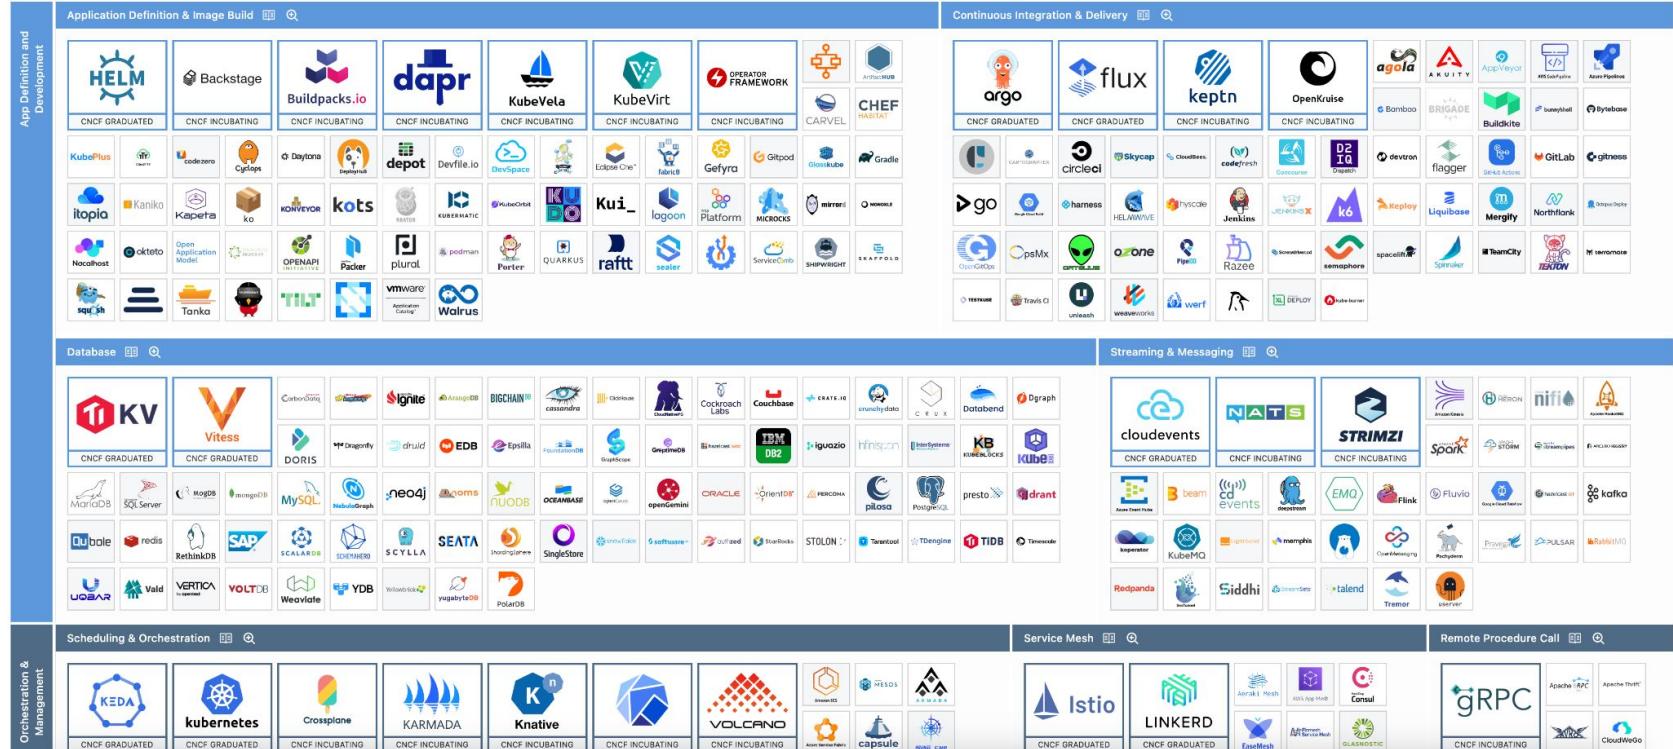


What is Kubernetes?

- Kubernetes in Greek means the Helmsman
- Also known as k8s
- Inspired and informed by Google's experiences and internal systems
- 100% opensource, written in GO
- Container orchestrator, runs and manages containers
- Kubernetes v1.0 released on July 2015, join CNCF on March 2016
- Container choose Docker, Container Orchestrator choose Kubernetes



CNCF Landscape



<https://landscape.cncf.io/>

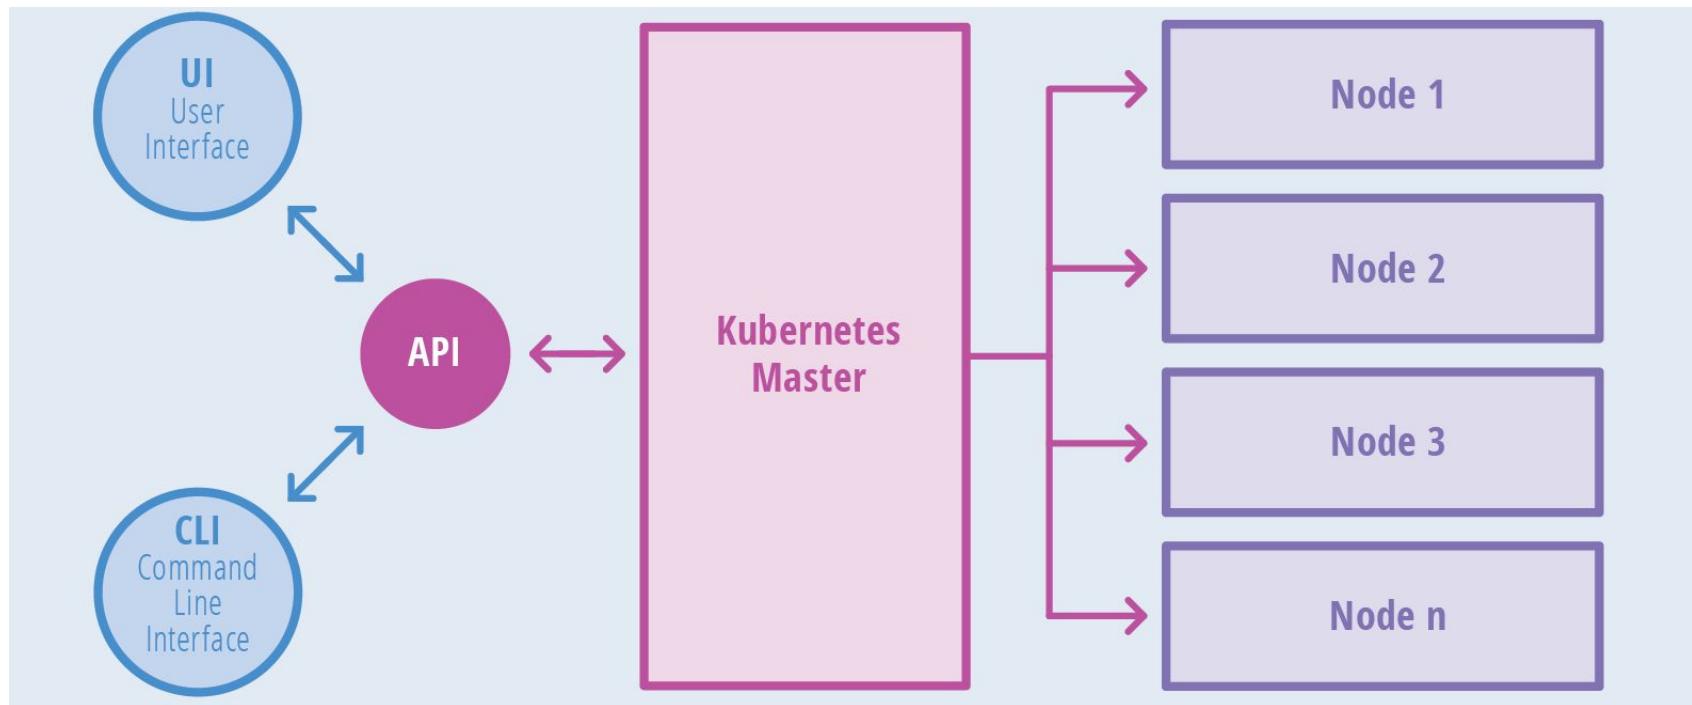
The features of Kubernetes

- Automatic bin packing
- Self healing
- Automated rollout & rollback
- Service Discovery & load Balancing
- Horizontal scaling
- Declarative Configuration

Kubernetes Architecture



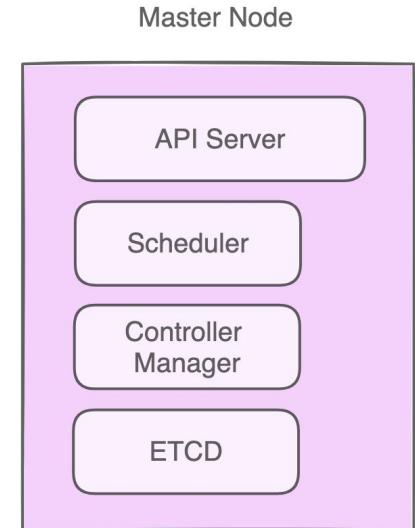
Kubernetes Architecture



Kubernetes Architecture

Master node

- Responsible for managing cluster
- Access master node via CLI, UI or API
- For fault tolerance, can be more than 1 master in cluster
- Include 4 components: **API Server, Scheduler, Controller manager and ETCD**

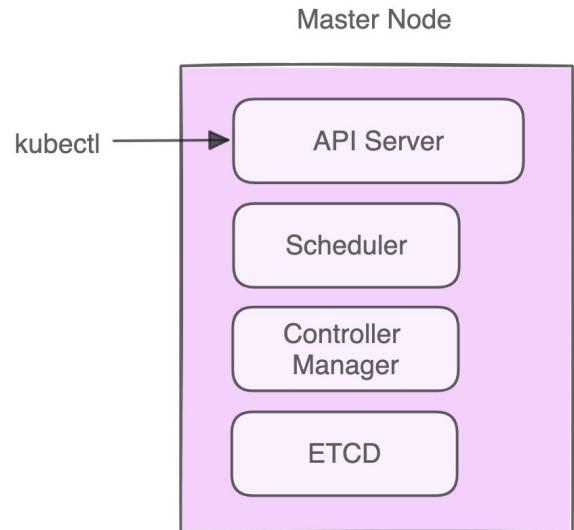


Kubernetes Architecture

Master node

API Server

- Centralized component where all the cluster components are communicated
- Execute and validate from user command

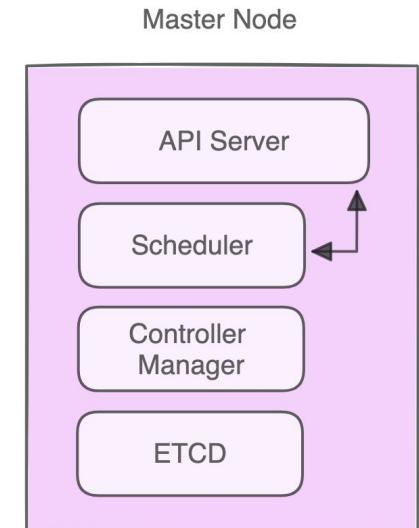


Kubernetes Architecture

Master node

Scheduler

- Assigning the application to worker node
- Automatically detect which pod to place on which node based on all factors

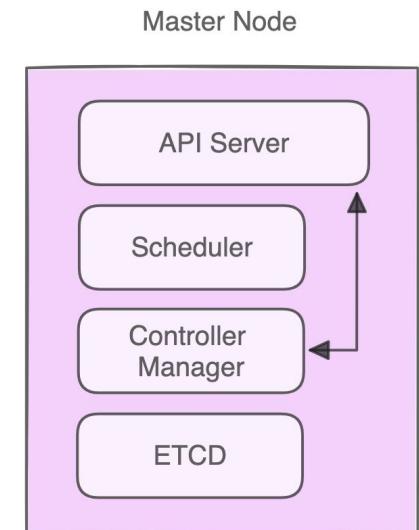


Kubernetes Architecture

Master node

Controller Manager

- The Controller Manager maintains the cluster
- Handles node failures, replicating components, maintaining the correct number of pods, etc.

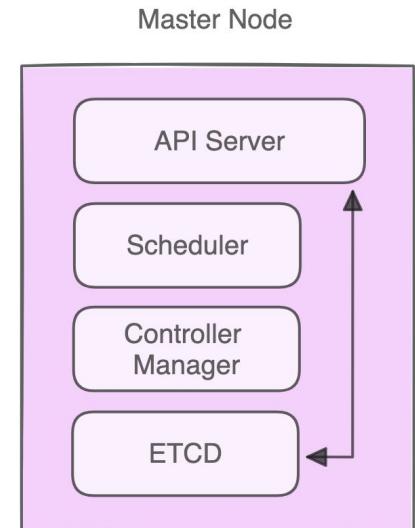


Kubernetes Architecture

Master node

ETCD

- ETCD is a distributed reliable key-value store used by Kubernetes to store all data used to manage the cluster
- When you have multi master node in cluster, ETCD stores all information in a distributed manner
- Application data doesn't store in ETCD

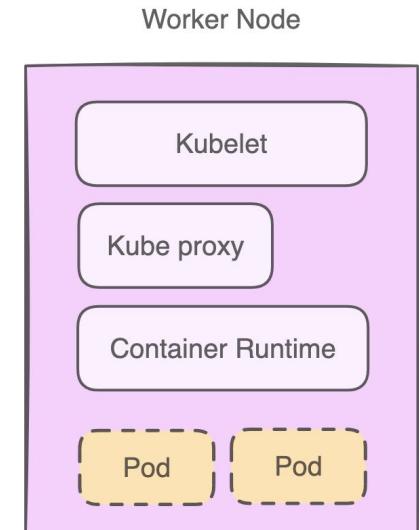


Kubernetes Architecture

Worker node

- Responsible about applications
- Include 3 components: **Container Runtime**, **Kubelet** and

Kube Proxy

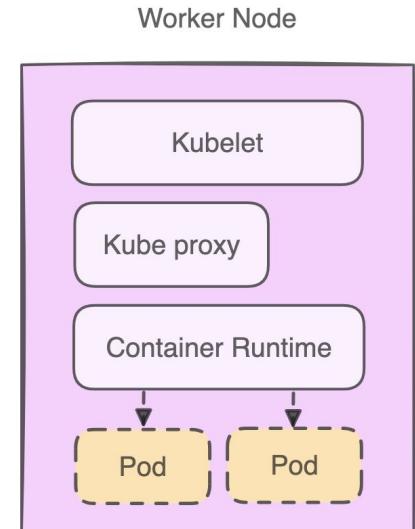


Kubernetes Architecture

Worker node

Container Runtime

- Container runtime which runs the containers like Docker, crio or containerd
- Pulling the images and run the containers

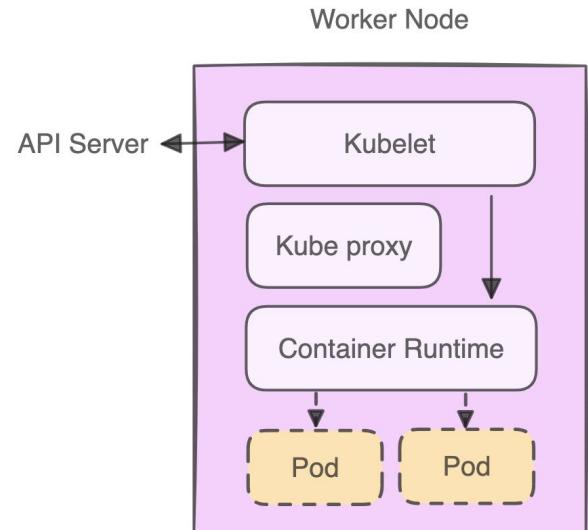


Kubernetes Architecture

Worker node

Kubelet

- Interacting with the master node
- Contact the container runtime for create pod depend on Pod Spec.
- Provide health information of the worker node

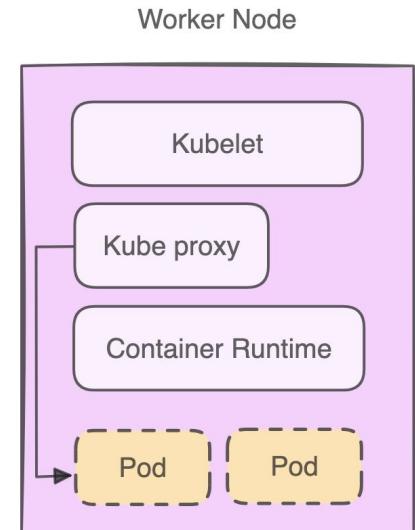


Kubernetes Architecture

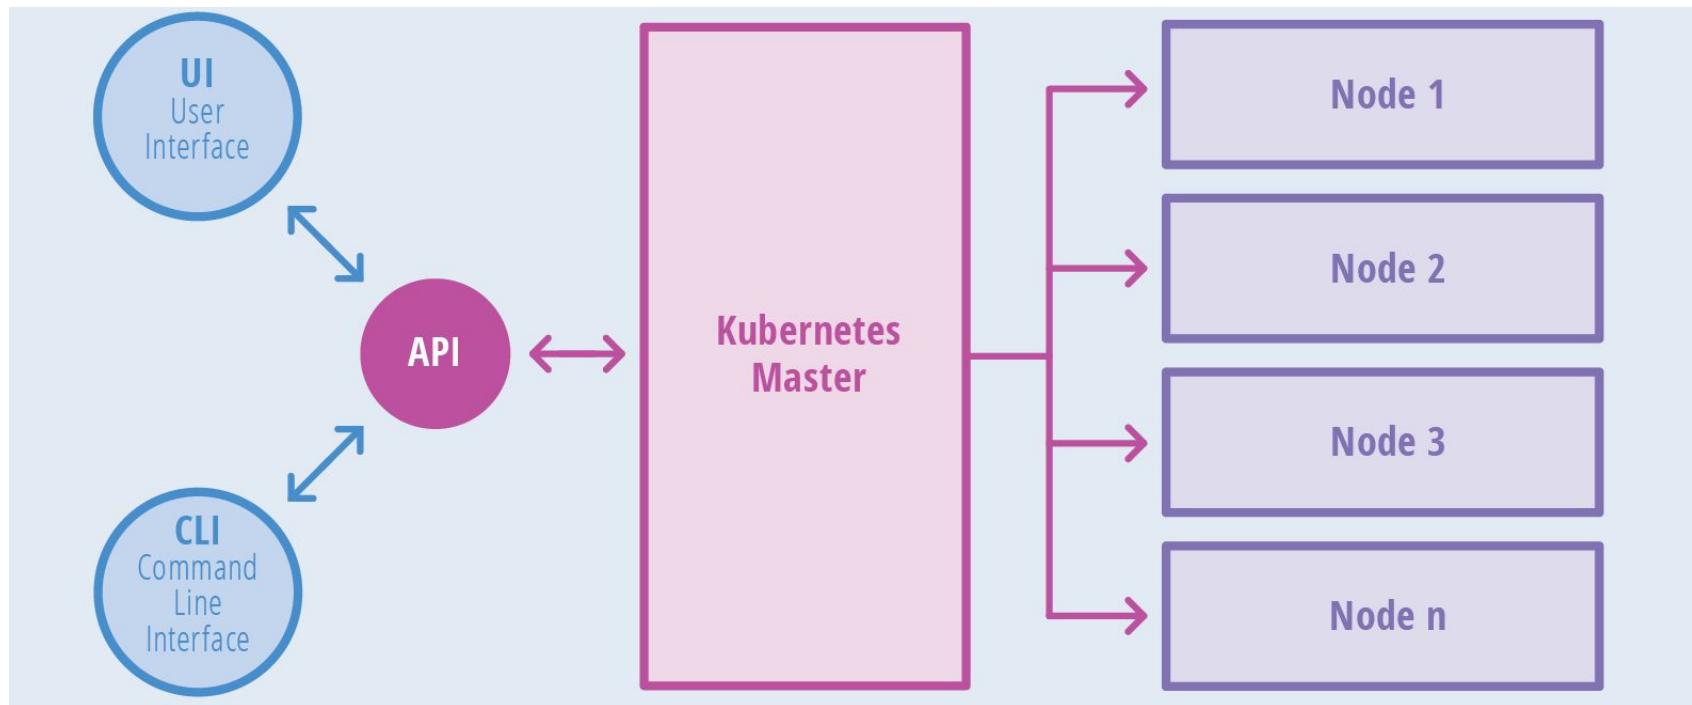
Worker node

Kube Proxy

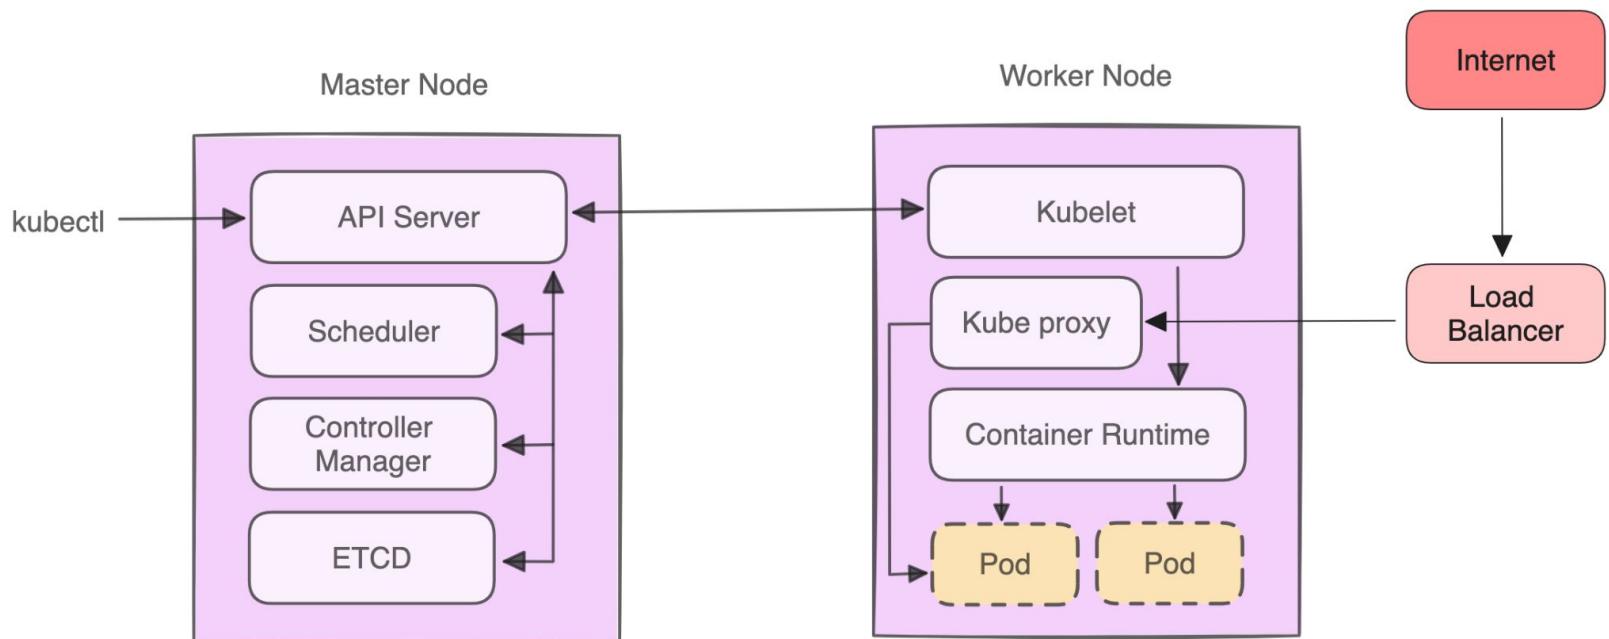
- Responsible for ensuring network traffic is routed properly to internal and external services
- Manage route for reduce latency time for example pod will communicate with another pod on same node first



Kubernetes Architecture



Kubernetes Architecture



Kubectl

Kubectl

- Command line
- Control the Kubernetes cluster manager
- CRUD Kubernetes resources

Kubectl

- \$ kubectl apply
- \$ kubectl create
- \$ kubectl get
- \$ kubectl describe
- \$ kubectl delete

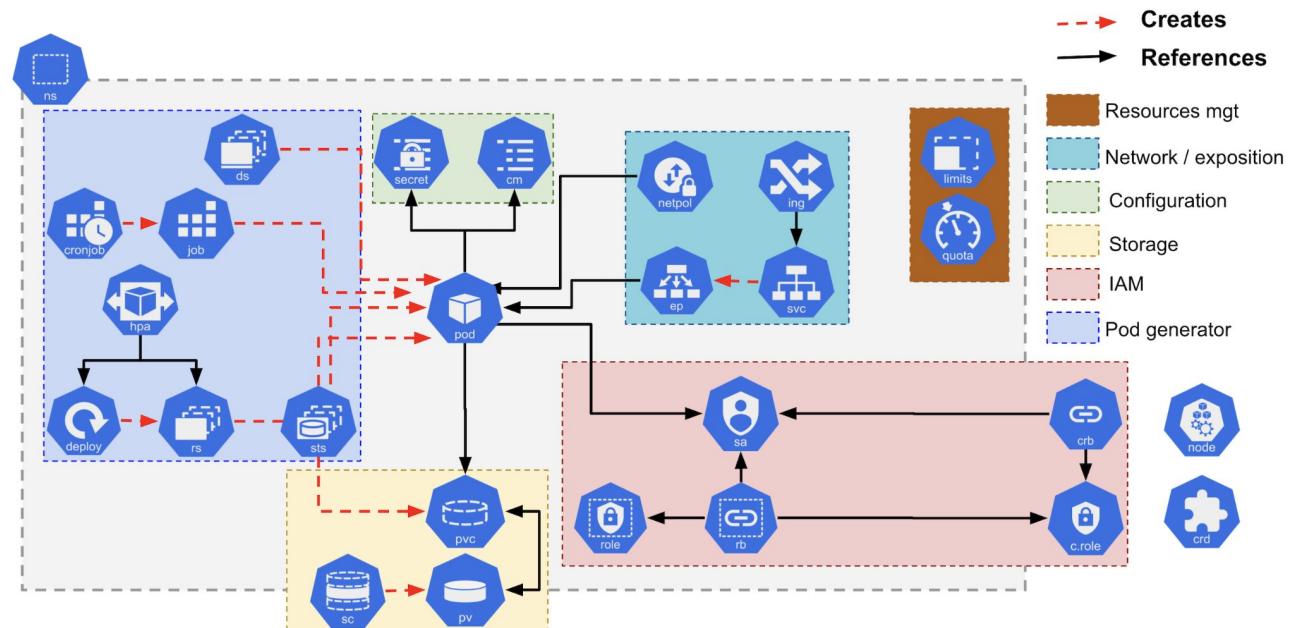
Resources / Objects



Resources / Objects

- Pod
- Deployment
- Service
- Ingress
- Config map
- Secret
- PV & PVC

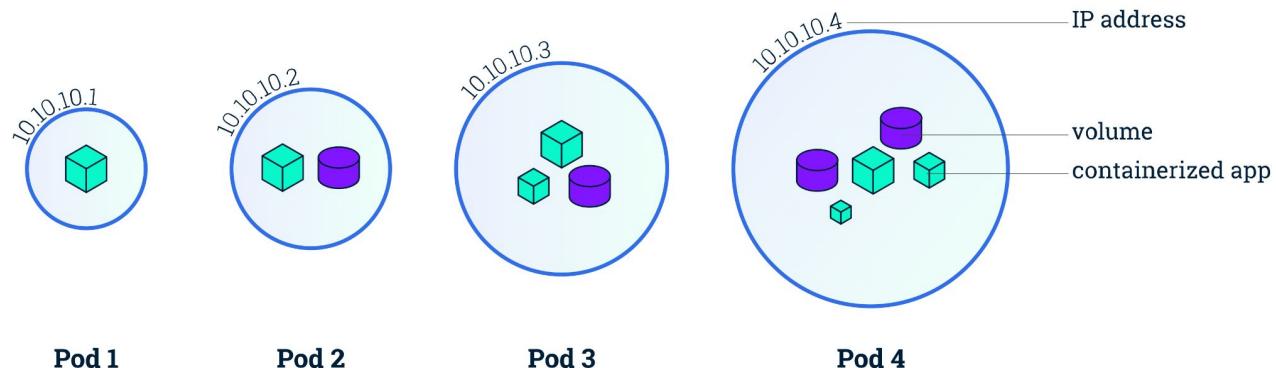
Kubernetes Resources Map



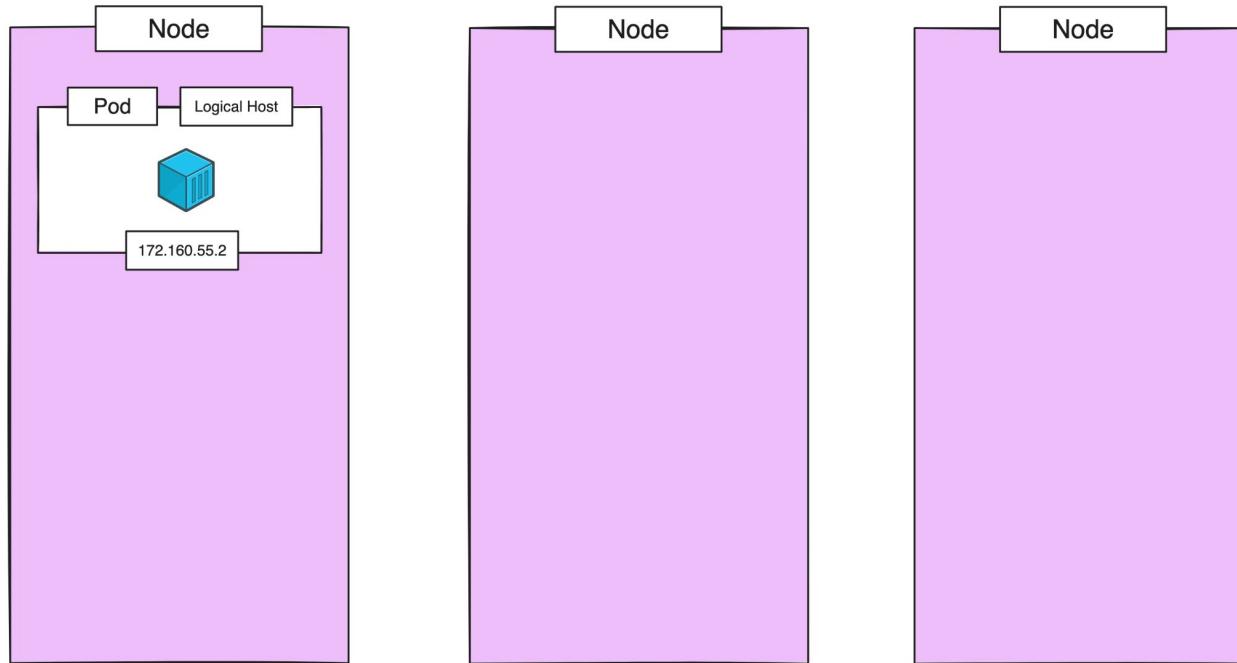
Pod

Pod

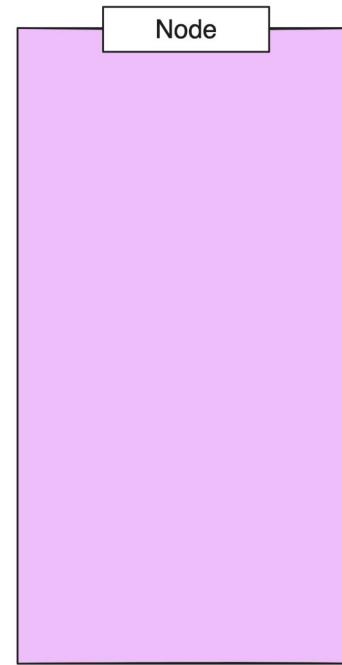
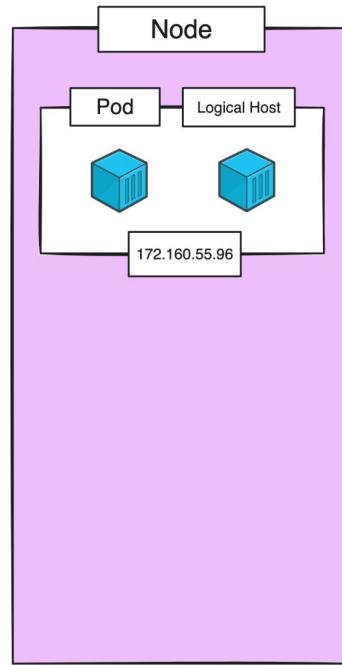
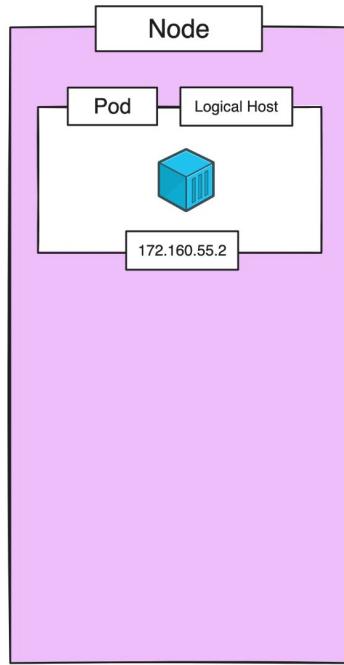
- Smallest deployable units of computing that you can create and manage in Kubernetes
- Logical Application
- One or more container in Pod (recommend 1 container per pod)
- Any containers in same pod will share storage volumes and network resources
- Each pod is assigned a unique ip



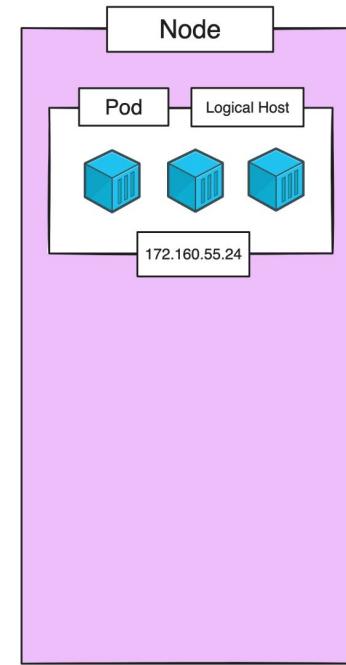
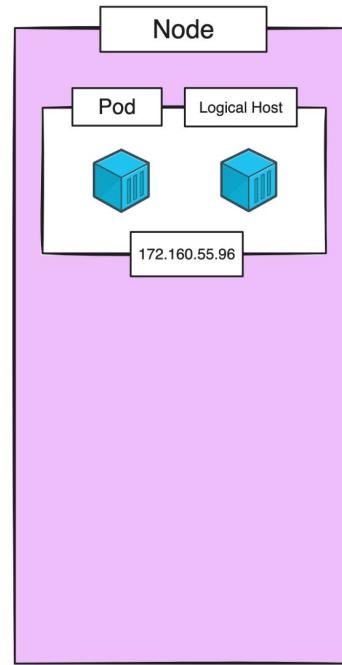
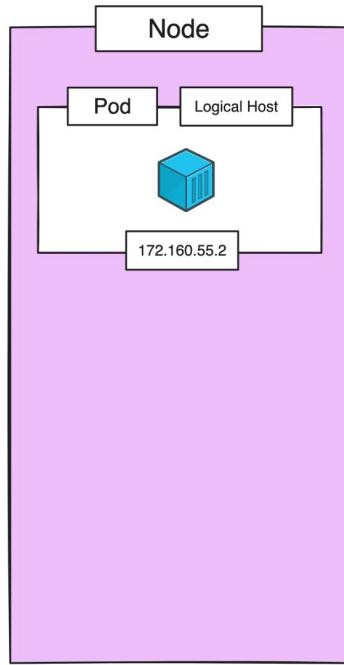
Node & Pod



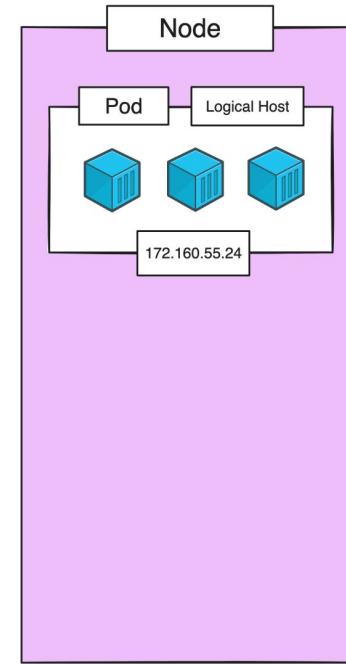
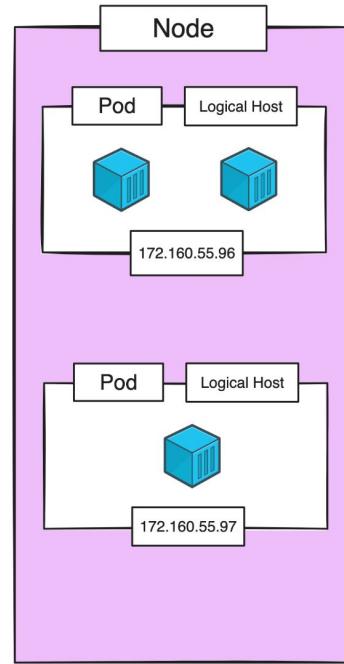
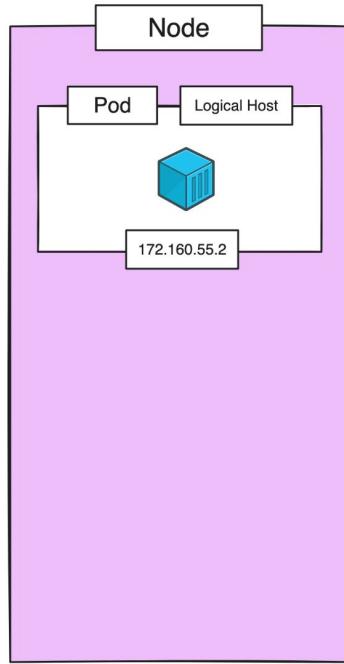
Node & Pod



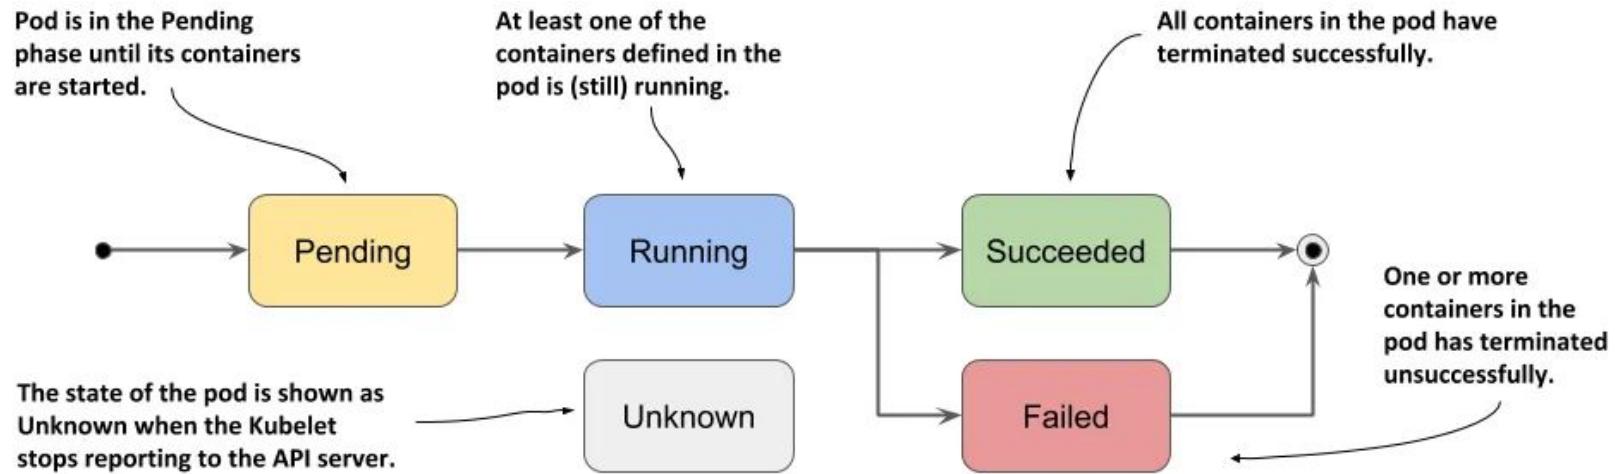
Node & Pod



Node & Pod

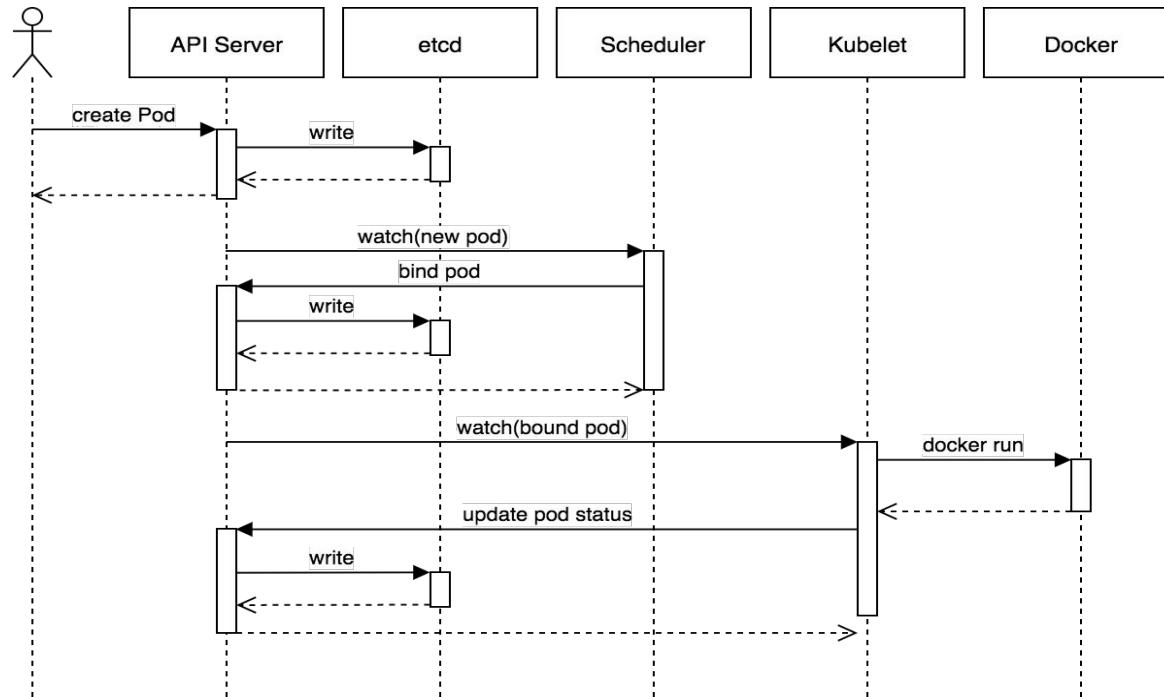


States of a Pod



<https://medium.com/@seifeddinerajhi/navigating-the-journey-of-a-pod-a-guide-to-the-explaining-world-of-pod-lifecycle-a1fbc2c98c55>

Creation of a Pod



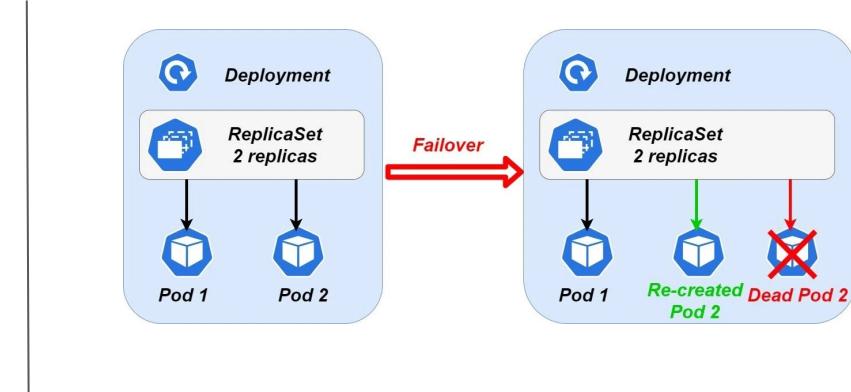
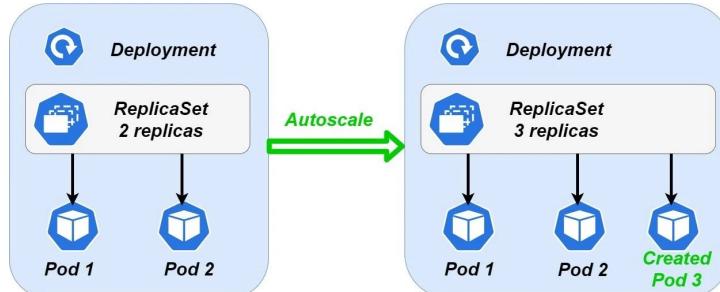
<https://medium.com/@seifeddinerajhi/navigating-the-journey-of-a-pod-a-guide-to-the-explaining-world-of-pod-lifecycle-a1fbc2c98c55>

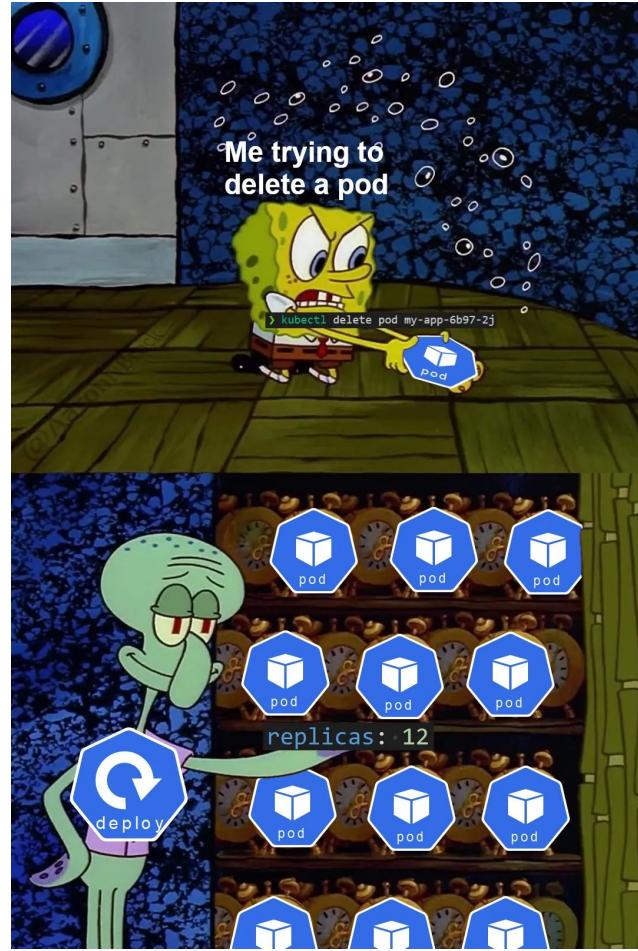
Workshop Pod

Deployment

Deployment

- A Deployment provides declarative updates for Pods and ReplicaSets
- You describe a **desired state** in a Deployment, and the Deployment Controller **changes the actual state to the desired state** at a controlled rate
- Manage for rollout and rollback
- Deployment => ReplicaSet => pod

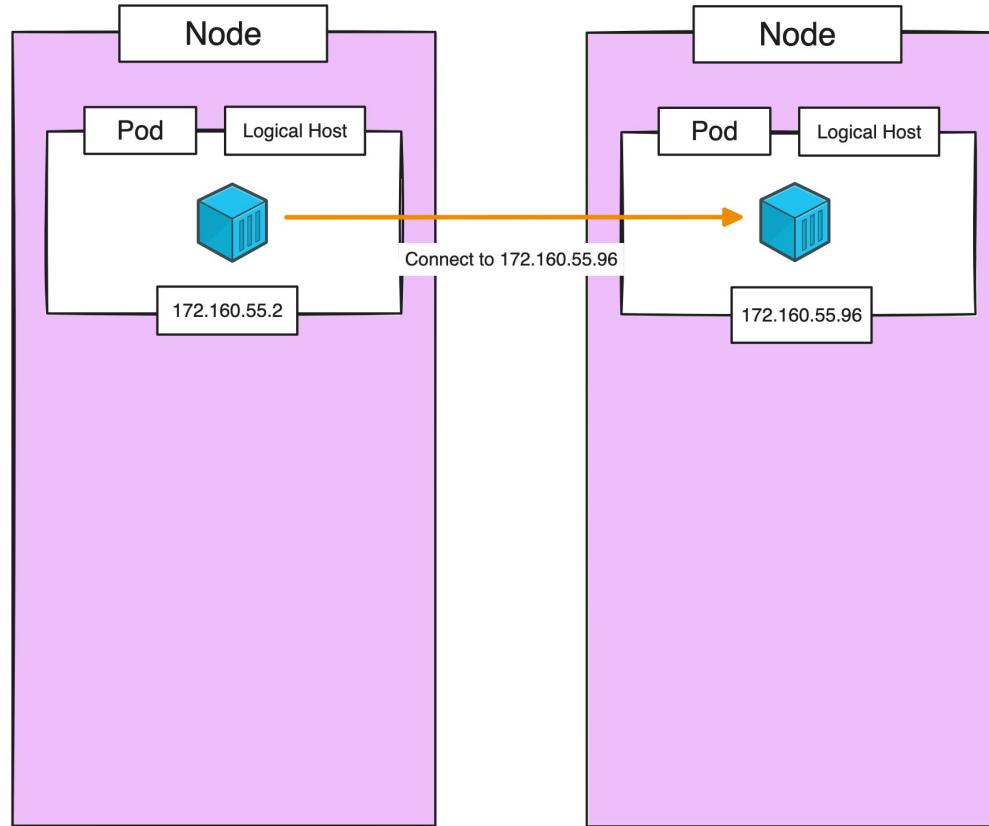




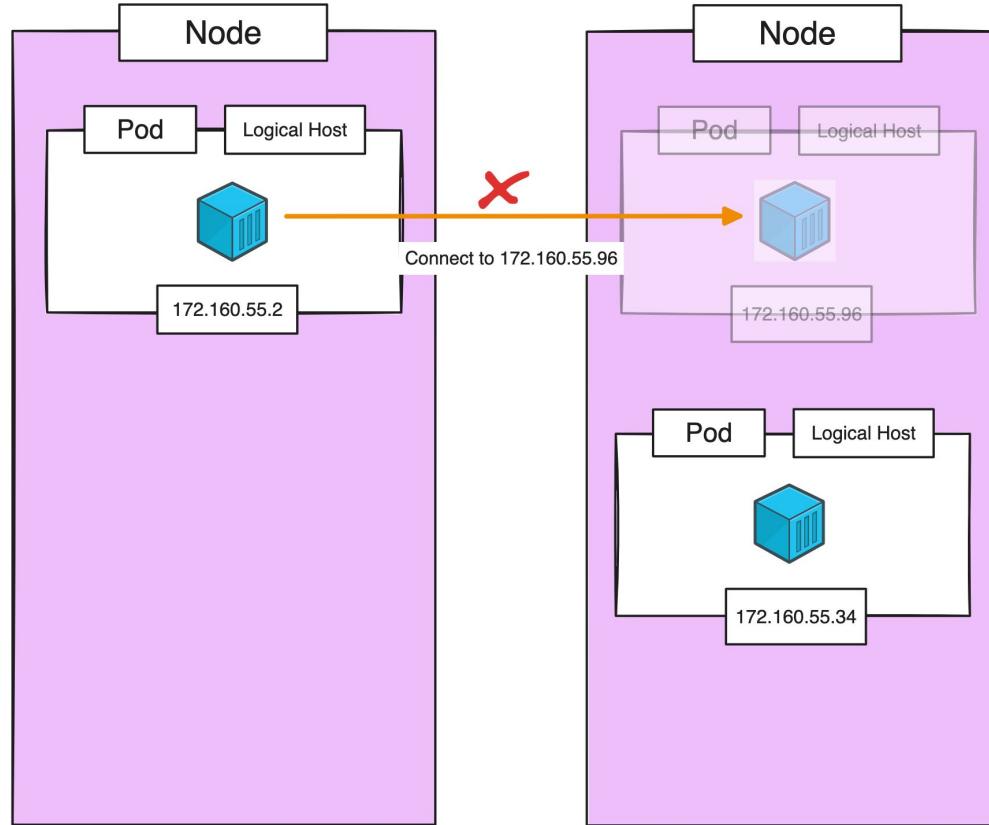
Workshop

Deployment

Problem



Problem

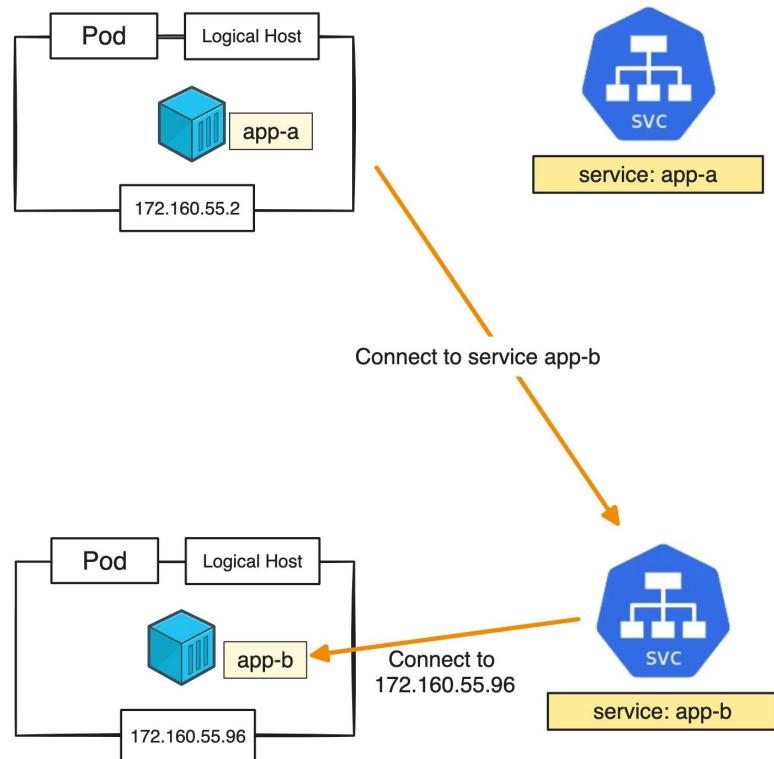


Service

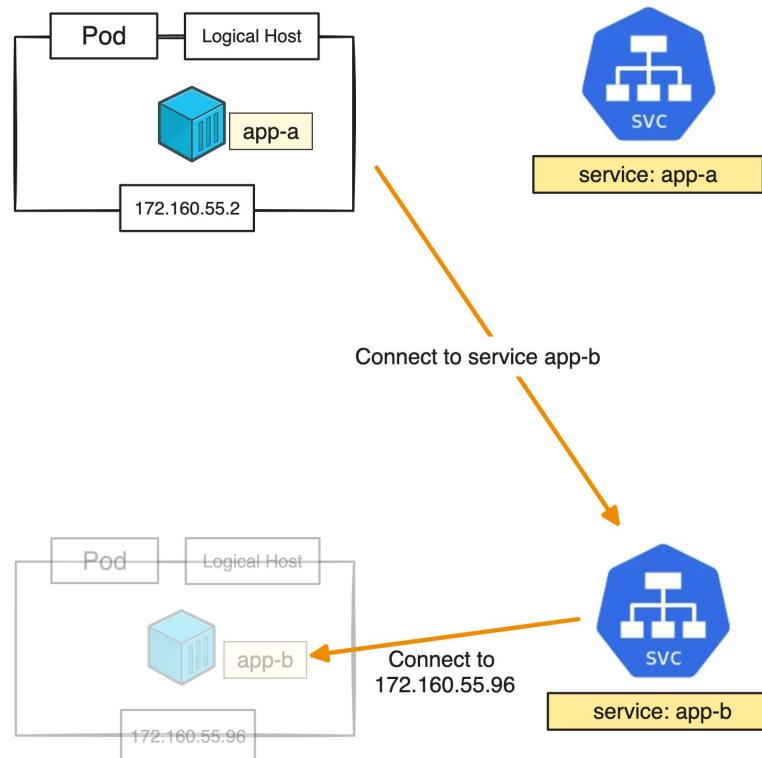
Service

- Acts as a **virtual endpoint** for your application
- Providing a stable and discoverable way to access your pods without worrying about their individual IP addresses or network details
- Logical load balancer
- Communicating for outside and inside of cluster
- 4 types:
 - ClusterIP
 - NodePort
 - LoadBalancer
 - ExternalName

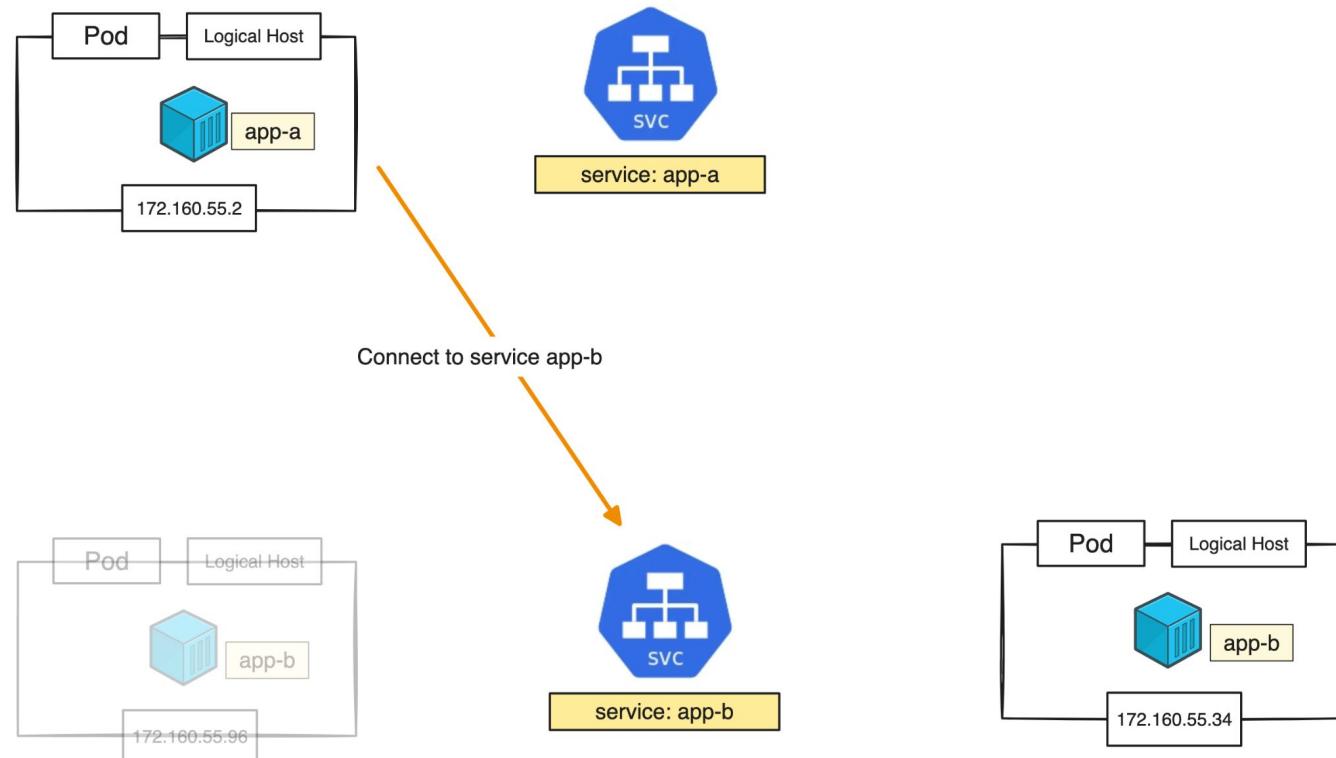
Service: Stable endpoint



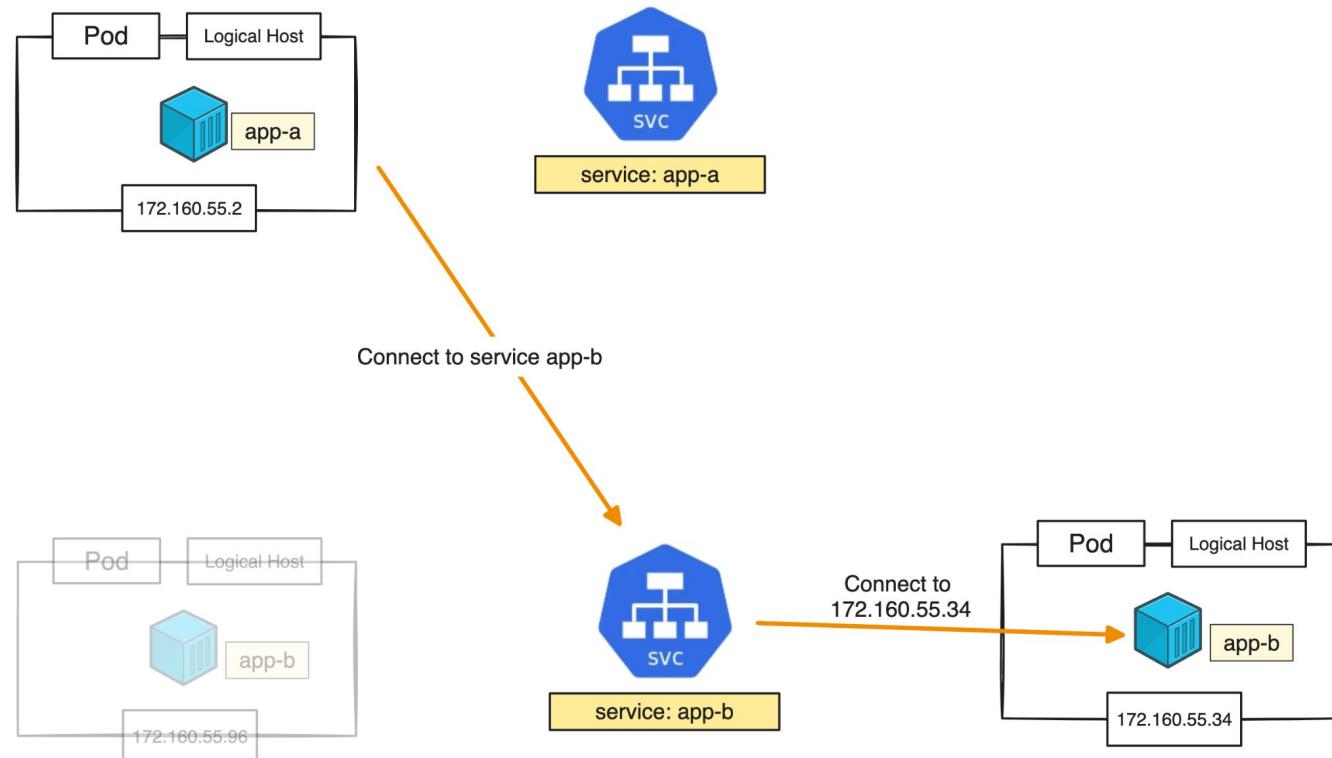
Service: Stable endpoint



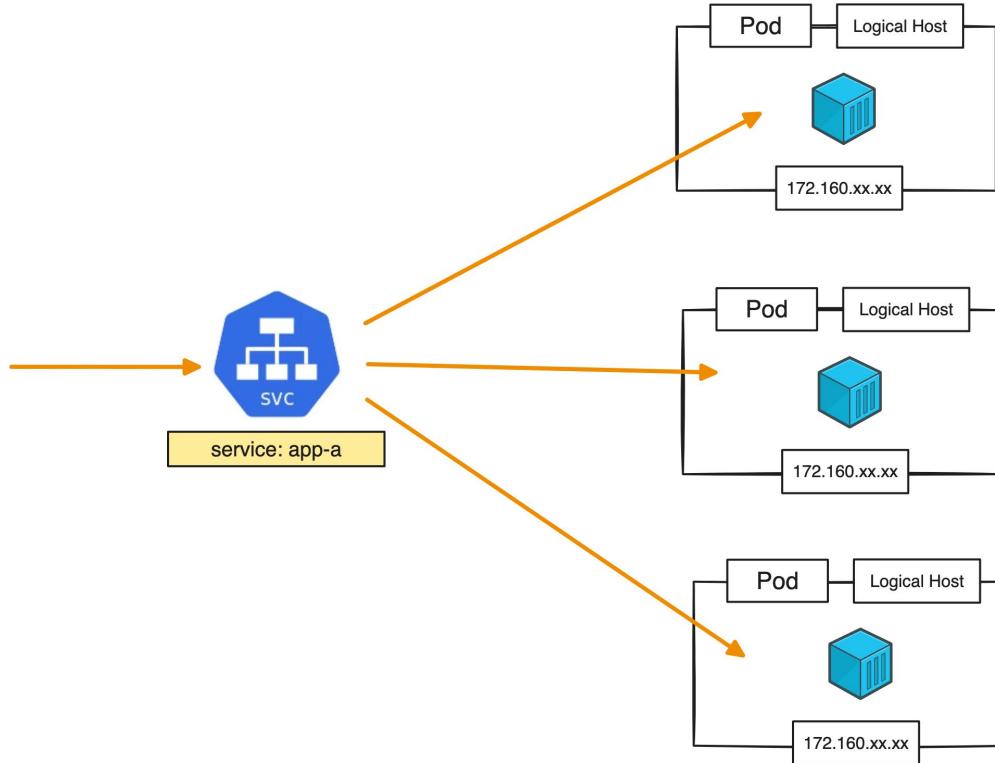
Service: Stable endpoint



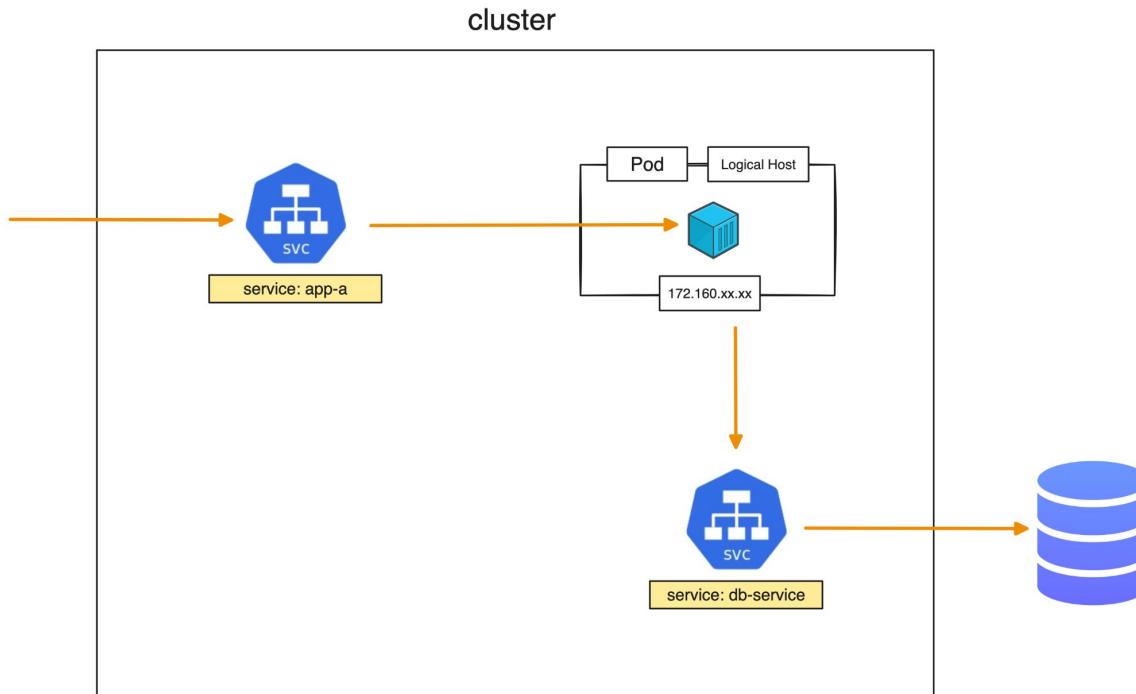
Service: Stable endpoint



Service: Load balancer



Service: Communicating inside & outside of cluster



Service type: ClusterIP

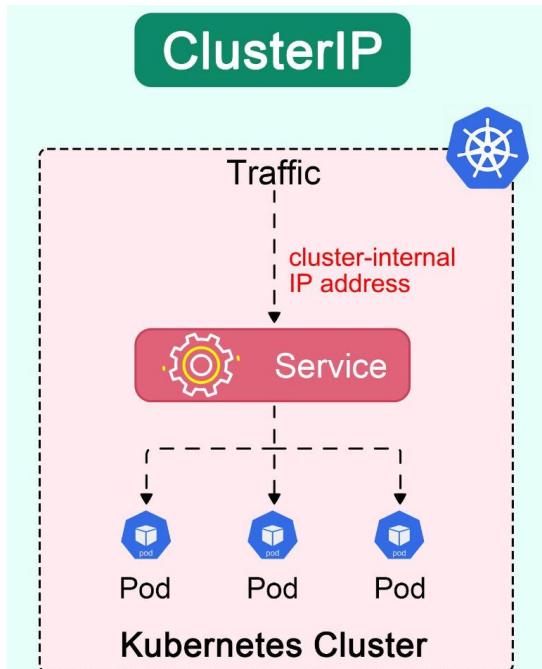
- ClusterIP is the default service type in Kubernetes
- It exposes the service on an internal IP address reachable **only within the cluster**

Advantages

- Suitable for inter-pod communication and internal services

Disadvantages

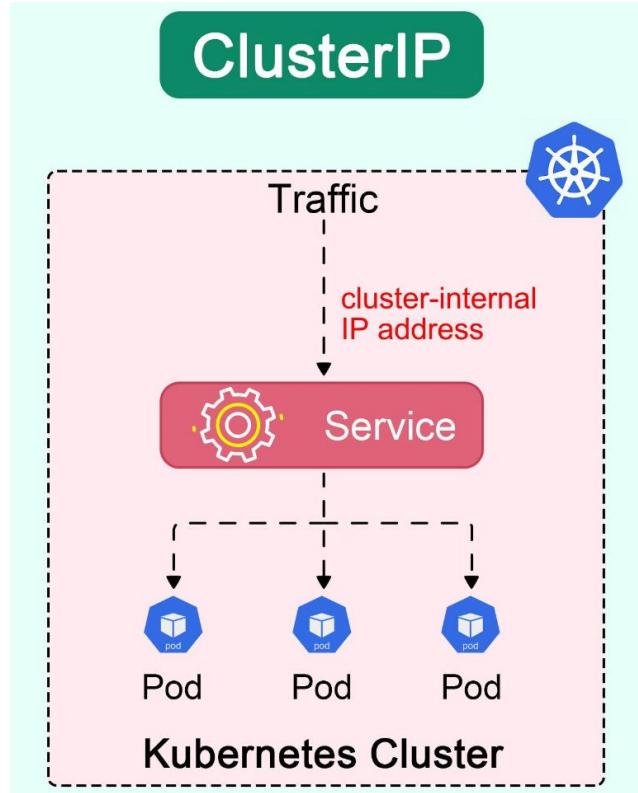
- Limited to internal cluster communication
- Not accessible externally



Service type: ClusterIP

Use Cases

- Inter-service communication within the cluster.
For example, communication between the front-end and back-end components of your app



Service type: NodePort

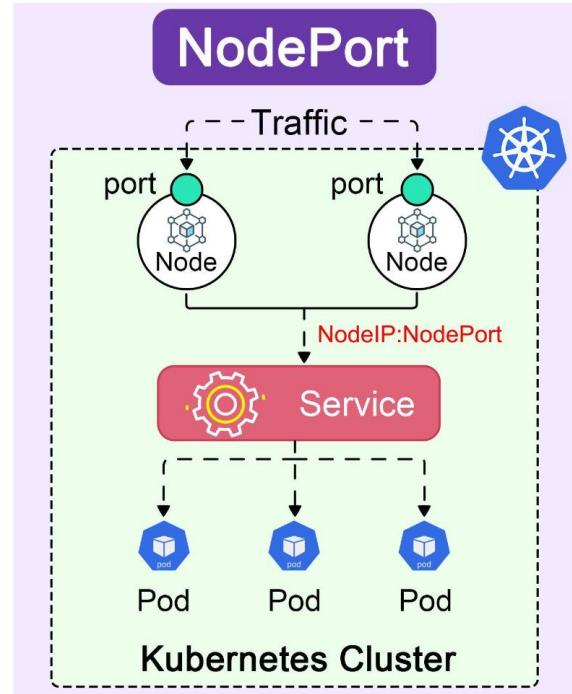
- It enables **external access** to the service using the cluster **nodes' IP address** and the assigned **NodePort**
- When a NodePort service is created, Kubernetes allocates a port from a predefined range (default: **30000-32767**)

Advantages

- Provides external access to the service
- Suitable for development and testing environments

Disadvantages

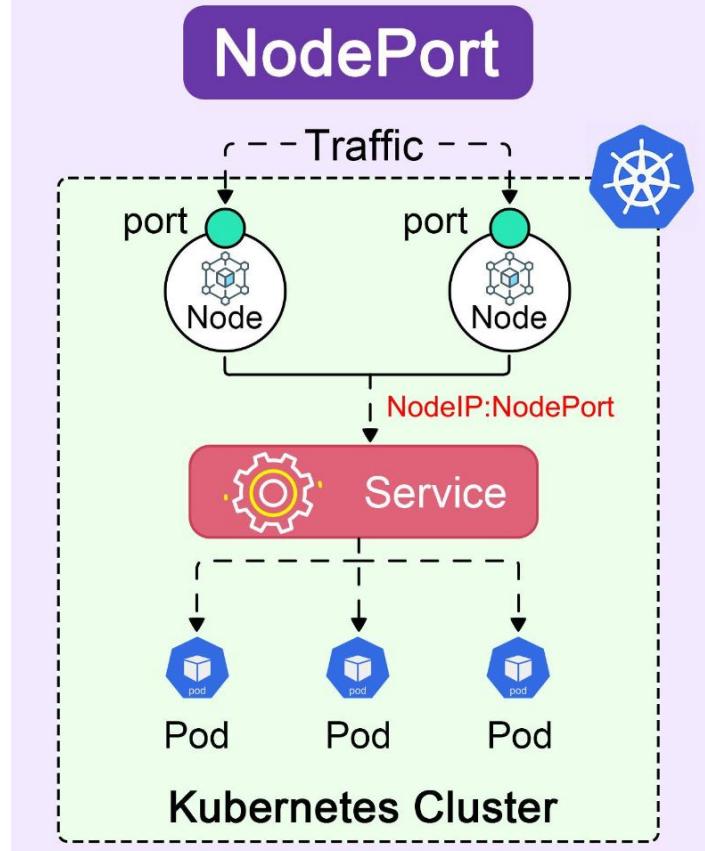
- Requires manual management of port allocations
- IP Node maybe have changed
- Should not use in production



Service type: NodePort

Use Cases

- When you want to enable external connectivity to your service
- Best for testing access for a moment time



Service type: LoadBalancer

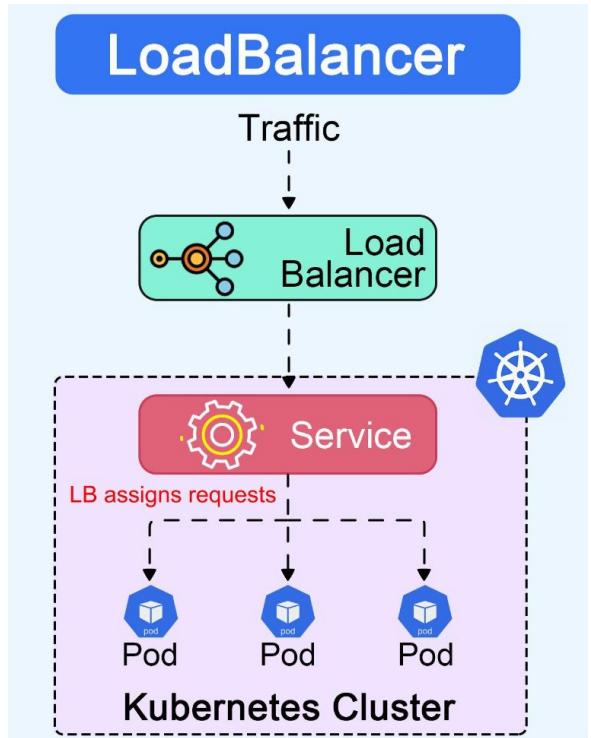
- LoadBalancer is a service type that **provisions an external load balancer** to distribute traffic across the pods
- Typically provided by a cloud provider's infrastructure
- When a LoadBalancer service is created, Kubernetes interacts with the cloud provider's API to an external load balancer

Advantages

- Provides external access to the service
- Automatically provisions an external load balancer for the service
- Supports scaling and high availability

Disadvantages

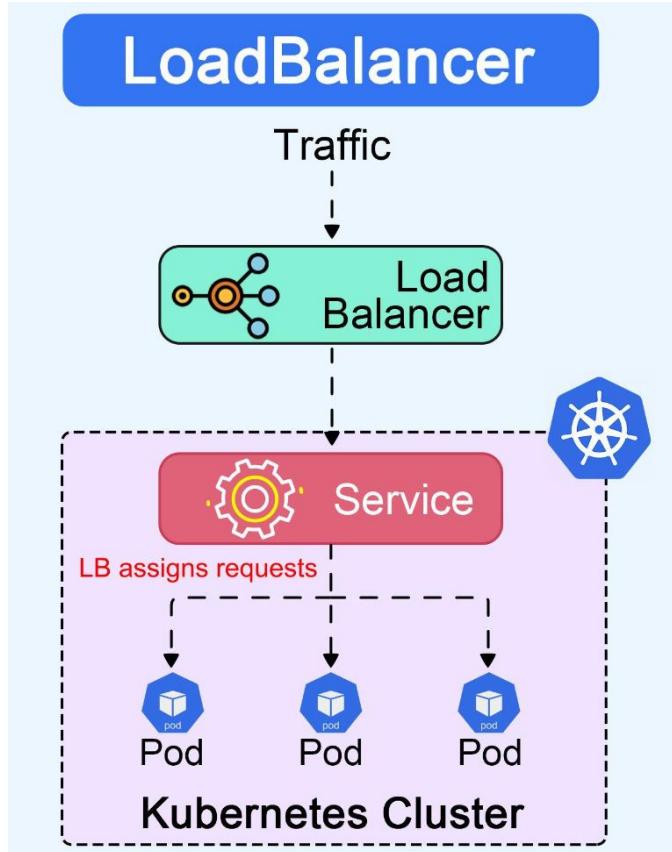
- Requires a cloud provider
- Additional costs for using the load balancer service on cloud provider



Service type: LoadBalancer

Use Cases

- When you want to enable external connectivity to your service
- When you are using a cloud provider to host your Kubernetes cluster

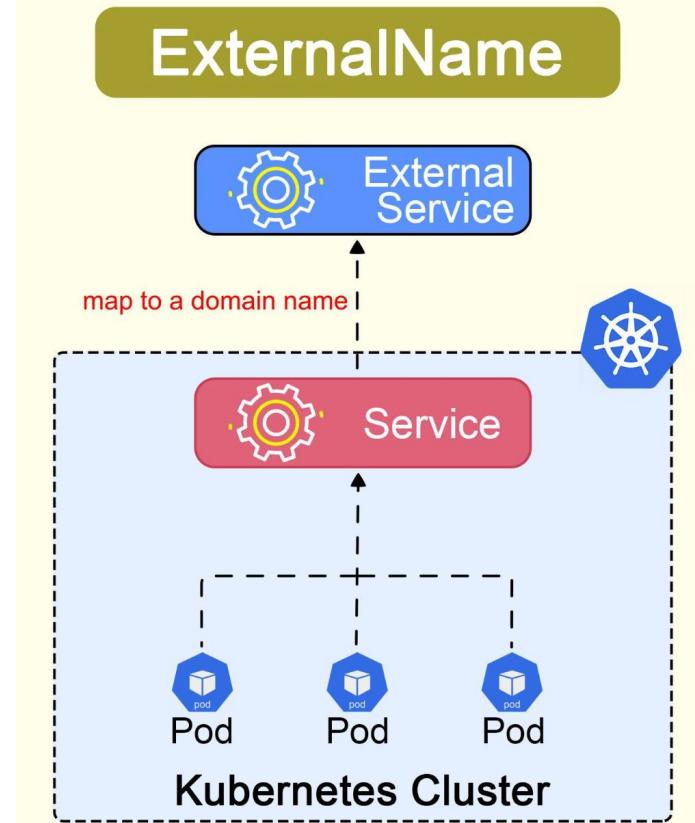


Service type: ExternalName

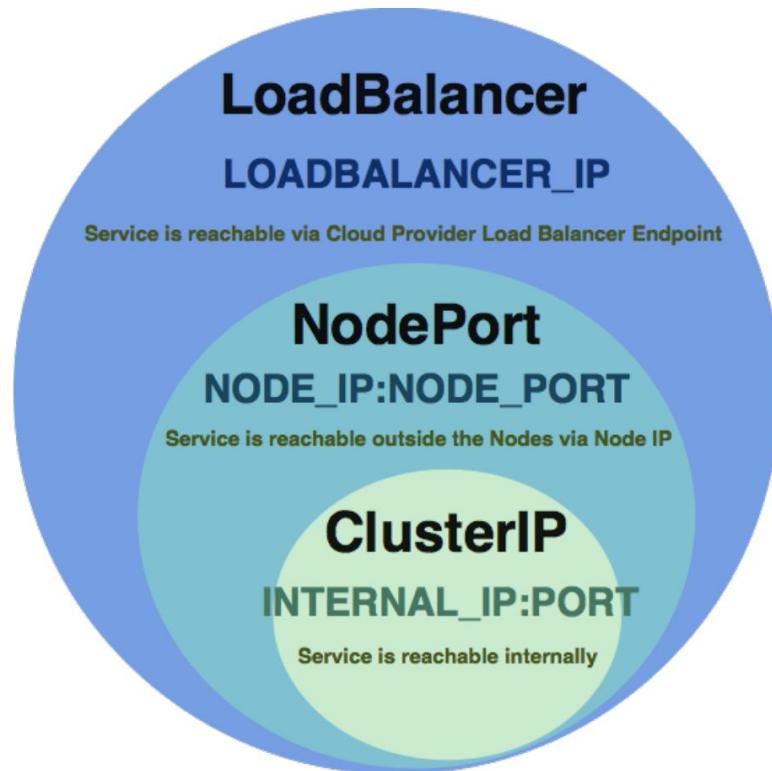
- An ExternalName Service is a special type of Kubernetes service that allows you to expose an existing external service (like a database or a third-party API) to your cluster without deploying any Pods

Use Cases

- Accessing an external database to your Kubernetes cluster
- Connecting a third-party API from within your cluster



Service type



Workshop

Service

Ingress

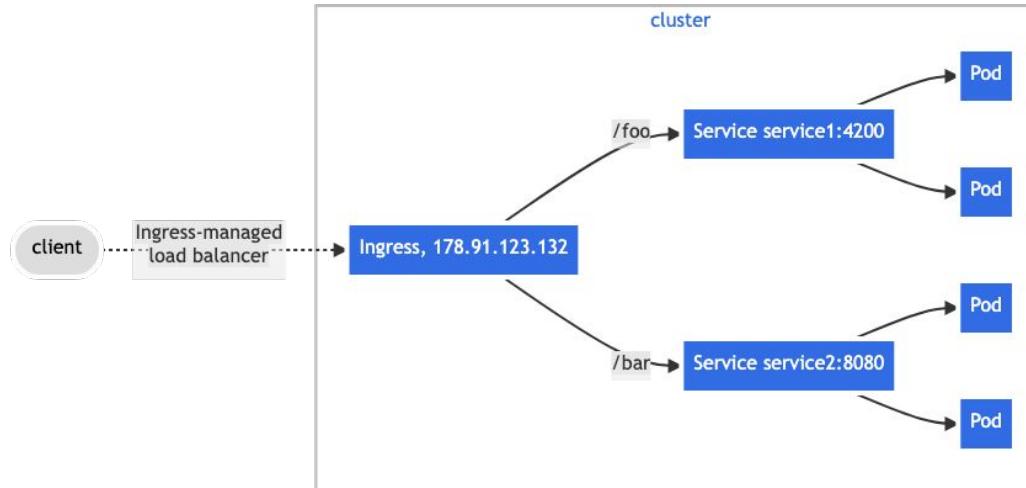
Ingress

- Routing from outside the cluster to services within the cluster
- Traffic routing is controlled by rules defined on the ingress resource
- An **Ingress controller** is responsible for fulfilling the Ingress



Ingress Controllers

- The Ingress Controller is an application hosted inside a Kubernetes cluster that actively **manages traffic** following Ingress Resources and pre-defined Ingress Rules
- Ingress controllers doesn't come with standard, have to be deployed separately
- Popular controllers include [ingress-nginx](#), [HAProxy](#), [Traefik](#)



Workshop

Ingress

ConfigMap & Secret

ConfigMap

- Store non-confidential data in key-value pairs
- Pods can consume ConfigMaps as environment variables
- The data stored in a ConfigMap cannot exceed 1 MiB
- If need to store larger than this limit, may want to consider mounting a volume or use a separate database or file service

Secret

- Secrets are similar to ConfigMaps but are specifically intended to hold confidential data
- Kubernetes Secrets are, by default, stored **unencrypted** in the API server's underlying data store (etcd)
- Anyone with API access can retrieve or modify a Secret
- In order to safely use Secrets, take at least the following steps
 1. Enable Encryption at Rest for Secrets
 2. Enable or configure RBAC rules with least-privilege access to Secrets
 3. Restrict Secret access to specific containers
 4. Consider using external Secret store providers

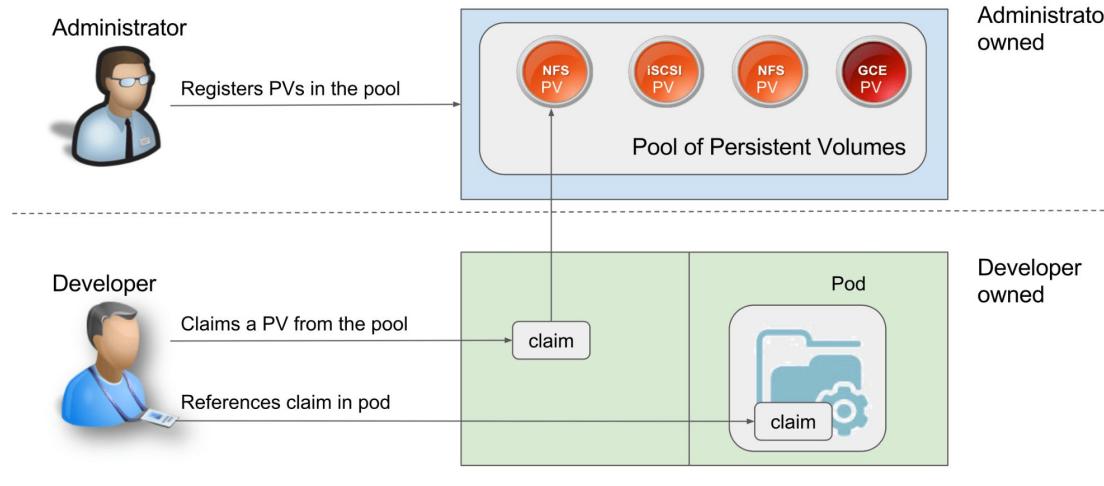
Workshop

ConfigMap & Secret

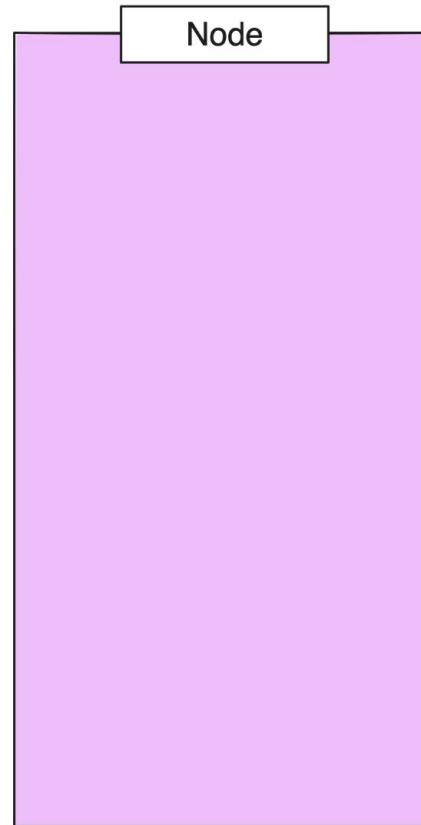
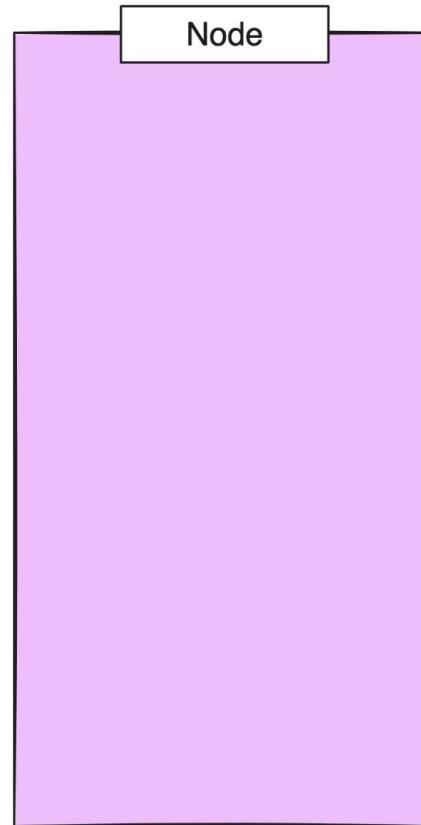
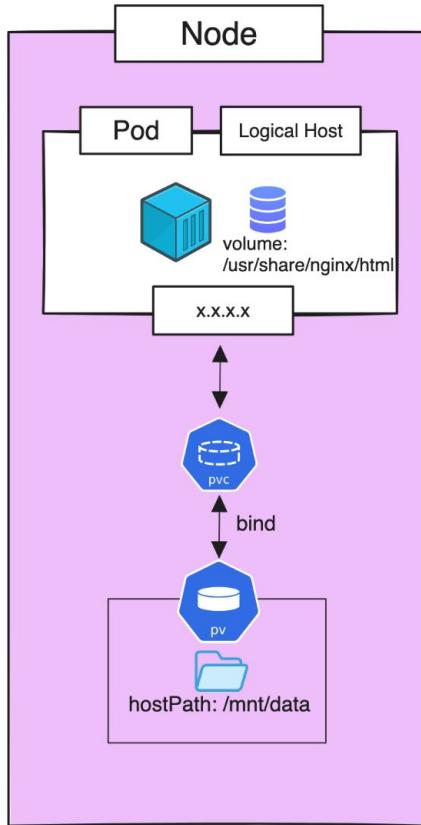
Persistent Volume

Persistent Volume

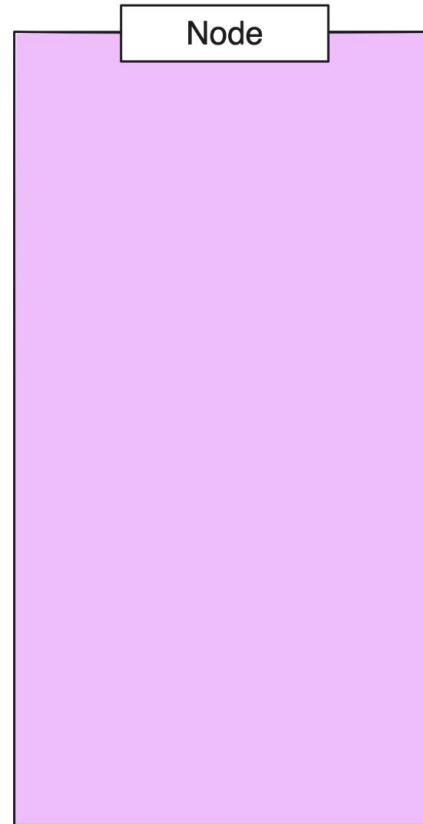
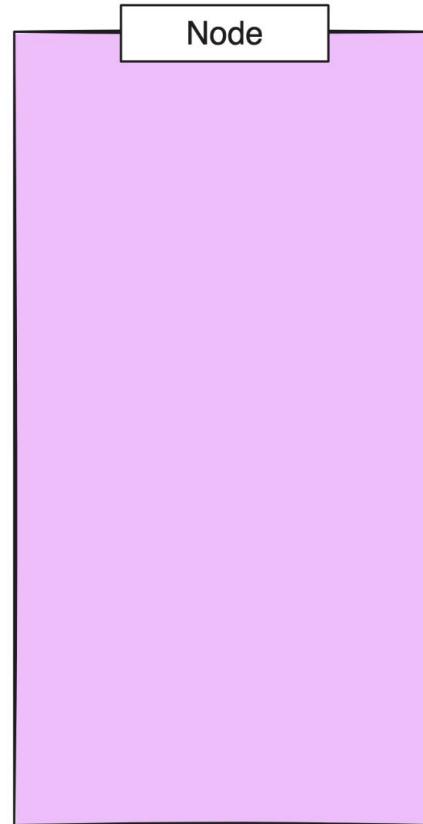
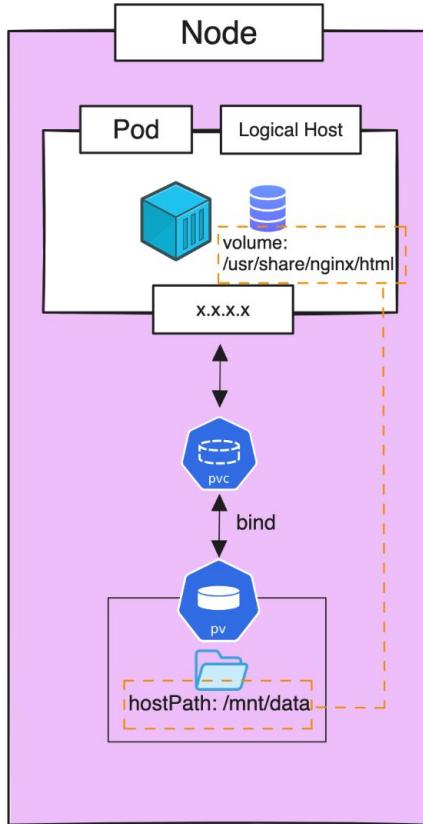
- Persistent Volume (PV) is a piece of storage in the cluster
- Can be provisioned from various sources like host local storage, cloud providers storage, or NFS.
- Persistent Volume Claims (PVC) allow a user to consume abstract storage resources
- A PVC to PV binding is a one-to-one mapping



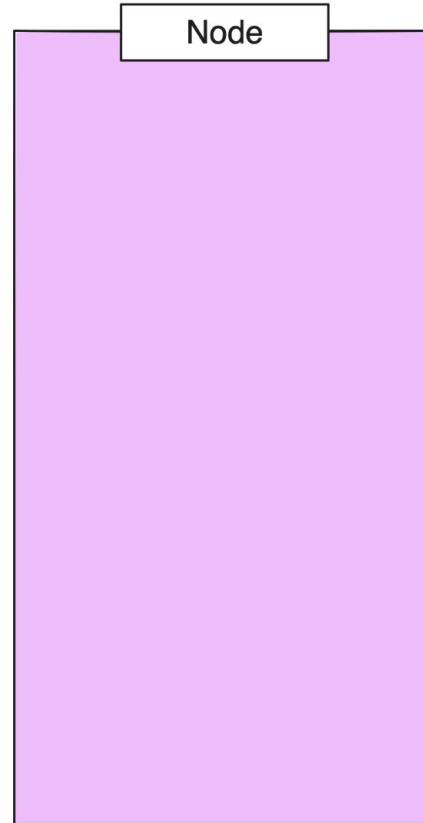
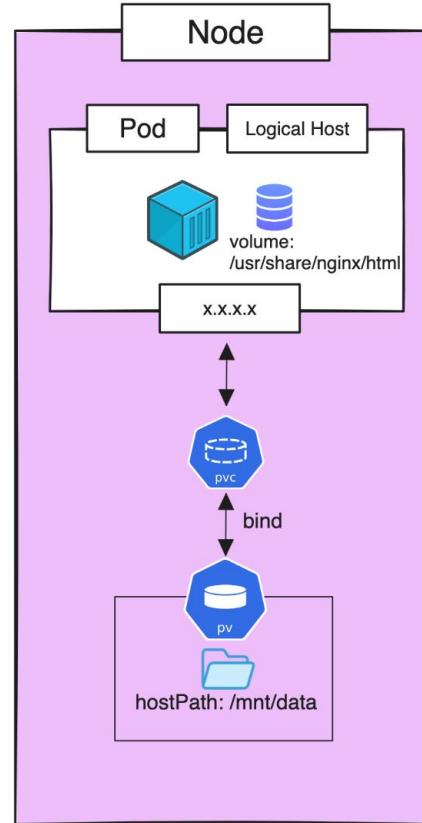
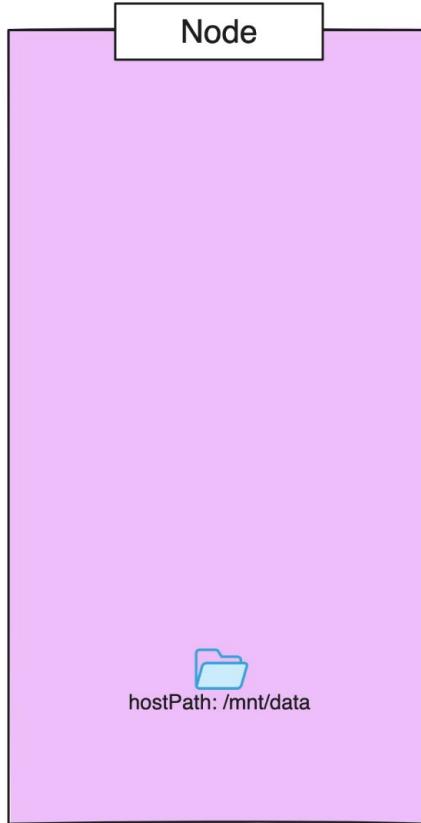
PV & PVC



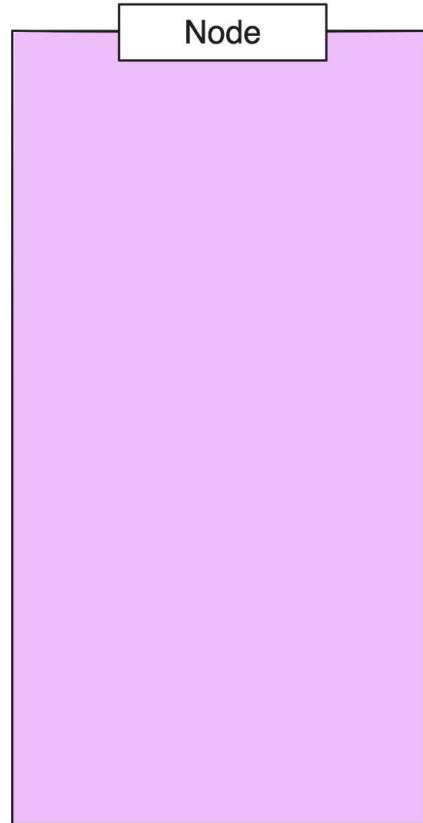
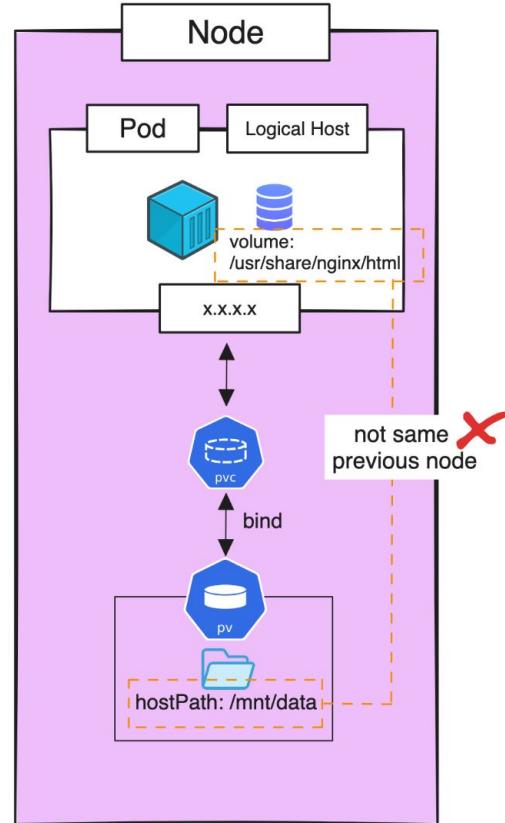
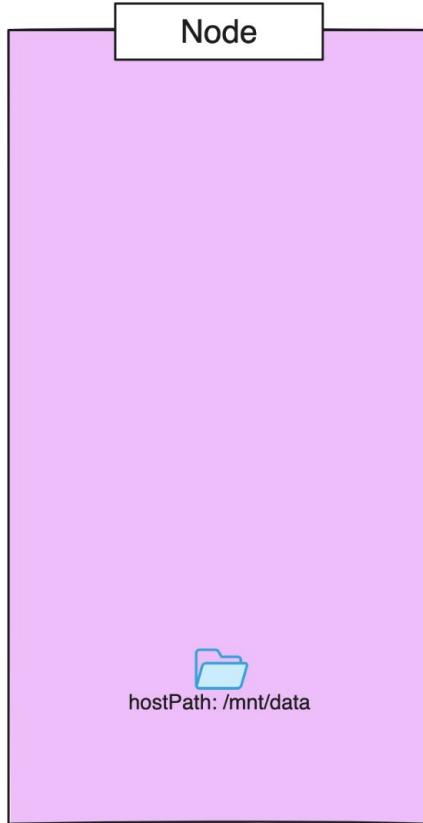
PV & PVC



PV & PVC



PV & PVC



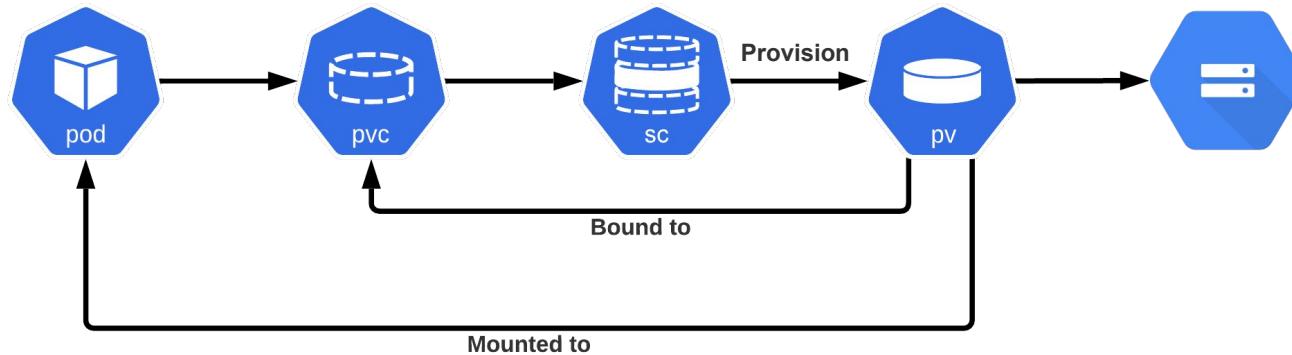
PV & PVC

Solutions

- Using selector node or nodeAffinity for ensure the target node is same
- Using shared storages like NFS or cloud storages

Storage Classes

- Template for Provisioning PersistentVolumes (PV)
- A StorageClass provides a way for administrators to describe the classes of storage
- Automatically provisions PV when defined PVC (Dynamic Volume Provisioning)



Workshop

Persistent Volume

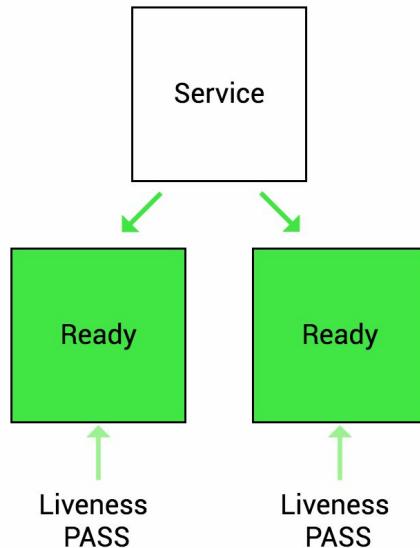
Day 3

Good to know



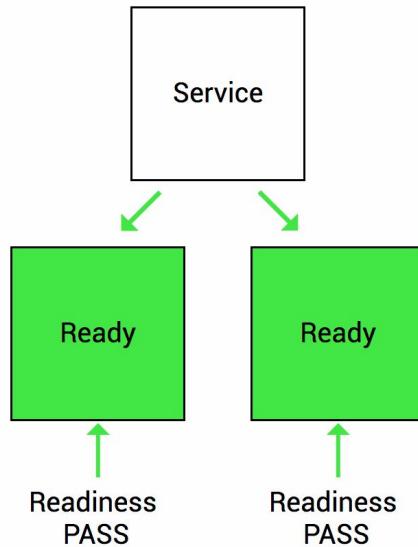
Pod liveness & readiness

- **Liveness:** Indicates whether the Container is running. and if for some reason the liveness probe fails, it restarts the container



Pod liveness & readiness

- **Readiness:** Indicates whether the Container is ready to service requests. If the condition inside readiness probe passes, only then our application can serve traffic.



Namespace

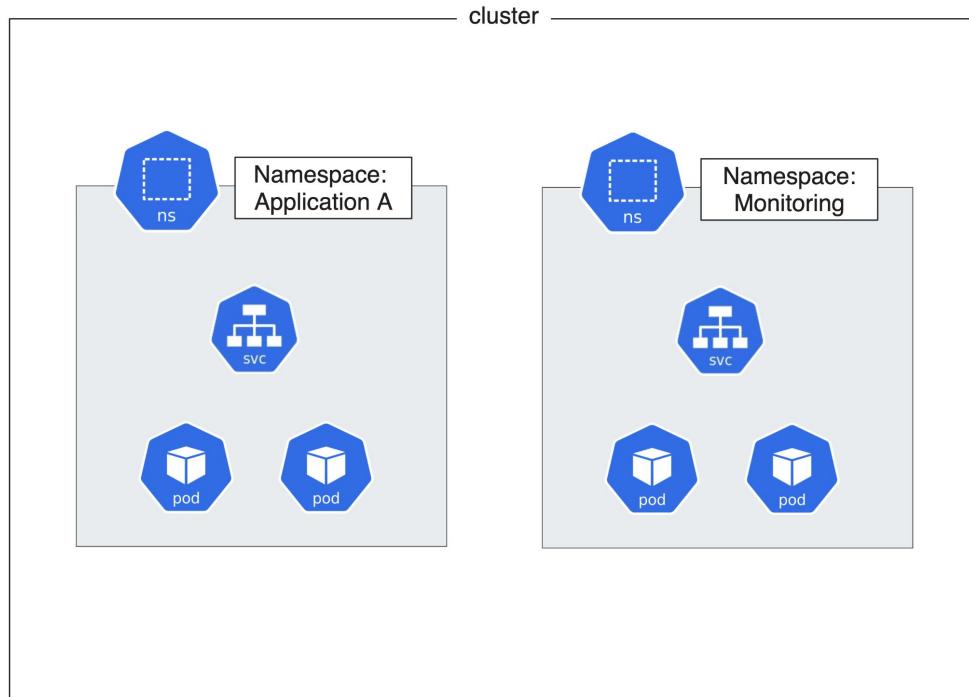


Namespace

- Partition a single kubernetes cluster into multiples clusters
- 4 defaults namespaces
 - **kube-system**: core system processes
 - **kube-public**: public accessible data, ex: ConfigMaps & Secrets
 - **kube-node-lease**: heartbeat of nodes, so that the control plane can detect node failure
 - **default**: your resources/objects

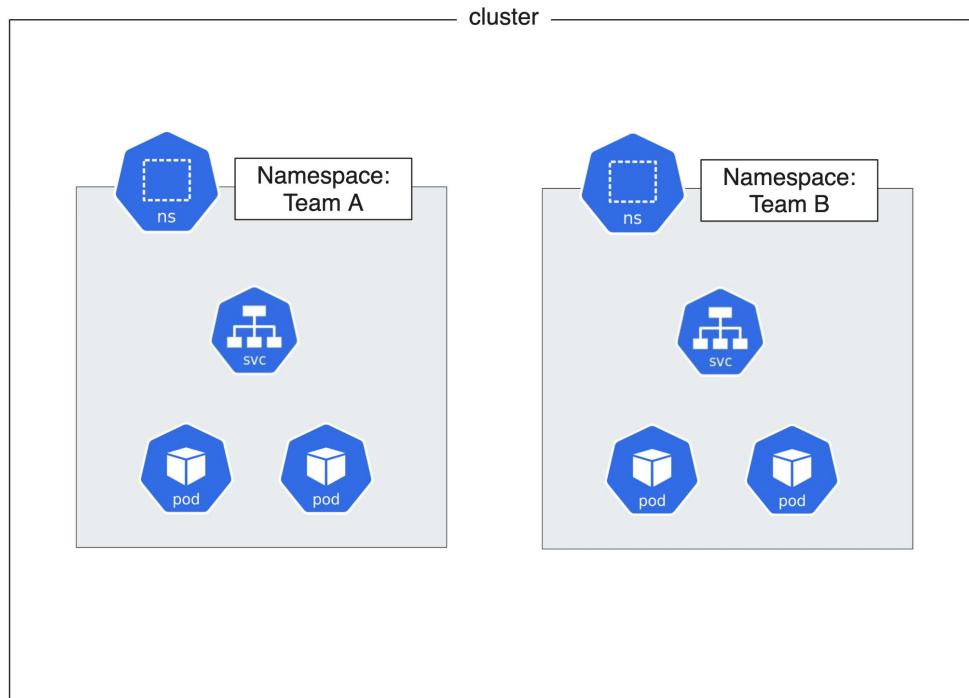
Namespace - use cases

- Application domain



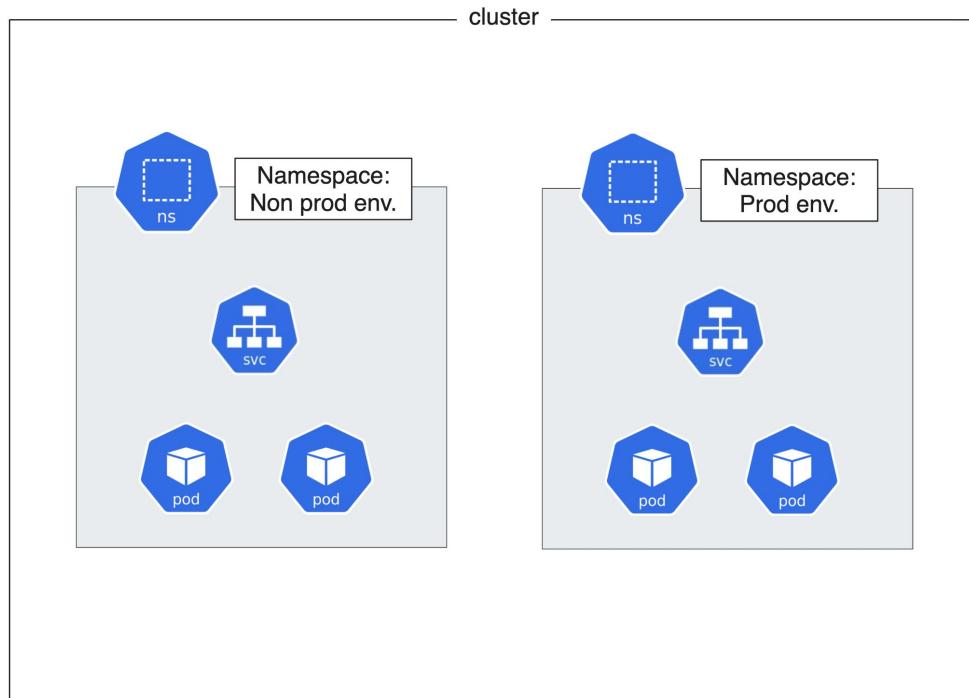
Namespace - use cases

- Teams driven



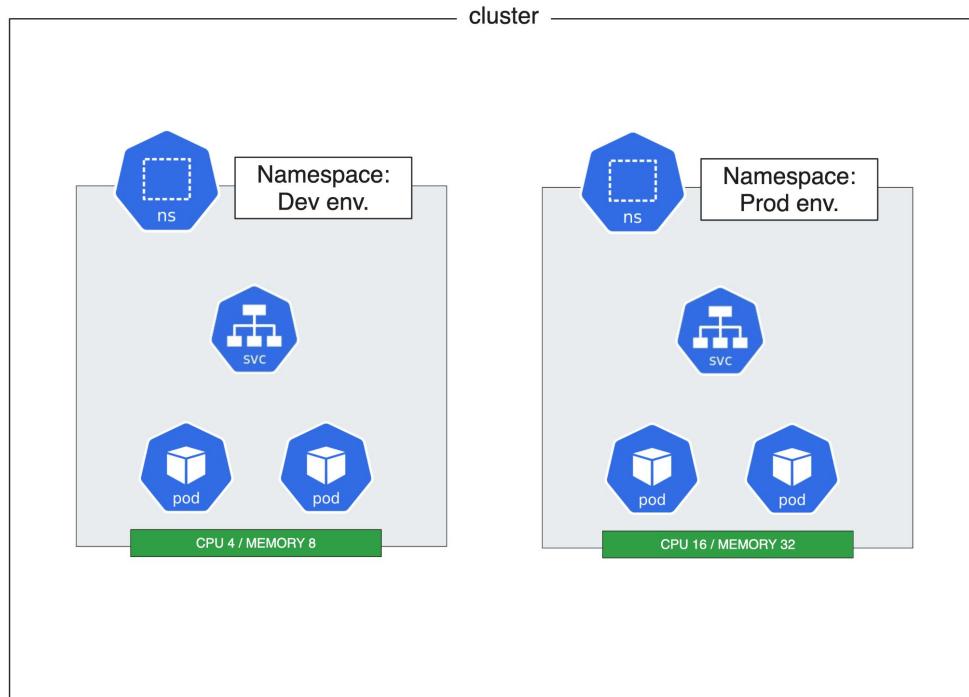
Namespace - use cases

- Environments



Namespace - use cases

- Resources quota limit



Workshop

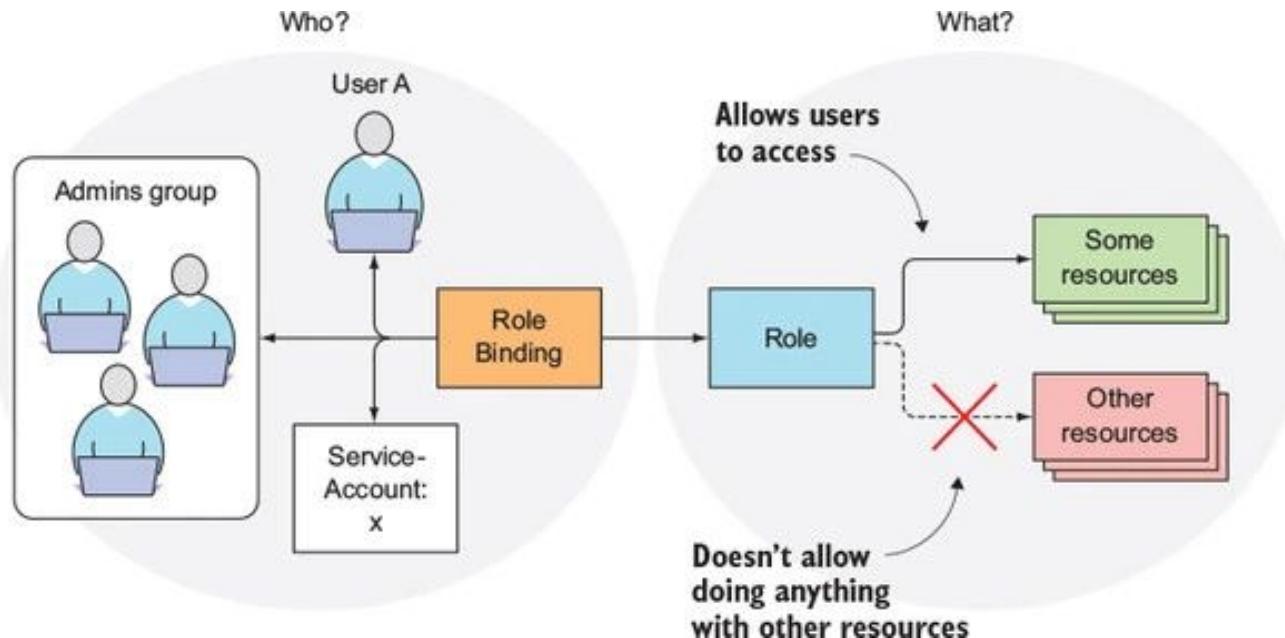
Namespace

RBAC (Role Based Access Control)



RBAC

- Who ?
- How ?
- What ?



RBAC

Who ? (Kubernetes Subject)

1. User
2. Group
3. Service account

RBAC

How ? (Kubernetes Verb)

- get
- list
- watch
- create
- update
- patch
- delete

<https://kubernetes.io/docs/reference/access-authn-authz/authorization/#determine-the-request-verb>

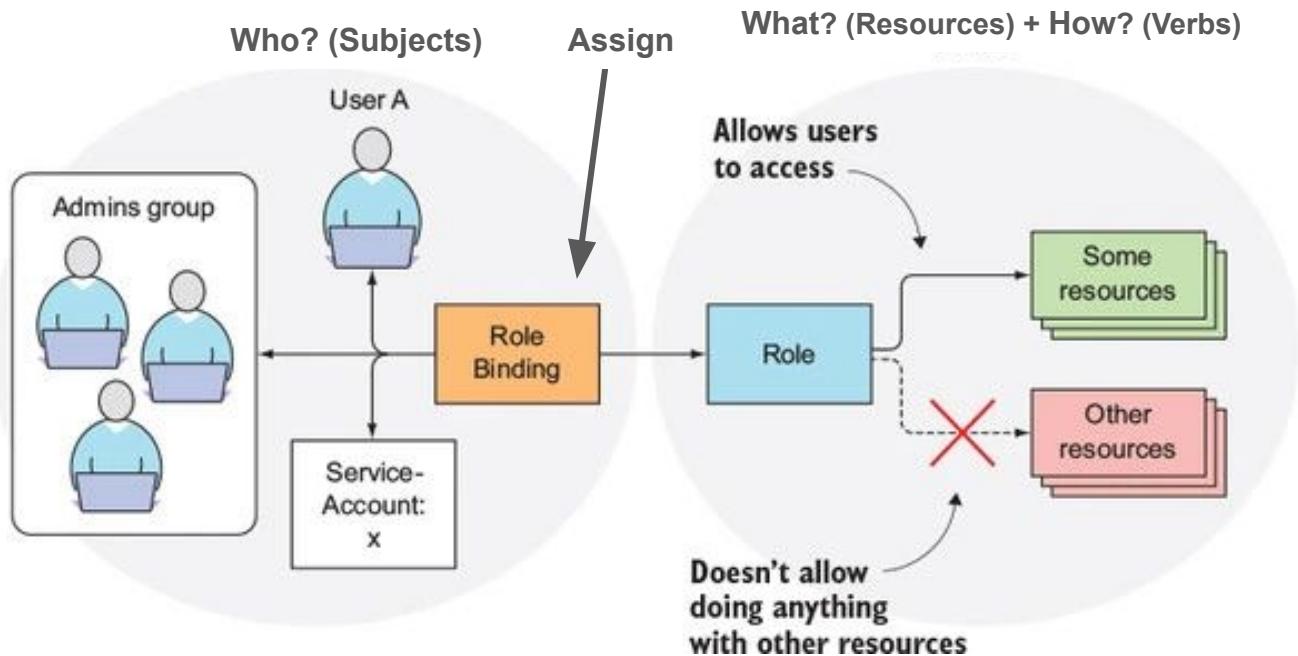
RBAC

What ? (Kubernetes Resource)

- pod
- deployment
- service
- etc.

<https://kubernetes.io/docs/reference/kubectl/#resource-types>

RBAC



RBAC

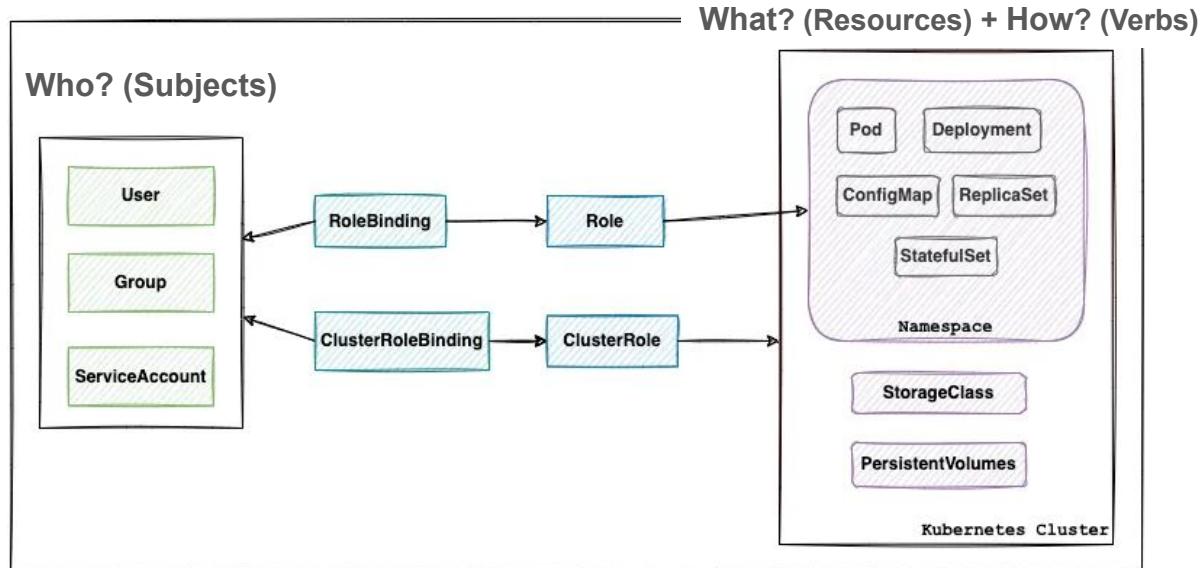
Permission Scopes

- Role
 - A Role can only be used to grant access to resources **within** a single namespace
- ClusterRole
 - A ClusterRole can be used to grant access to resources **across** all namespaces in the cluster

RBAC

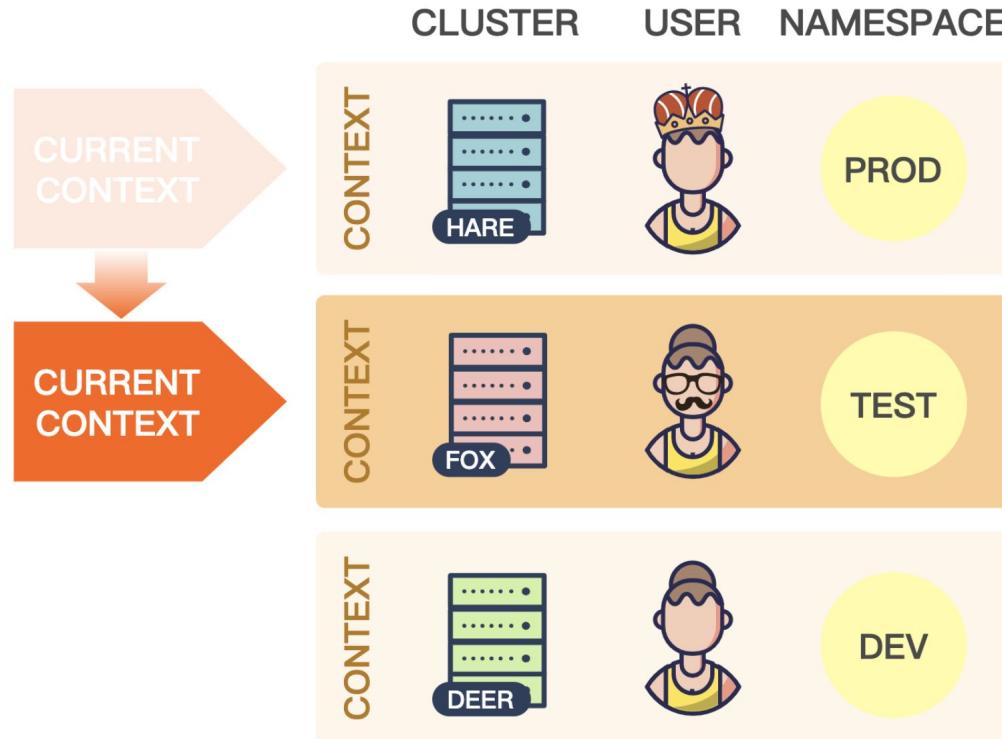
Binding

- RoleBinding: use with Role
- ClusterRoleBinding: use with ClusterRole



Context

- Access to Cluster(s)



Workshop

RBAC

HELM



HELM

HELM is the package manager for Kubernetes



HELM

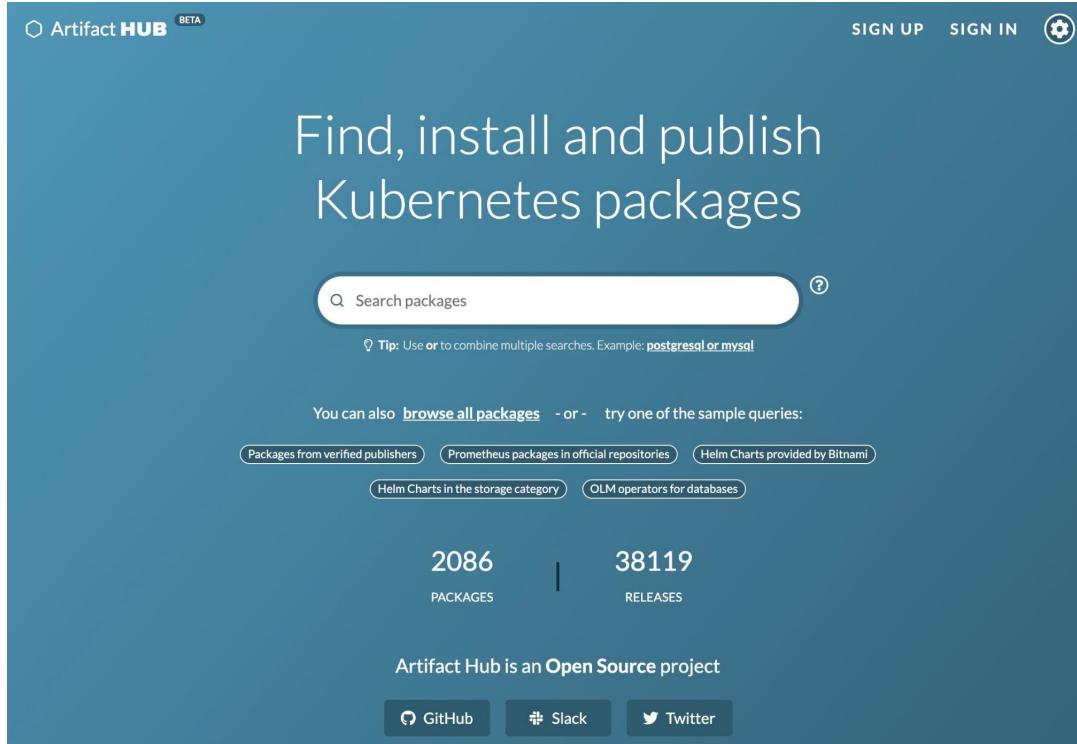
HELM is the package manager for Kubernetes

- Boosts productivity
- Reduces complexity
- Simple sharing
- Versioning and Upgrades

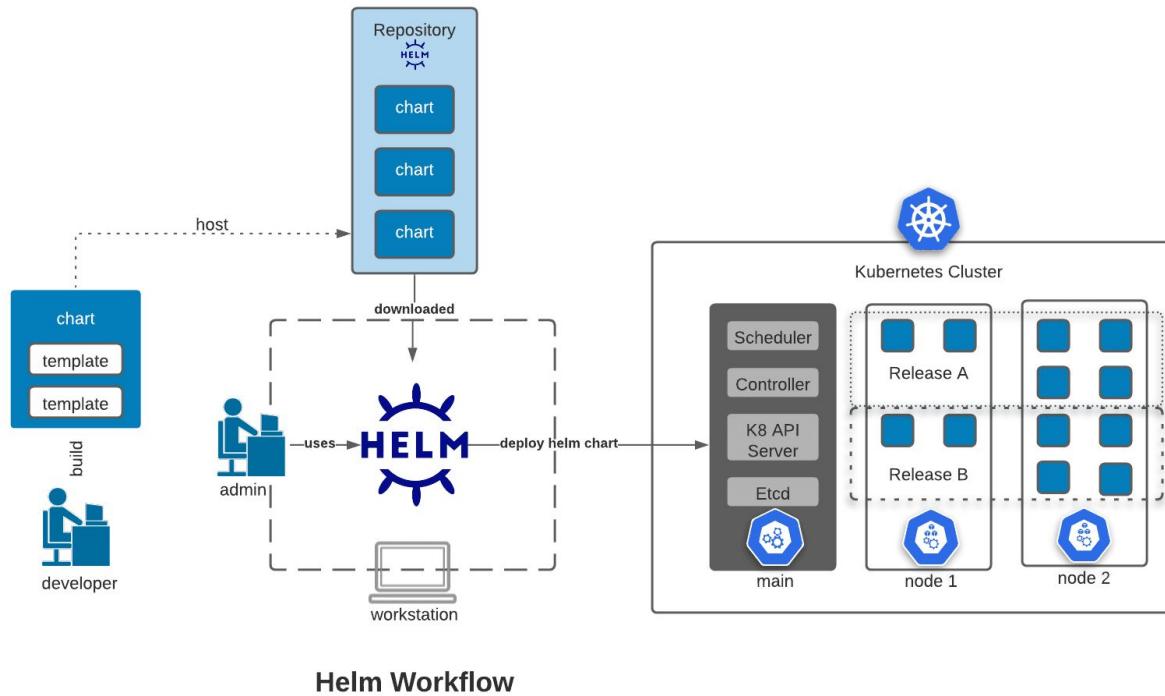


HELM Charts

- Public HELM Chart: <https://artifacthub.io/>



HELM Charts



Workshop

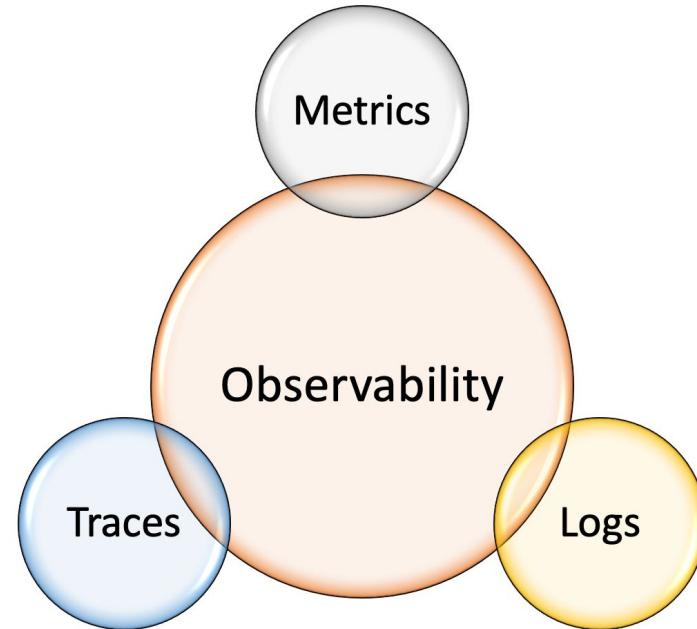
HELM

MONITORING



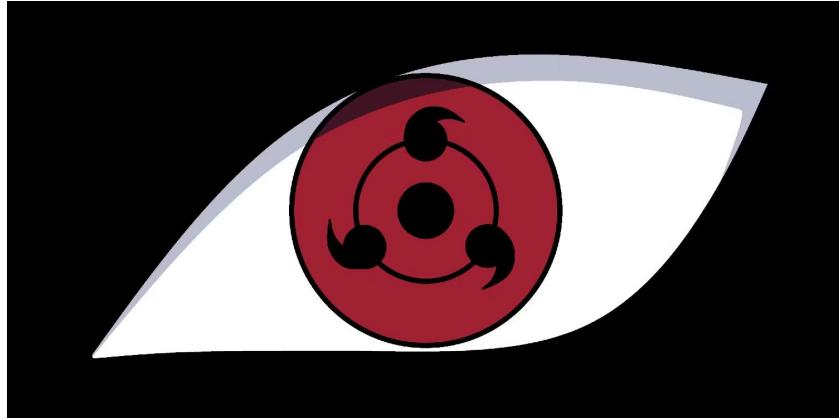
Three pillars of Observability

- Metrics
- Logs
- Traces

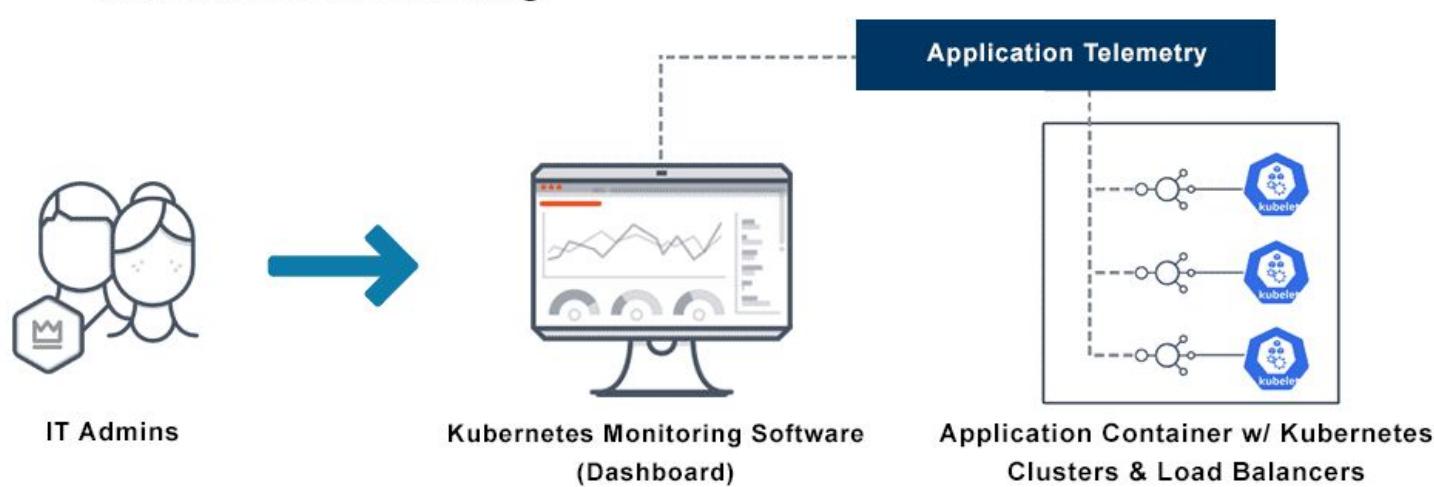


Three pillars of Observability

- Metrics
- Logs
- Traces



Metric Monitoring



Metric Monitoring

- API request latency, the lower the better
- Cluster state metrics, including the availability and health of pods
- CPU utilization
- Disk utilization
- Node status, including disk or processor overload, memory, network availability, and readiness
- Pod availability (unavailable pods may indicate poorly designed readiness probes or configuration issues)

Kubernetes Dashboard

The screenshot shows the Kubernetes Dashboard interface. At the top, there is a navigation bar with a logo, a dropdown menu set to "default", a search bar, and several icons for creating, deleting, and filtering.

The main area has a blue header bar labeled "Workloads". Below it, there are three large green circles representing different workload types:

- Deployments:** Running: 2
- Pods:** Running: 3
- Replica Sets:** Running: 2

On the left side, there is a sidebar with the following sections and their sub-items:

- Workloads**: Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets.
- Service**: Ingresses, Ingress Classes, Services.
- Config and Storage**: Config Maps, Persistent Volume Claims.
- Secrets**: Secrets, Storage Classes.
- Cluster**: Cluster Role Bindings, Cluster Roles.
- Events**: Events.
- Namespaces**: Namespaces.

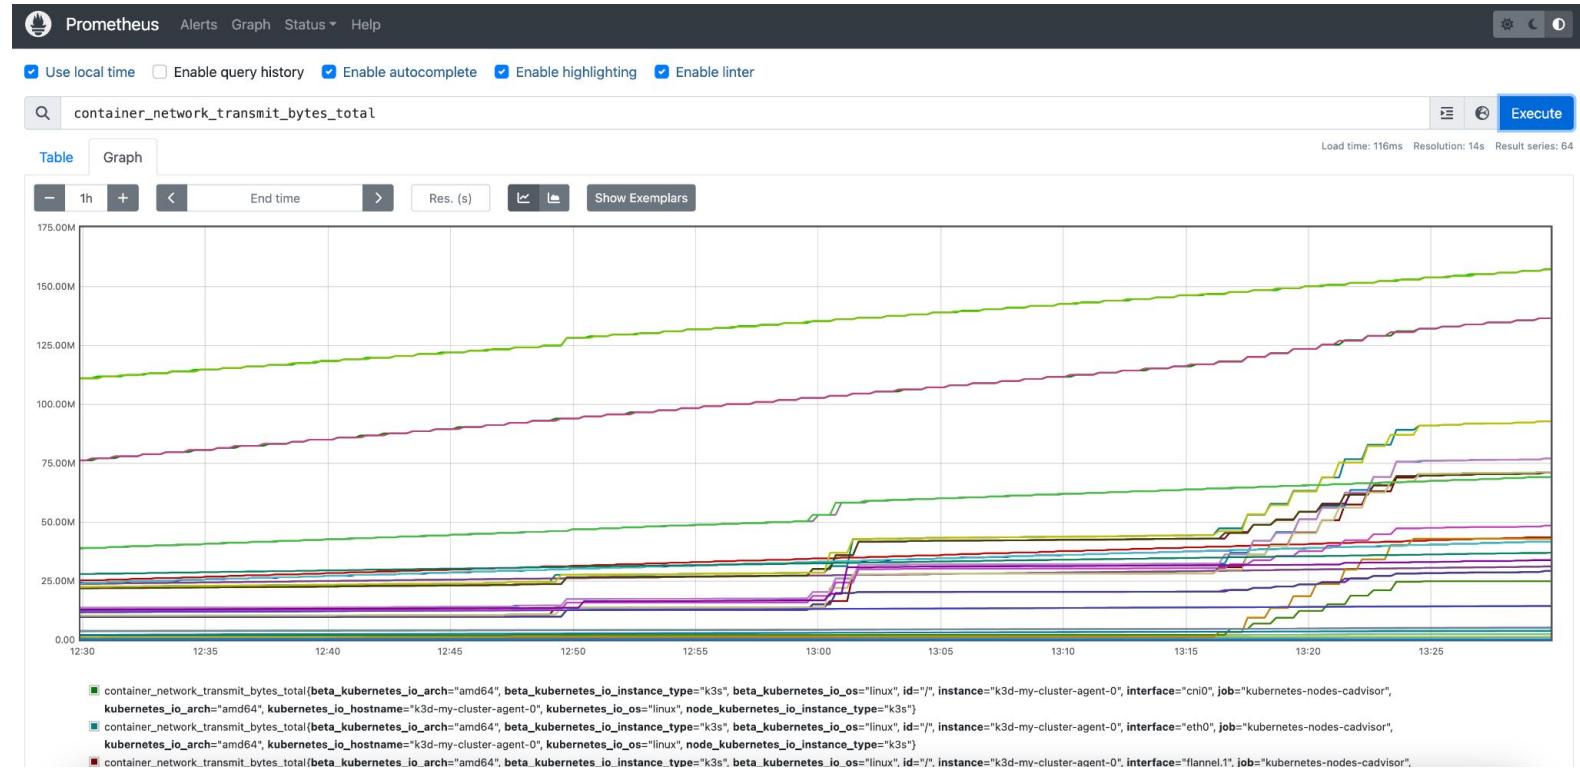
The "Deployments" section contains two entries:

| Name | Images | Labels | Pods | Created |
|----------------------|------------------------|--------|-------|-------------|
| mysql-deployment | mysql:8.0 | - | 1 / 1 | 3 hours ago |
| wordpress-deployment | wordpress:6.5.0-php8.1 | - | 2 / 2 | 3 hours ago |

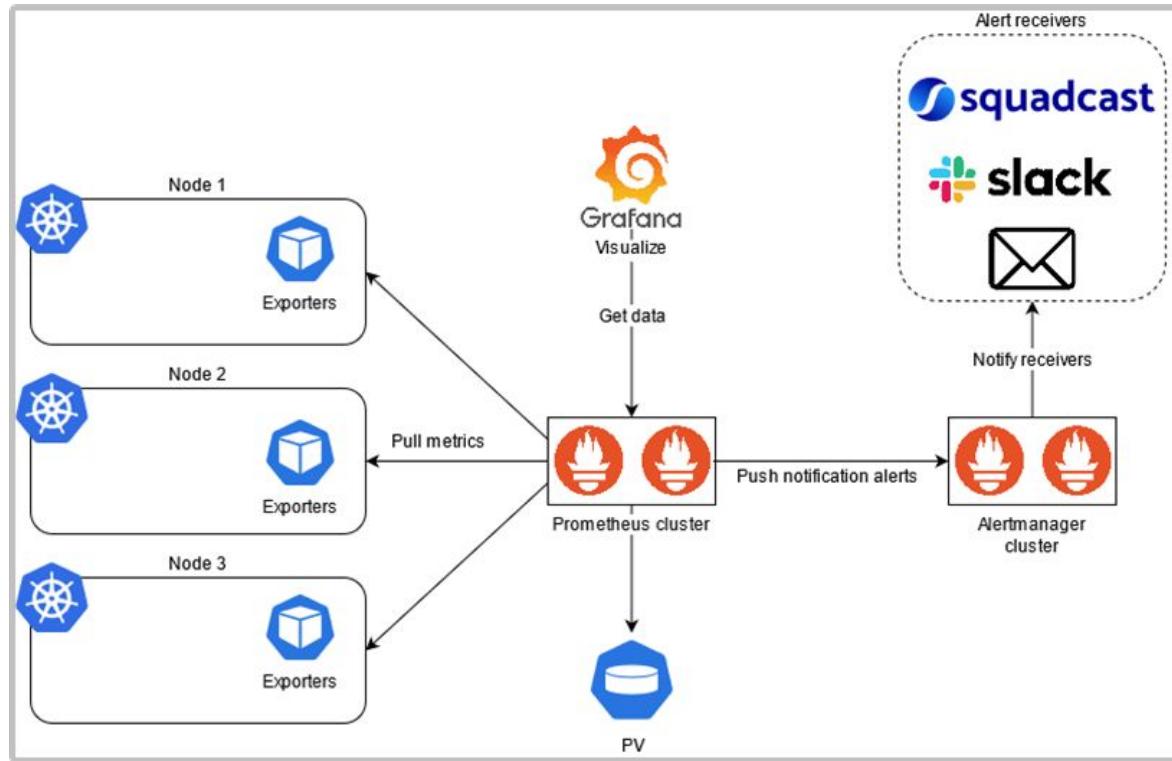
The "Pods" section contains three entries:

| Name | Images | Labels | Node | Status | Restarts | CPU Usage (cores) | Memory Usage (bytes) | Created |
|---------------------------------------|------------------------|--|------------------------|---------|----------|---|---|--------------|
| wordpress-deployment-5cd7bf5cff-7xmnc | wordpress:6.5.0-php8.1 | app: my-wordpress
pod-template-hash: 5cd7bf5cff | k3d-my-cluster-agent-0 | Running | 0 | <div style="width: 1.00m">1.00m</div> | <div style="width: 10.89Mi">10.89Mi</div> | a minute ago |
| wordpress-deployment-5cd7bf5cff-Szpc5 | wordpress:6.5.0-php8.1 | app: my-wordpress
pod-template-hash: 5cd7bf5cff | k3d-my-cluster-agent-1 | Running | 0 | <div style="width: 1.00m">1.00m</div> | <div style="width: 107.11Mi">107.11Mi</div> | 3 hours ago |
| mysql-deployment-57d5f9946c-jq4rb | mysql:8.0 | app: my-mysql
pod-template-hash: 57d5f9946c | k3d-my-cluster-agent-1 | Running | 0 | <div style="width: 18.00m">18.00m</div> | <div style="width: 370.61Mi">370.61Mi</div> | 3 hours ago |

Prometheus



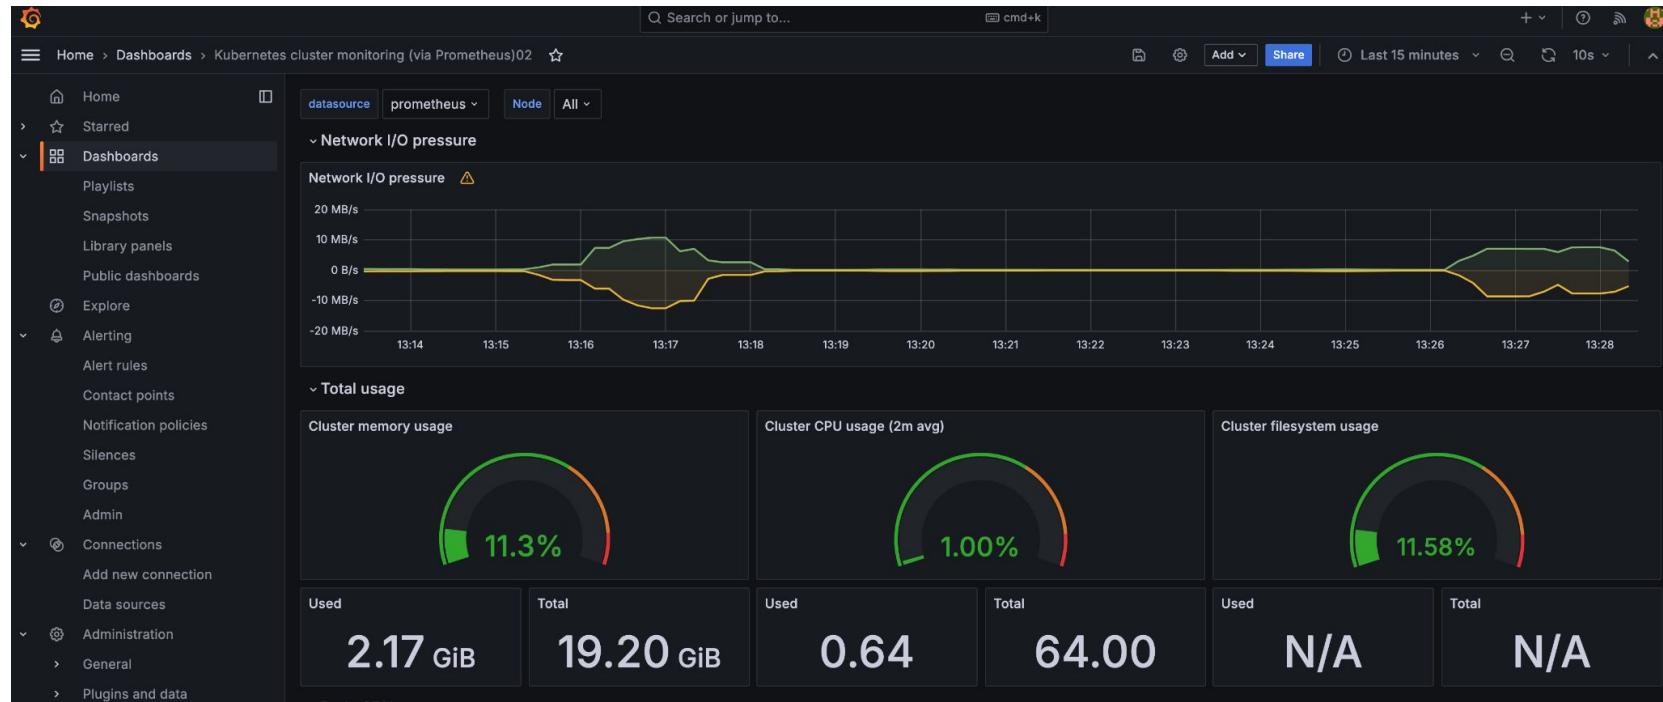
Prometheus Alert



Prometheus Alert Rule Examples

- <https://samber.github.io/awesome-prometheus-alerts/rules>

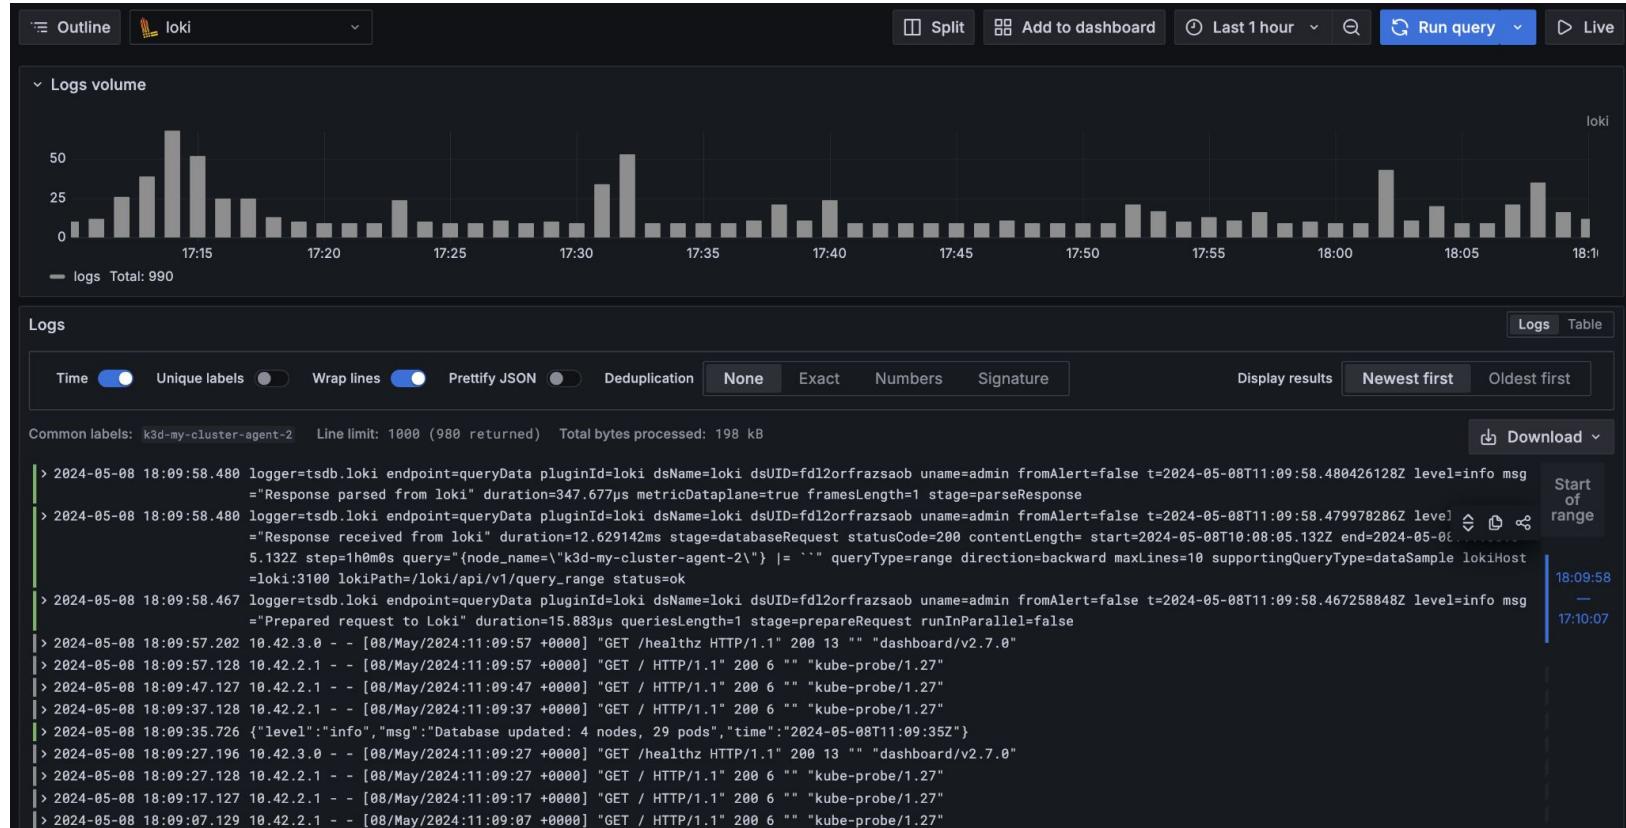
Grafana



Logs Monitoring

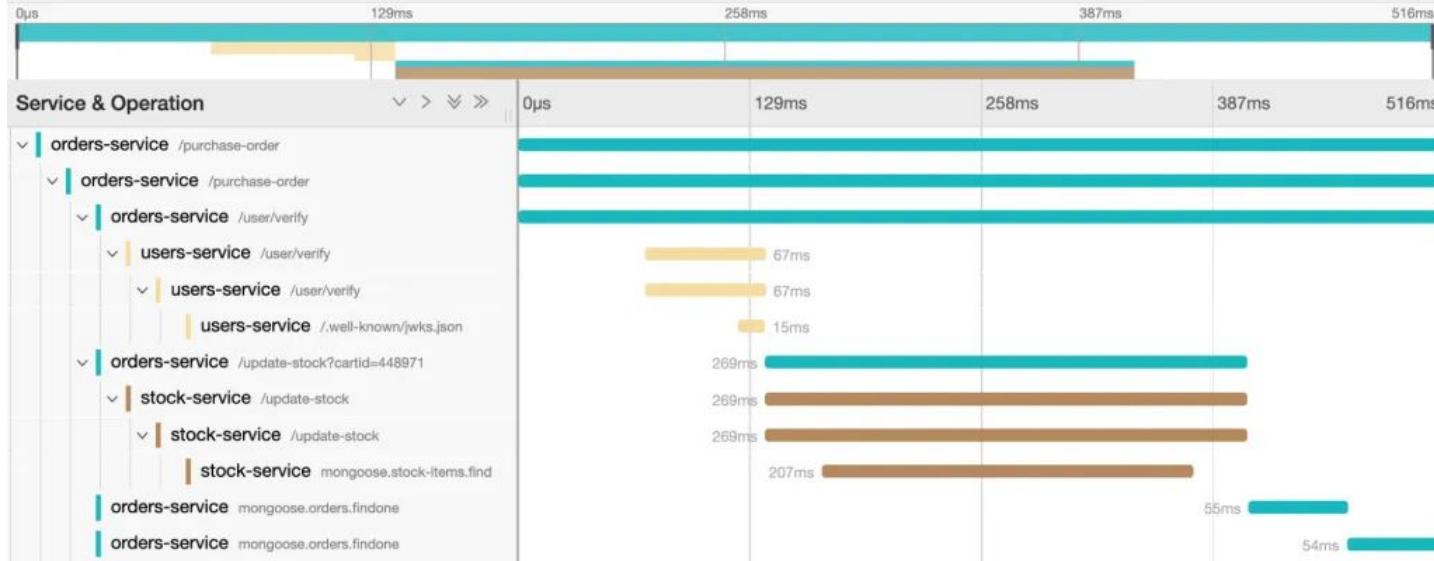
- Cluster logs
- Node logs
- Container logs
- Kubernetes Events

Loki

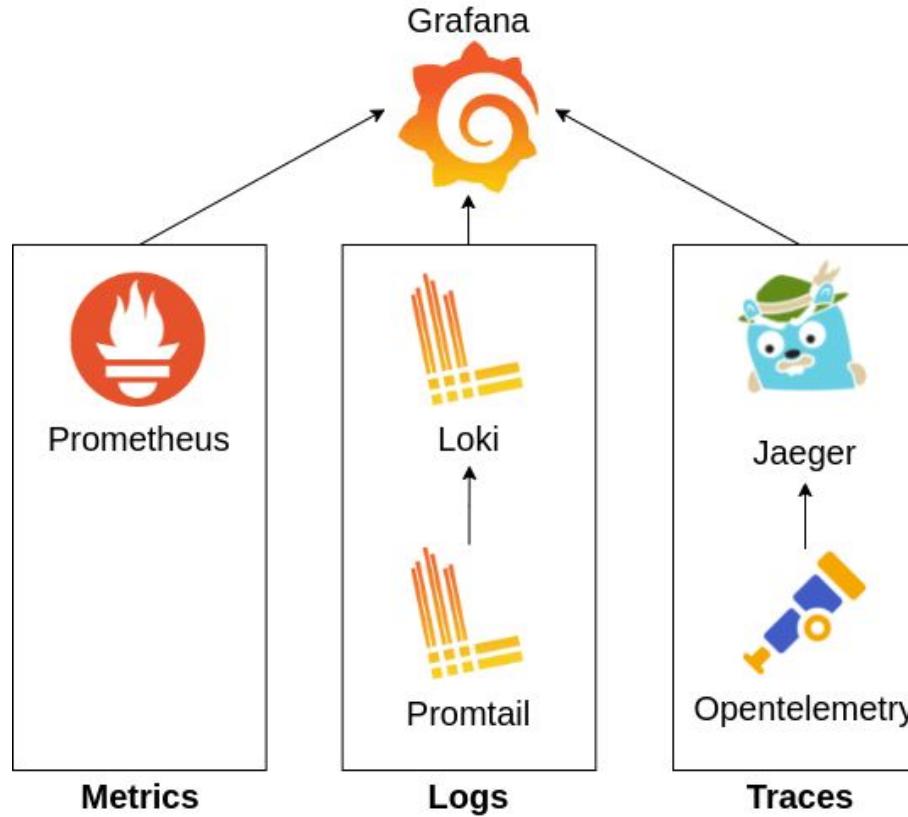


Tracing Monitoring

- A trace represents the entire journey of a request or action as it moves through all the nodes of a distributed system



Monitoring



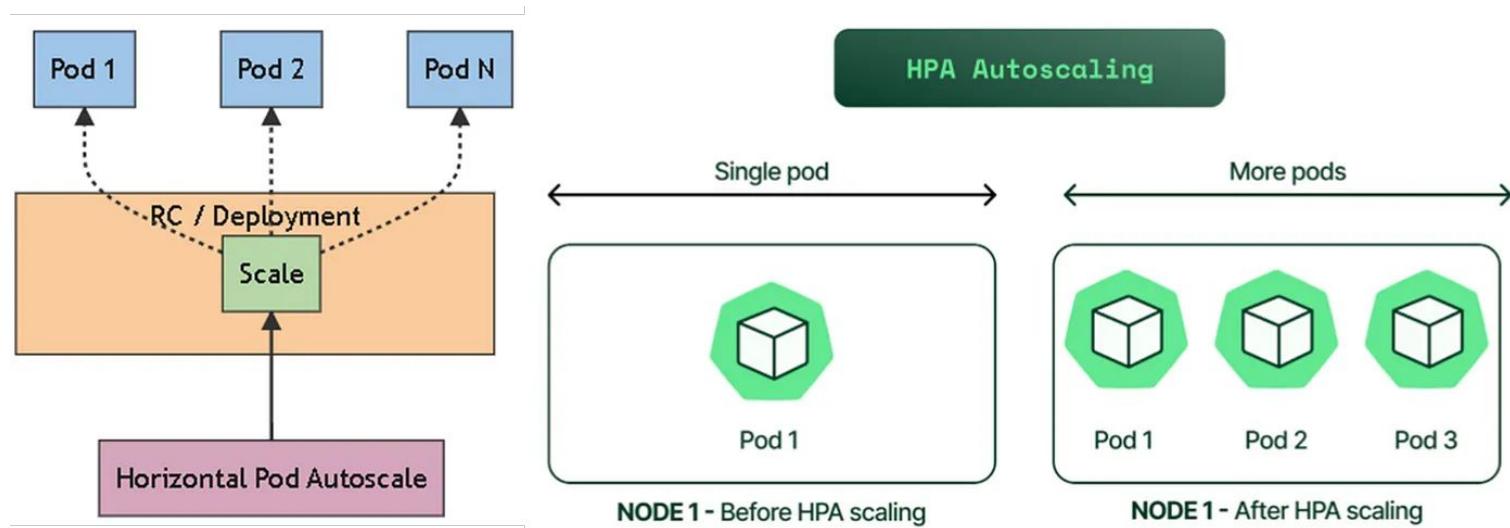
Workshop Monitoring

HPA (Horizontal Pod Autoscaler)

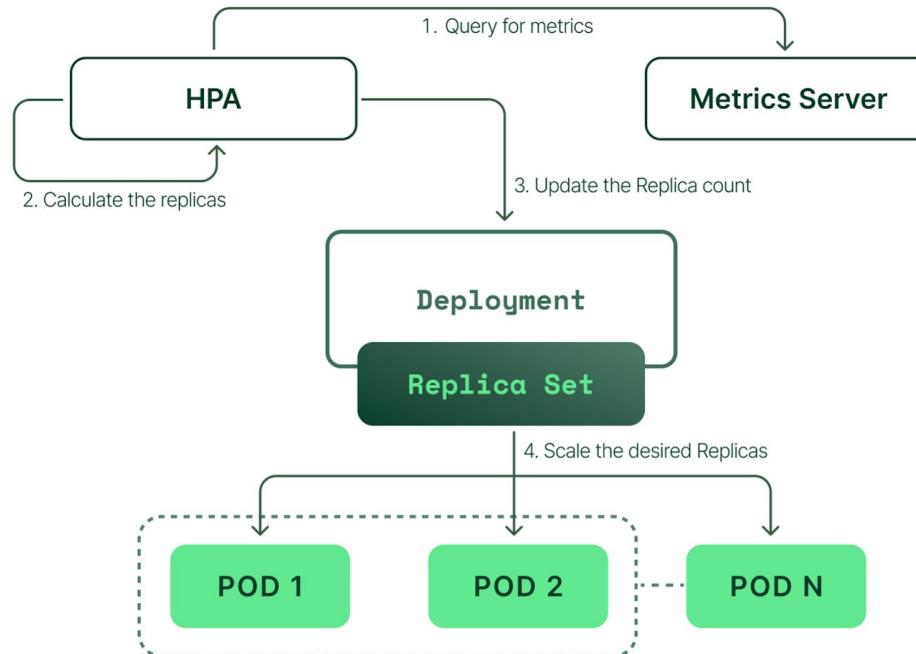


HPA

- A HorizontalPodAutoscaler (HPA for short) automatically updates a workload resource, with the aim of automatically scaling the workload to match demand



How does HPA work?



Workshop

HPA

Deployment & Pipelines



Deployment strategies

Deployment strategies

- Recreate deployment
- Ramped (Rolling deployment)
- Blue/green deployment
- Canary deployment
- A/B testing
- Shadow deployment

Recreate deployment

- Terminating all the running instances then recreate them with the newer version

Advantage

- Application state entirely renewed

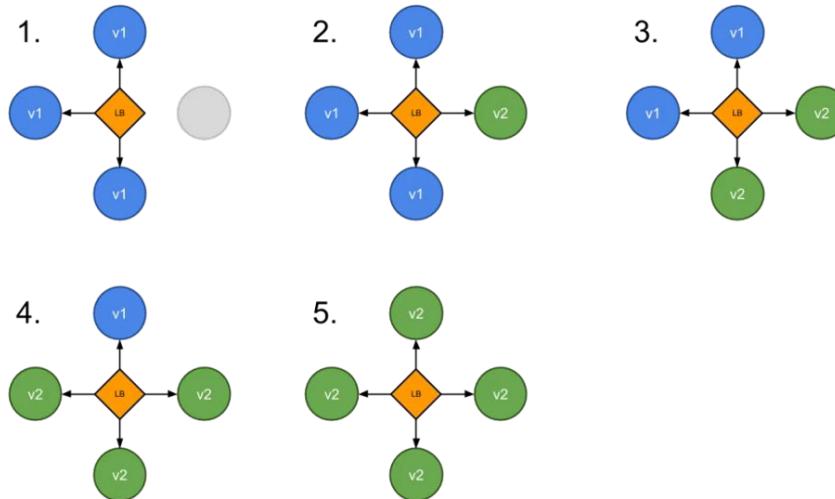
Disadvantage

- Downtime that depends on both shutdown and boot duration of the application



Ramped (Rolling deployment)

- The rolling update is the **default** deployment strategy in Kubernetes that update the pods one by one without causing cluster downtime. the deployment objects create new pods using the new pod specifications and shut down the older ones when the application starts running successfully



Ramped (Rolling deployment)

Advantages

- Enables zero-downtime deployments
- Well-suited for large deployments and stateful applications

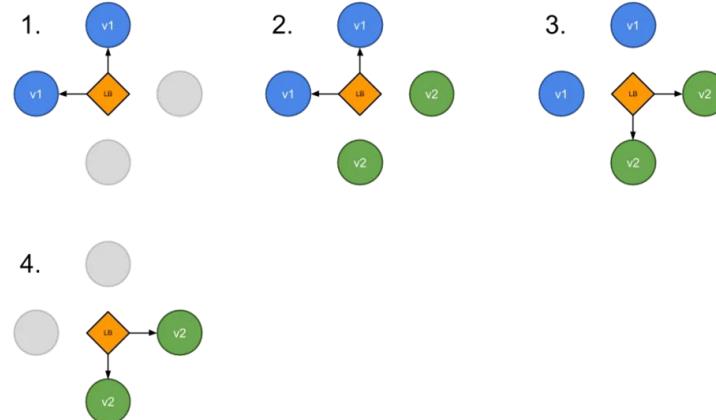
Disadvantages

- Rollout/rollback can take time
- Supporting multiple APIs is hard
- No control over traffic



Blue/green deployment

- While the “Blue” version is running, engineers deploy the “green” version in another deployment and run basic tests to ensure that the pods within the deployments are stable, then using services, you switch to the newer version of the application
- Best to avoid API versioning issues



Blue/green deployment

Advantages

- Instant rollout/rollback
- Clear separation between old and new versions

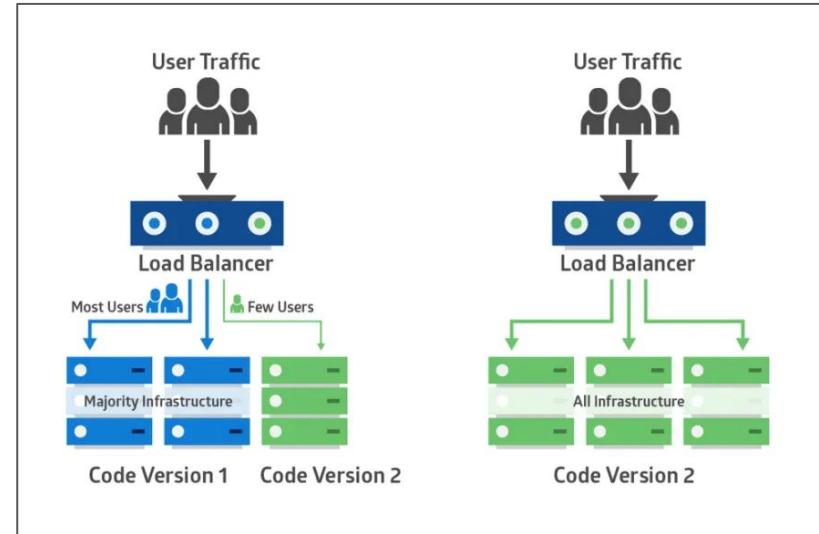
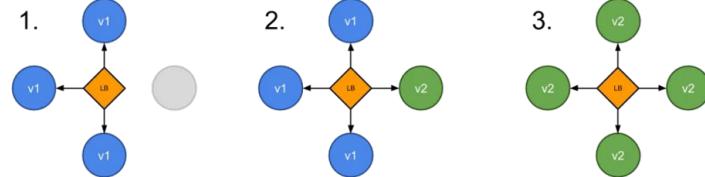
Disadvantages

- Requires double the resources
- More complex configuration



Canary deployment

- Similar to blue/green deployments, but only a **small subset** of traffic is directed to the new version (canary) initially. This allows for early detection of issues before a full rollout



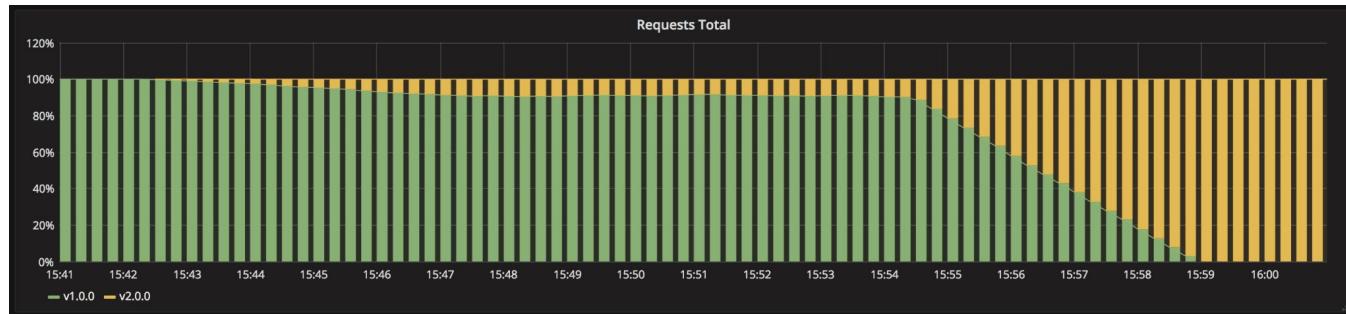
Canary deployment

Advantages

- Lowers risk by gradually rolling out new version to a small subset of users
- Enables early detection of issues

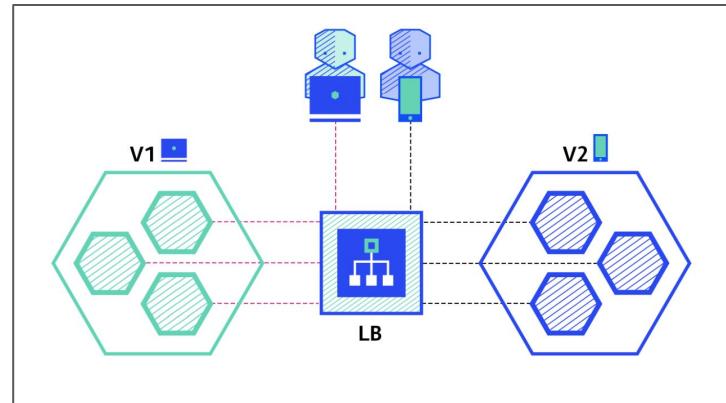
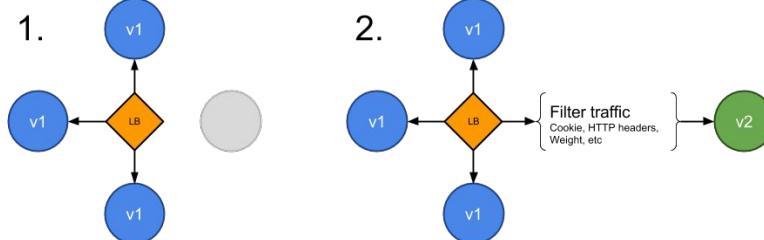
Disadvantages

- Requires additional configuration for traffic routing
- Slow rollout, most of users cannot use new versions
- Application must support both of old & new versions



A/B testing

- A/B testing deployments consists of routing a subset of users to a new functionality under **specific conditions**. It is usually a technique for making business decisions based on statistics, rather than a deployment strategy.
- Best for feature testing on a subset of users



A/B testing

Advantages

- Several versions run in parallel
- Full control over the traffic distribution

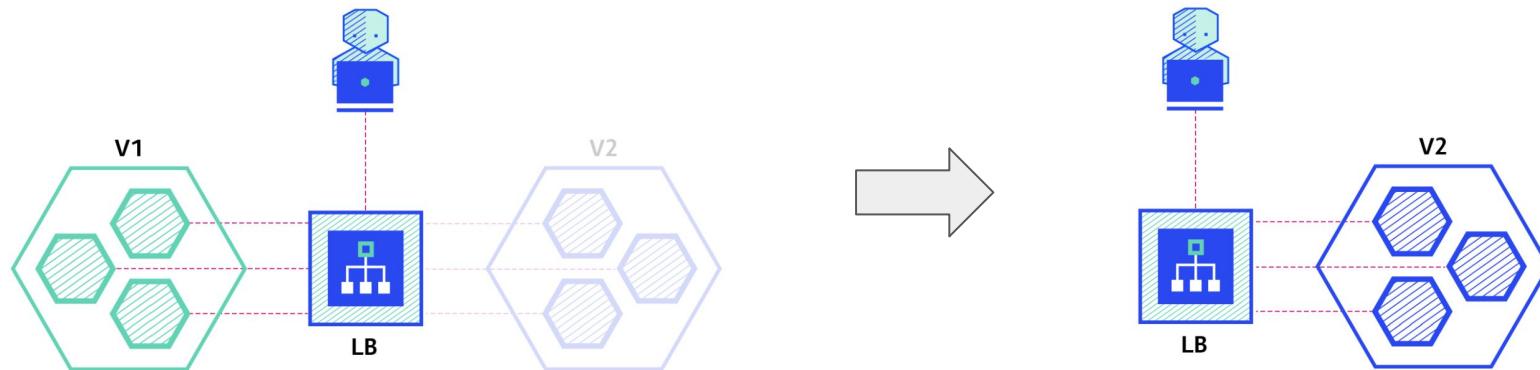
Disadvantages

- Requires intelligent load balancer
- Not straightforward, you need to setup additional tools
- Hard to troubleshoot errors for a given session



Shadow deployment

- Consists of releasing version B alongside version A, fork version A's incoming requests and send them to version B as well without impacting production traffic. This is particularly useful to test production load on a new feature



Shadow deployment

Advantages

- Performance testing, auto scale of the application with production traffic,
- No impact on the user

Disadvantages

- Expensive as it requires double the resources.
- Complex to setup



DEPLOYMENT STRATEGIES

When it comes to production, a ramped or blue/green deployment is usually a good fit, but proper testing of the new platform is necessary.

Blue/green and shadow strategies have more impact on the budget as it requires double resource capacity. If the application lacks in tests or if there is little confidence about the impact/stability of the software, then a canary, a/b testing or shadow release can be used.

If your business requires testing of a new feature amongst a specific pool of users that can be filtered depending on some parameters like geolocation, language, operating system or browser features, then you may want to use the a/b testing technique.



| Strategy | ZERO DOWNTIME | REAL TRAFFIC TESTING | TARGETED USERS | CLOUD COST | ROLLBACK DURATION | NEGATIVE IMPACT ON USER | COMPLEXITY OF SETUP |
|--|---------------|----------------------|----------------|------------|-------------------|-------------------------|---------------------|
| RECREATE
version A is terminated then version B is rolled out | ✗ | ✗ | ✗ | ■□□ | ■■■ | ■■■ | □□□ |
| RAMPED
version B is slowly rolled out and replacing version A | ✓ | ✗ | ✗ | ■□□ | ■■■ | ■□□ | □□□ |
| BLUE/GREEN
version B is released alongside version A, then the traffic is switched to version B | ✓ | ✗ | ✗ | ■■■ | □□□ | ■■■ | □□□ |
| CANARY
version B is released to a subset of users, then proceed to a full rollout | ✓ | ✓ | ✗ | ■□□ | ■□□ | ■□□ | □□□ |
| A/B TESTING
version B is released to a subset of users under specific condition | ✓ | ✓ | ✓ | ■□□ | ■□□ | ■□□ | ■■■ |
| SHADOW
version B receives real world traffic alongside version A and doesn't impact the response | ✓ | ✓ | ✗ | ■■■ | □□□ | □□□ | ■■■ |

CI / CD Pipelines

What is a CI/CD pipelines ?

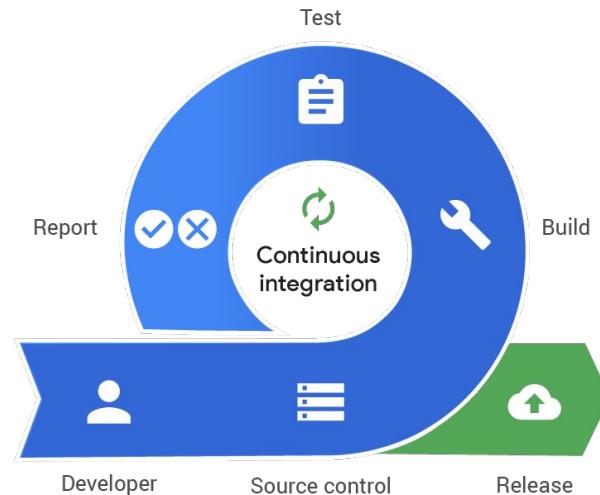
- Continuous Integration and Continuous Deployment (CI/CD) pipeline represents an **automatic workflow** that continuously integrates code developed by software developers and deploys them into a target environment with less human intervention.



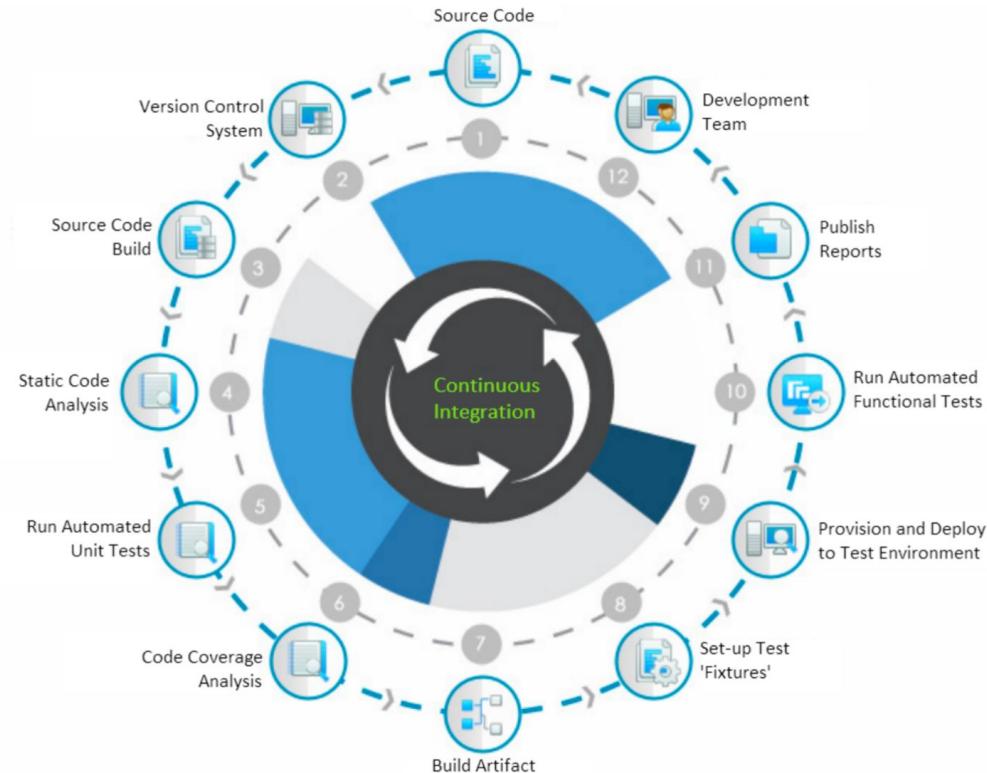
Continuous Integration or CI pipeline

CI is part of the development process. The central idea in this stage is to automatically build the software whenever new software is **committed** by developers

- Focuses on **frequently** merging code changes into a central repository
- Each code change, trigger automated builds and tests
- Early detection
- This ensures that the code compiles correctly and passes all the essential tests



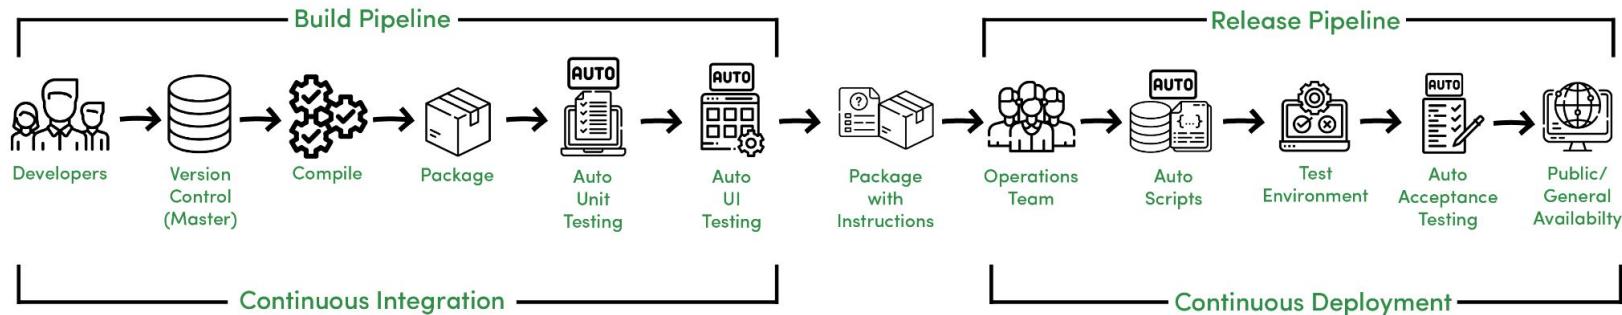
Continuous Integration or CI pipeline



Continuous Deployment or CD pipeline

Continuous deployment stage refers to pulling the artifact from the repository and deploying it frequently and safely. Automate the deployment into the each environments with less human intervention

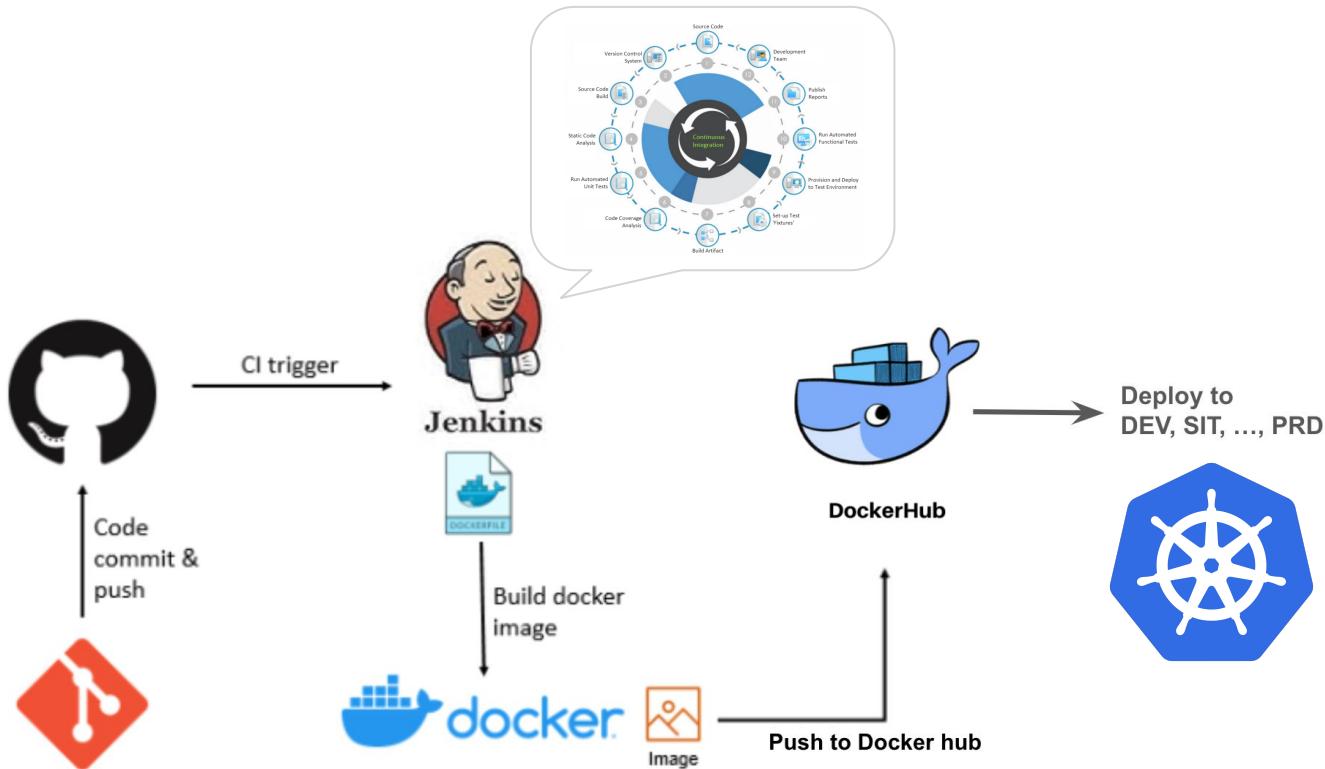
- Once the code passes CI checks, CD takes over
- Some approaches call **Continuous Delivery**: Makes the built code readily available for deployment, allowing for manual review before pushing it live



Benefits of CI/CD Pipelines

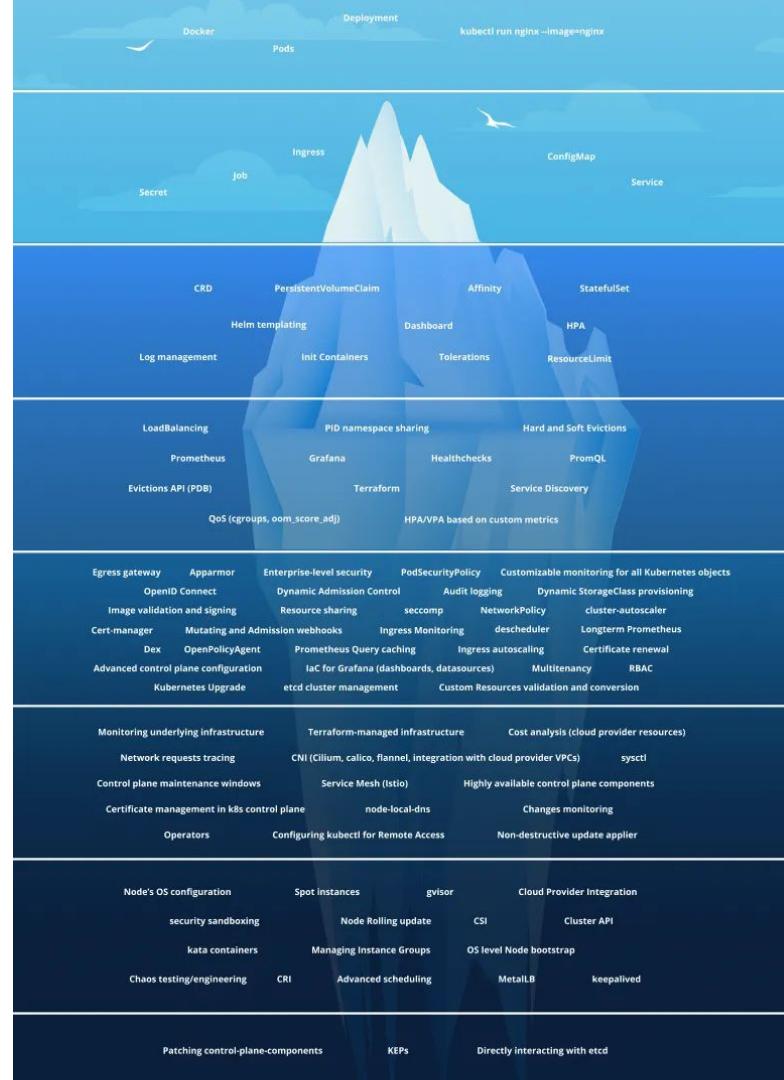
- Faster Software Delivery
- Improved Software Quality
- Reduced Risk
- Increased Collaboration
- Reliable & cheaper deployments

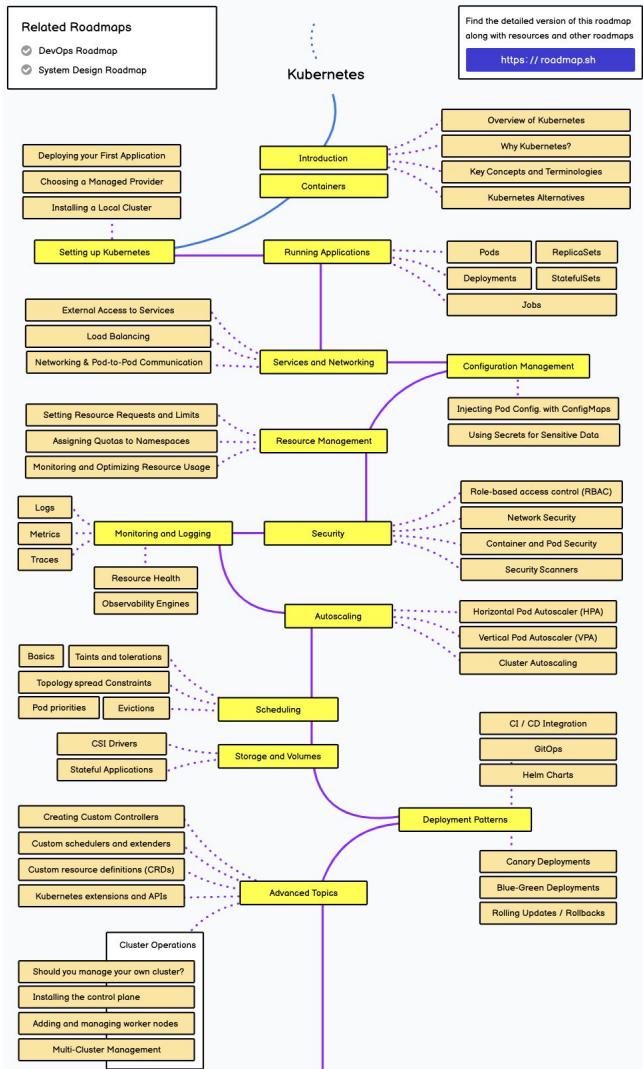
Apply Docker & k8s with pipelines



Workshop

CI / CD Pipelines



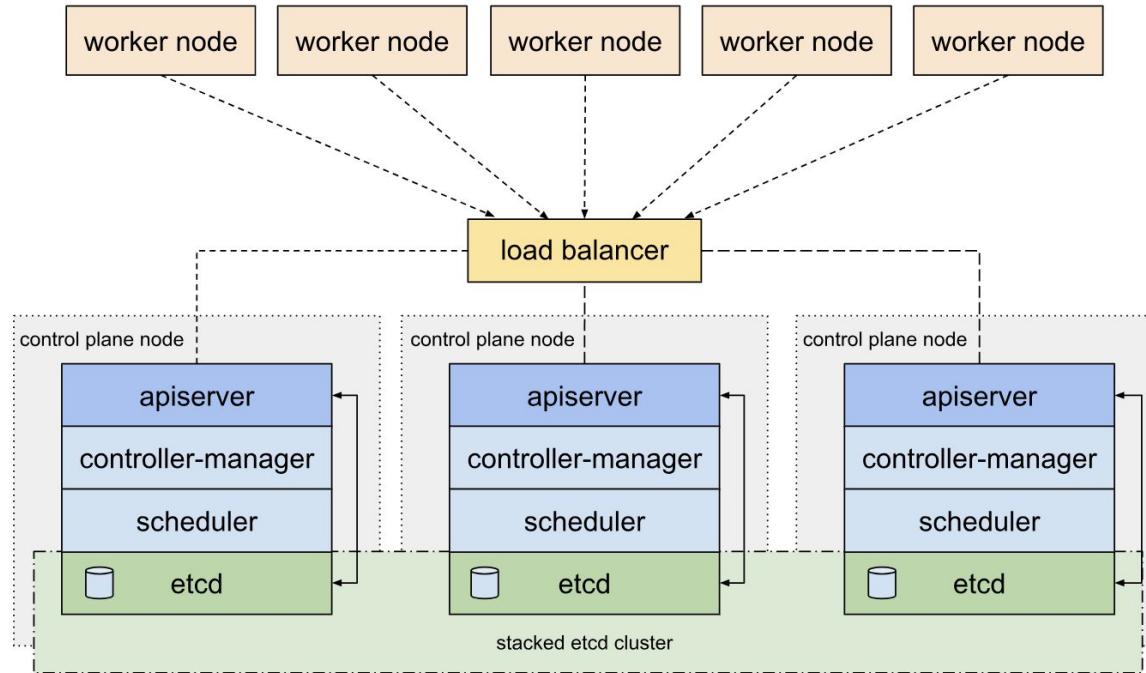


Cluster Setup

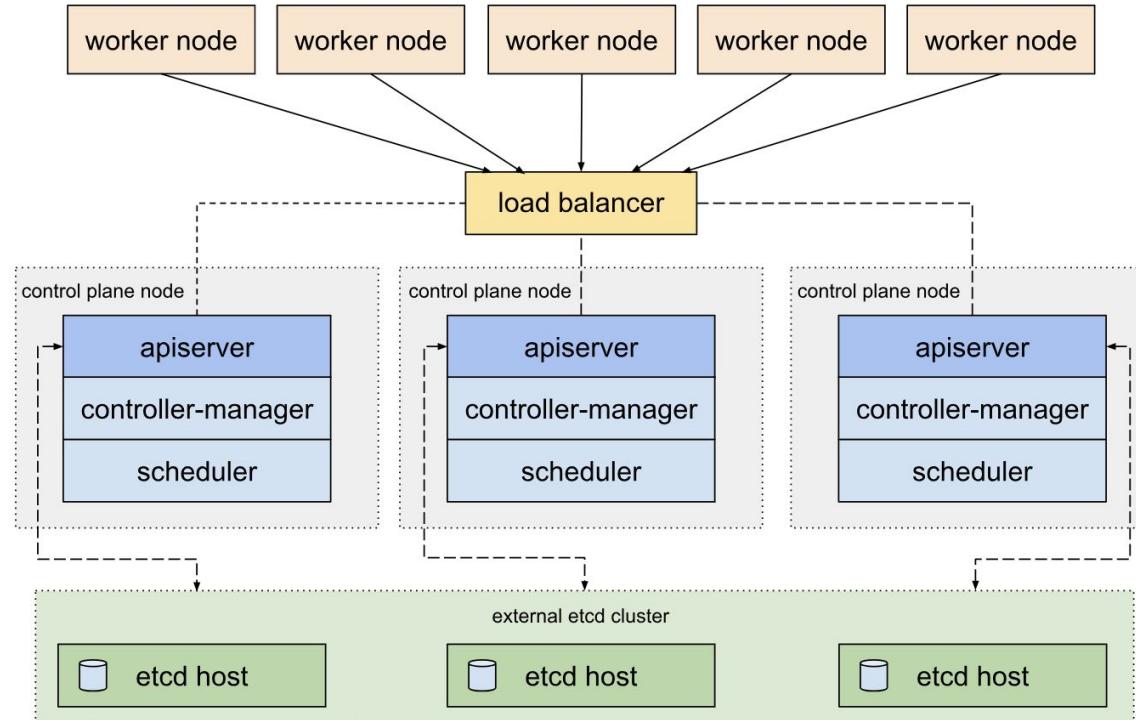
Options for Highly Available Topology

- Stacked etcd topology
- External etcd topology

Stacked etcd topology



External etcd topology



Fault Tolerance Table

| Cluster Size | Majority | Failure Tolerance |
|--------------|----------|-------------------|
| 1 | 1 | 0 |
| 2 | 2 | 0 |
| 3 | 2 | 1 |
| 4 | 3 | 1 |
| 5 | 3 | 2 |
| 6 | 4 | 2 |
| 7 | 4 | 3 |
| 8 | 5 | 3 |
| 9 | 5 | 4 |

Q & A

Feedback